

Algorithm Design and Analysis

CS
330

LECTURE 1

Analysis of Algorithms

- Course information
- Why study algorithms?
- Stable matching problem

Dora Erdos
Adam Smith

Algorithm Design and Analysis

Theoretical study of how to solve computational problems

- sorting a list of numbers
- finding a shortest route on a map
- scheduling when to work on homework
- answering web search queries

(*Computational problem*: precisely defined set of **inputs** and, for each input, **acceptable outputs**)

Algorithms

- Definition: Finite set of unambiguous instructions for solving a problem.
 - An algorithm is **correct** if on all legitimate inputs, it outputs the right answer in a finite amount of time
- Can be expressed as
 - pseudocode
 - flow charts
 - text in a natural language (e.g. English)
 - computer code

Etymology of “Algorithm”

*Abu Abdullah Muhammad ibn Musa
al-Khwarizmi (c. 780 -- 850 AD)*

- Persian (?) astronomer and mathematician
- “On calculating with hindu numerals”
a treatise in Arabic, 825
- “Agoritmi de numero Indorum”
translation into Latin, 12th century

Author’s name, mistaken for a plural noun, came to mean
“calculation methods”



Data Structures

- **Data structures** are ways to store information for which there are **algorithms** for performing particular operations (retrieving/manipulating information),
- e.g.
 - linked lists
 - hash tables
 - arrays
 - trees
 - heaps

Course information

- 1. Staff**
- 2. Lectures vs
labs**
- 3. Textbook**
- 4. Piazza & Gradescope**
- 5. Grading scheme**
- 6. Homework**
- 7. Programming exercises**
- 8. Name plate**
- 9. Collaboration policy**

Collaboration Policy

- You may discuss homework problems with up to 5 other students in the class
 - No outside help (including web sources)
 - **Write your own solution**
 - You must **acknowledge collaborators**
 - Or write “Collaborators: None”
 - Discussion is fun!
- You may spur discussion on Piazza by posting questions
 - The more specific your question, the better.
 - If in doubt, post to instructors.
- Absolutely no collaboration on exams.

Collaboration Tips

- Give everyone a chance to speak
 - If you're shy, speak up!
 - If you're outgoing, give others the floor.
- Listen
- Ask each other questions
- Be aware of your own biases, *e.g.*
 - Men have a tendency to talk over women
 - Nonnative speakers have it harder

Important Dates

- Wednesday, September 11: Homework 1 due
 - Weekly thereafter
- Thursday, October 24, 2019: Midterm
- Thursday, November 28: No class ☺
- Tuesday, December 9: Last class ☹
- Tuesday, December 17 (tentative): Final exam

Course Objectives

- classical algorithms and data structures
- analysis of algorithms
- standard design techniques

Why study algorithms?

- a *language* for talking about program behavior
- standard set of algorithms and design techniques
- feasibility (what can and cannot be done)
 - halting problem, NP-completeness
- analyzing correctness and resource usage
- successful companies (Google, Mapquest, Akamai)
- computation is fundamental to understanding the world
 - cells, brains, social networks, physical systems all can be viewed as computational devices
- IT IS **FUN!!!**

Example: Integer Addition

- Input: Positive integers $a, b \in \mathbb{N}$
- Output: $a + b$

Preschool-age children often use the following algorithm

```
Fingers( $a, b$ ):  
     $i \leftarrow 0$   
    while( $i \leq b$ ):  
         $a \leftarrow a + 1$   
         $i \leftarrow i + 1$ 
```

Running time scales with b .

If b has 100 digits, this could take as many as $10^{100} - 1$ executions of the while loop!

(The sun will die out in about 10^{27} cycles of a typical PC's processor)

But there is a much better algorithm!

Better Integer Addition

- You learned a better algorithm in grade school!
- Write a, b in decimal form
 - Suppose a has m digits and b has n digits
- Write one over the other and add columns from right to left, taking the carry with you as you go

$$\begin{array}{r} a_1 a_2 a_3 a_4 \cdots a_m \\ + b_1 b_2 b_3 \cdots b_n \\ \hline a + b \end{array}$$

- Can we analyze the running time?
 - Suppose operations on digits take one unit of time.
 - The algorithm will take something like $3(m + n)$ time units
 - (It depends on how exactly you handle carries.)
 - Adding two 100-digits numbers only takes humans a few minutes
 - You'll be done well before the sun dies out
 - Total time scales with $\log_{10} a + \log_{10} b$
 - It is exponentially faster than the fingers algorithm!

Performance isn't everything

- Typical goal: Find most space- and time-efficient algorithm for given problem.
- What else is important?
 - modularity
 - maintainability
 - functionality
 - robustness
 - user-friendliness
 - programmer time
 - simplicity
 - extensibility
 - reliability

Course Objectives

Material

- Classical algorithms
- Analysis of algorithms
- Design techniques

Skills

- Algorithmic Thinking
- Problem-solving & mathematical skills
- Technical writing

Prerequisites

- CS 112 (basic programming, data structures, sorting algorithms, “big O” notation)
- CS 131, MA 293 or equivalent (proofs, proofs, proofs, counting, and probability)

CS majors usually complete Group B coursework (any two of CS 132, CS 235 and CS 237) before taking this class

- (You may need to fill in some material if you haven’t.)

A first problem: Stable Matching

Matching Residents to Hospitals

- **Goal:** Given a set of preferences among hospitals and medical school students, design a **self-reinforcing** admissions process.
- **Unstable pair:** applicant x and hospital y are **unstable** if
 - x prefers y to its assigned hospital, and
 - y prefers x to one of its admitted students
- **Stable assignment:** no unstable pairs.
 - Individual self-interest will prevent any applicant/hospital deal from being made.
- 2012 Nobel prize in economics for work on matching algorithms (“mechanisms”)

Stable Matching Problem

- **Goal:** Given n men and n women, find a "suitable" matching.
 - Participants rate members of opposite sex.*
 - Each man lists women in order of preference from best to worst.
 - Each woman lists men in order of preference from best to worst.

	favorite ↓		least favorite ↓
	1 st	2 nd	3 rd
Xavier	Amy	Bertha	Clare
Yancey	Bertha	Amy	Clare
Zeus	Amy	Bertha	Clare

Men's Preference Profile

	favorite ↓		least favorite ↓
	1 st	2 nd	3 rd
Amy	Yancey	Xavier	Zeus
Bertha	Xavier	Yancey	Zeus
Clare	Xavier	Yancey	Zeus

Women's Preference Profile

* NOT a faithful model of real life relationships.

Stable Matching Problem

- **Unstable pair:** man m and woman w are **unstable** if
 - m prefers w to his assigned match, and
 - w prefers m to her assigned match
- Unstable pairs have an incentive to elope
- **Stable matching:** no unstable pairs.

	favorite ↓		least favorite ↓
	1 st	2 nd	3 rd
Xavier	Amy	Bertha	Clare
Yancey	Bertha	Amy	Clare
Zeus	Amy	Bertha	Clare

Men's Preference Profile

	favorite ↓		least favorite ↓
	1 st	2 nd	3 rd
Amy	Yancey	Xavier	Zeus
Bertha	Xavier	Yancey	Zeus
Clare	Xavier	Yancey	Zeus

Women's Preference Profile

Stable Matching Problem

- **Input:** preference lists of n men and n women
- **Goal:** find a stable matching if one exists

	favorite ↓		least favorite ↓
	1 st	2 nd	3 rd
Xavier	Amy	Bertha	Clare
Yancey	Bertha	Amy	Clare
Zeus	Amy	Bertha	Clare

Men's Preference Profile

	favorite ↓		least favorite ↓
	1 st	2 nd	3 rd
Amy	Yancey	Xavier	Zeus
Bertha	Xavier	Yancey	Zeus
Clare	Xavier	Yancey	Zeus

Women's Preference Profile

Stable Matching Problem

- Q. Is assignment X-C, Y-B, Z-A stable?

	favorite ↓	least favorite ↓	
	1 st	2 nd	3 rd
Xavier	Amy	Bertha	Clare
Yancey	Bertha	Amy	Clare
Zeus	Amy	Bertha	Clare

Men's Preference Profile

	favorite ↓	least favorite ↓	
	1 st	2 nd	3 rd
Amy	Yancey	Xavier	Zeus
Bertha	Xavier	Yancey	Zeus
Clare	Xavier	Yancey	Zeus

Women's Preference Profile

Stable Matching Problem

- Q. Is assignment X-C, Y-B, Z-A stable?
- A. No. Bertha and Xavier will hook up.

	favorite ↓	least favorite ↓	
	1 st	2 nd	3 rd
Xavier	Amy	Bertha	Clare
Yancey	Bertha	Amy	Clare
Zeus	Amy	Bertha	Clare

Men's Preference Profile

	favorite ↓	least favorite ↓	
	1 st	2 nd	3 rd
Amy	Yancey	Xavier	Zeus
Bertha	Xavier	Yancey	Zeus
Clare	Xavier	Yancey	Zeus

Women's Preference Profile

Stable Matching Problem

- Q. Is assignment X-A, Y-B, Z-C stable?

favorite
↓

least favorite
↓

	1 st	2 nd	3 rd
Xavier	Amy	Bertha	Clare
Yancey	Bertha	Amy	Clare
Zeus	Amy	Bertha	Clare

Men's Preference Profile

favorite
↓

least favorite
↓

	1 st	2 nd	3 rd
Amy	Yancey	Xavier	Zeus
Bertha	Xavier	Yancey	Zeus
Clare	Xavier	Yancey	Zeus

Women's Preference Profile

Stable Matching Problem

- Q. Is assignment X-A, Y-B, Z-C stable?
- A. Yes. X and Y got their first choice; Z is the last choice for every woman. No man can participate in an unstable pair.

	favorite ↓		least favorite ↓
	1 st	2 nd	3 rd
Xavier	Amy	Bertha	Clare
Yancey	Bertha	Amy	Clare
Zeus	Amy	Bertha	Clare

Men's Preference Profile

	favorite ↓	least favorite ↓	
	1 st	2 nd	3 rd
Amy	Yancey	Xavier	Zeus
Bertha	Xavier	Yancey	Zeus
Clare	Xavier	Yancey	Zeus

Women's Preference Profile

Existence of Stable Matching

- Q. Do stable matchings always exist?
- A. Not obvious a priori.

Stable Roommate Problem

- **Stable roommate problem**

- **2n** people; each person ranks others from **1** to **2n-1**.
- Assign roommate pairs so that no unstable pairs.

	<i>1st</i>	<i>2nd</i>	<i>3rd</i>
<i>Adam</i>	B	C	D
<i>Bob</i>	C	A	D
<i>Chris</i>	A	B	D
<i>Doofus</i>	A	B	C

A-B, C-D \Rightarrow B-C unstable
A-C, B-D \Rightarrow A-B unstable
A-D, B-C \Rightarrow A-C unstable

- **Observation.** Stable matchings do not always exist for stable roommate problem.

An Algorithm for Stable Matching

- Propose-and-reject algorithm. [Gale-Shapley 1962]

```
Initialize each person to be free.  
while (some man is free and hasn't proposed to every woman) {  
    Choose such a man m  
    w = 1st woman on m's list to whom m has not yet proposed  
    if (w is free)  
        assign m and w to be engaged  
    else if (w prefers m to her fiancé m')  
        assign m and w to be engaged, and m' to be free  
    else  
        w rejects m  
}
```

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Victor proposes to Bertha.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Victor proposes to Bertha.

- Bertha accepts since previously unmatched.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Wyatt proposes to Diane.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Wyatt proposes to Diane.

- Diane accepts since previously unmatched.

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Xavier proposes to Bertha.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Xavier proposes to Bertha.

- Bertha dumps Victor and accepts Xavier.

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Victor proposes to Amy.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Victor proposes to Amy.

- Amy accepts since previously unmatched.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Yancey proposes to Amy.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Yancey proposes to Amy.

- Amy rejects since she prefers Victor.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Yancey proposes to Diane.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Yancey proposes to Diane.

- Diane dumps Wyatt and accepts Yancey.

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Wyatt proposes to Bertha.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Wyatt proposes to Bertha.

- Bertha rejects since she prefers Xavier.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Wyatt proposes to Amy.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Wyatt proposes to Amy.

- Amy rejects since she prefers Victor.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Wyatt proposes to Clare.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Wyatt proposes to Clare.

- Clare accepts since previously unmatched.

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Zeus proposes to Bertha.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Zeus proposes to Bertha.

- Bertha rejects since she prefers Xavier.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Zeus proposes to Diane.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Zeus proposes to Diane.

- Diane rejects Yancey and accepts Zeus.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Yancey proposes to Clare.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Victor	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Yancey proposes to Clare.

- Clare rejects since she prefers Wyatt.

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Yancey proposes to Bertha.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Yancey proposes to Bertha.

- Bertha rejects since she prefers Xavier.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Victor	Wyatt
Erika	Yancey	Wyatt	Zeus	Victor	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Yancey proposes to Erika.

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Victor	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Yancey proposes to Erika.

- Erika accepts since previously unmatched.

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

STOP

- Everyone matched.
- Stable matching!

Propose-and-Reject Algorithm

- Propose-and-reject algorithm. [Gale-Shapley 1962]

```
Initialize each person to be free.  
while (some man is free and hasn't proposed to every woman) {  
    Choose such a man m  
    w = 1st woman on m's list to whom m has not yet proposed  
    if (w is free)  
        assign m and w to be engaged  
    else if (w prefers m to her fiancé m')  
        assign m and w to be engaged, and m' to be free  
    else  
        w rejects m  
}
```

Proving correctness of the algorithm

What do we need to prove?

For all correctly-formed inputs,

- Algorithm terminates and outputs a matching
- Matching is perfect
- Matching is stable

Proof of Correctness: Termination

- **Claim 1:** Algorithm terminates after at most n^2 iterations of while loop.

Proof:

- *Invariant:* Each time through the loop a man proposes to a new woman.
- There are only n^2 possible proposals. ▀

	1 st	2 nd	3 rd	4 th	5 th
Victor	A	B	C	D	E
Wyatt	B	C	D	A	E
Xavier	C	D	A	B	E
Yancey	D	A	B	C	E
Zeus	A	B	C	D	E

	1 st	2 nd	3 rd	4 th	5 th
Amy	W	X	Y	Z	V
Bertha	X	Y	Z	V	W
Clare	Y	Z	V	W	X
Diane	Z	V	W	X	Y
Erika	V	W	X	Y	Z

An instance where $n(n-1) + 1$ proposals required

Propose-and-Reject Algorithm

- **Claim 2.** Men propose to women in decreasing order of preference.
- **Claim 3.** Once a woman is matched, she never becomes unmatched; she only "trades up."

Men's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Victor	Bertha	Amy	Diane	Erika	Clare
Wyatt	Diane	Bertha	Amy	Clare	Erika
Xavier	Bertha	Erika	Clare	Diane	Amy
Yancey	Amy	Diane	Clare	Bertha	Erika
Zeus	Bertha	Diane	Amy	Erika	Clare

Women's Preference Profile

	0 th	1 st	2 nd	3 rd	4 th
Amy	Zeus	Victor	Wyatt	Yancey	Xavier
Bertha	Xavier	Wyatt	Yancey	Victor	Zeus
Clare	Wyatt	Xavier	Yancey	Zeus	Victor
Diane	Victor	Zeus	Yancey	Xavier	Wyatt
Erika	Yancey	Wyatt	Zeus	Xavier	Victor

Proof of Correctness: Perfection

- **Claim 4:** All men and women get matched.
- **Proof:** (by contradiction)
 - Suppose, for sake of contradiction, some guy, say Zeus, is not matched upon termination of algorithm.
 - Then some woman, say Amy, is not matched upon termination because the algorithm maintains a partial matching (one-to-one correspondence) at all times.
 - By Claim 3, Amy was never proposed to.
 - But Zeus proposes to everyone, since he ends up unmatched.
 - Thus, there cannot exist unmatched people. ▀

Proof of Correctness: Stability

- **Claim 5:** No unstable pairs.

- **Proof: (by contradiction)**

- Suppose A-Z is an unstable pair: they prefer each other to their partners in Gale-Shapley matching S^* .

- *Case 1:* Z never proposed to A.

- Z prefers his GS partner to A. ← men propose in decreasing order of preference

- A-Z is stable.

- *Case 2:* Z proposed to A.

- A rejected Z (right away or later)

- A prefers her GS partner to Z. ← women only trade up

- A-Z is stable.

- In either case A-Z is stable, a contradiction. ▀

Stable Roommate Problem

- **Stable roommate problem**

- **$2n$** people; each person ranks others from **1** to **$2n-1$** .
- Assign roommate pairs so that no unstable pairs.

	<i>1st</i>	<i>2nd</i>	<i>3rd</i>
<i>Adam</i>	B	C	D
<i>Bob</i>	C	A	D
<i>Chris</i>	A	B	D
<i>Doofus</i>	A	B	C

$A-B, C-D \Rightarrow B-C$ unstable
 $A-C, B-D \Rightarrow A-B$ unstable
 $A-D, B-C \Rightarrow A-C$ unstable

- **Exercise.** Where does the correctness proof break down for the roommates version?

Two more exercises

- Does the order in which unmatched men are selected change the final matching?
- If women propose (instead of men), what changes?
Is it better to propose or to receive proposals?

Algorithm Design and Analysis

CS
330

LECTURES 2 & 3 **Analysis of Algorithms**

- Asymptotic notation
- Basic data structures

**Dora Erdos
Adam Smith**

Useful Functions and Asymptotics

Permutations and combinations

- Factorial: “ n factorial”

$$n! = n \cdot (n - 1) \cdot \dots \cdot 2 \cdot 1$$

= number of permutations of $\{1, \dots, n\}$

- Combinations: “ n choose k ”

$$\binom{n}{k} = \frac{n \times (n - 1) \times \dots \times (n - k + 1)}{k \times (k - 1) \times \dots \times 2 \times 1} = \frac{n!}{k!(n - k)!}$$

= number of ways of choosing an unordered subset of k items in $\{1, \dots, n\}$ without repetition

Review Question

- In how many ways can we select two disjoint subsets of $\{1, \dots, n\}$, of size k and m , respectively?
- **Answer:**

$$\binom{n}{k} \binom{n-k}{m} = \binom{n}{m} \binom{n-m}{k} = \binom{n}{m+k} \binom{m+k}{k}$$

Asymptotic notation

O -notation (upper bounds):

$f(n) = O(g(n))$ means

there exist constants $c > 0$, $n_0 > 0$ such that $0 \leq f(n) \leq cg(n)$ for all $n \geq n_0$.

EXAMPLE: $2n^2 = O(n^3)$ $(c = 1, n_0 = 2)$



*functions,
not values*

Asymptotic Notation

- **One-sided equality:** $T(n) = O(f(n))$.

➤ Not transitive:

- $f(n) = 5n^3; g(n) = 3n^2$
- $f(n) = O(n^3) = g(n)$
- but $f(n) \neq g(n)$.

➤ Alternative notation: $T(n) \in O(f(n))$.

Set Definition

$O(g(n)) = \{ f(n) : \text{there exist constants } c > 0, n_0 > 0 \text{ such that } 0 \leq f(n) \leq cg(n) \text{ for all } n \geq n_0 \}$

EXAMPLE: $2n^2 \in O(n^3)$

(*Logicians:* $\lambda n.2n^2 \in O(\lambda n.n^3)$, but it's convenient to be sloppy, as long as we understand what's *really* going on.)

Examples

- $10^6 n^3 + 2n^2 - n + 10 = O(n^3)$
- $n^{1/2} + \log n = O(n^{1/2})$
- $n (\log n + \sqrt{n}) = O(n^{3/2})$
- $n = O(n^2)$

Ω -notation (lower bounds)

O -notation is an *upper-bound* notation. It makes no sense to say $f(n)$ is at least $O(n^2)$.

$\Omega(g(n)) = \{f(n) : \text{there exist constants } c > 0, n_0 > 0 \text{ such that } 0 \leq cg(n) \leq f(n) \text{ for all } n \geq n_0\}$

EXAMPLE: $\sqrt{n} = \Omega(\log n)$ ($c = 1, n_0 = 16$)

Ω -notation (lower bounds)

- **Be careful:** “Any comparison-based sorting algorithm requires at least $O(n \log n)$ comparisons.”
 - Meaningless!
 - Use Ω for lower bounds.

Θ -notation (tight bounds)

$$\Theta(g(n)) = O(g(n)) \cap \Omega(g(n))$$

EXAMPLE: $\frac{1}{2}n^2 - 2n = \Theta(n^2)$

Polynomials are simple:

$$a_d n^d + a_{d-1} n^{d-1} + \dots + a_1 n + a_0 = \Theta(n^d)$$

\mathbf{o} -notation and ω -notation

O -notation and Ω -notation are like \leq and \geq .
 o -notation and ω -notation are like $<$ and $>$.

$o(g(n)) = \{ f(n) : \text{for every constant } c > 0, \text{ there is a constant } n_0 > 0 \text{ such that } 0 \leq f(n) < cg(n) \text{ for all } n \geq n_0 \}$

EXAMPLE: $2n^2 = o(n^3)$ ($n_0 = 2/c$)

Overview of Asymptotic Notation

Notation	... means ...	Think...	E.g.	$\lim f(n)/g(n)$
$f(n)=O(g(n))$	$\exists c > 0, n_0 > 0$ $\forall n > n_0 :$ $0 \leq f(n) < cg(n)$	Upper bound	$100n^2 = O(n^3)$	If it exists, it is $< \infty$
$f(n)=\Omega(g(n))$	$\exists c > 0, n_0 > 0, \forall n > n_0 :$ $0 \leq cg(n) < f(n)$	Lower bound	$2^n = \Omega(n^{100})$	If it exists, it is > 0
$f(n)=\Theta(g(n))$	both of the above: $f=\Omega(g)$ and $f=O(g)$	Tight bound	$\log(n!) = \Theta(n \log n)$	If it exists, it is > 0 and $< \infty$
$f(n)=o(g(n))$	$\forall c > 0, \exists n_0 > 0, \forall n > n_0 :$ $0 \leq f(n) < cg(n)$	Strict upper bound	$n^2 = o(2^n)$	Limit exists, $= 0$
$f(n)=\omega(g(n))$	$\forall c > 0, \exists n_0 > 0, \forall n > n_0 :$ $0 \leq cg(n) < f(n)$	Strict lower bound	$n^2 = \omega(\log n)$	Limit exists, $= \infty$

Common Functions: Asymptotic Bounds

- **Polynomials.** $a_0 + a_1n + \dots + a_dn^d$ is $\Theta(n^d)$ if $a_d > 0$.
- **Polynomial time.** Running time is $O(n^d)$ for some constant d independent of the input size n .
- **Logarithms.** $\log_a n = \Theta(\log_b n)$ for all constants $a, b > 0$.

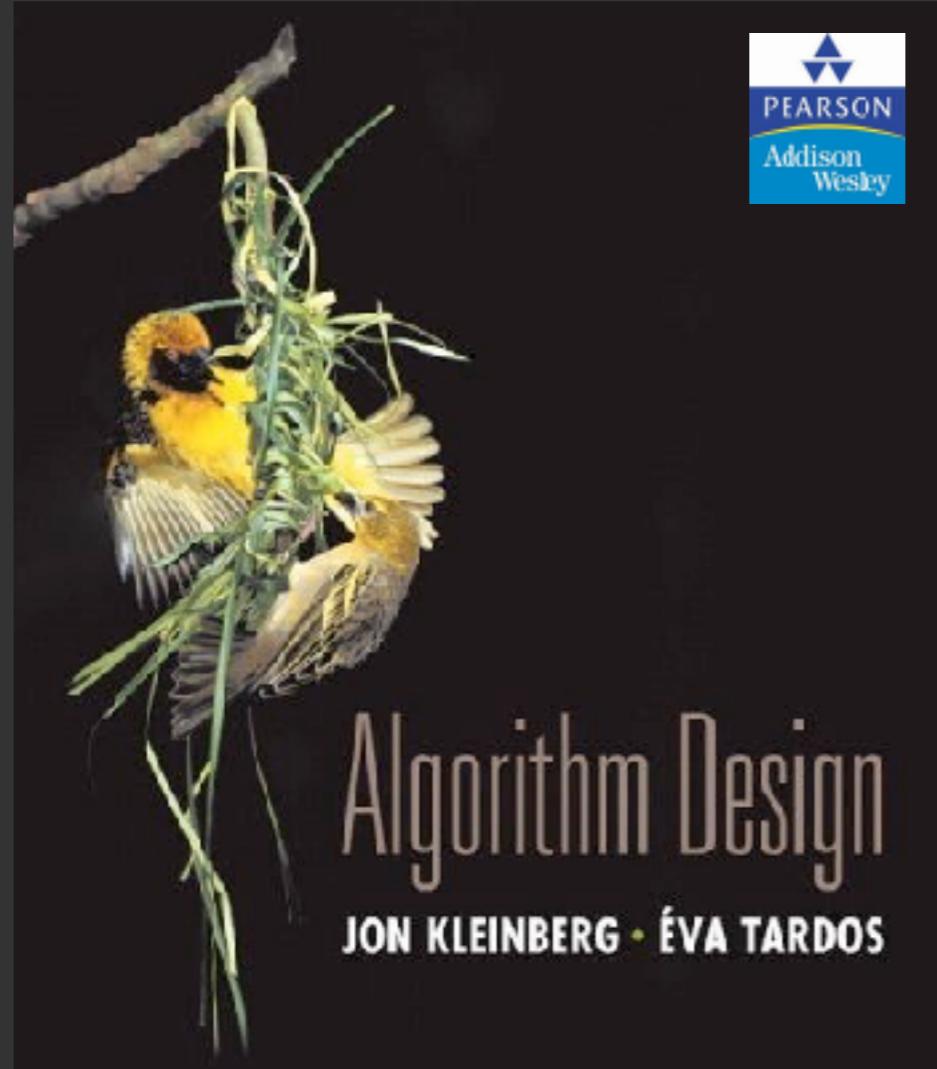
$$\text{For every } x > 0, \log n = o(n^x).$$

↑
can avoid specifying the base ↓
log grows slower than every polynomial

- **Exponentials.** For all $r > 1$ and all $d > 0$, $n^d = o(r^n)$.
- **Factorial.** By Sterling's formula,

$$n! = (\sqrt{2\pi n}) \left(\frac{n}{e}\right)^n (1 + o(1)) = 2^{\Theta(n \log n)}$$

↑
Every polynomial grows slower than every exponential
grows faster than every exponential



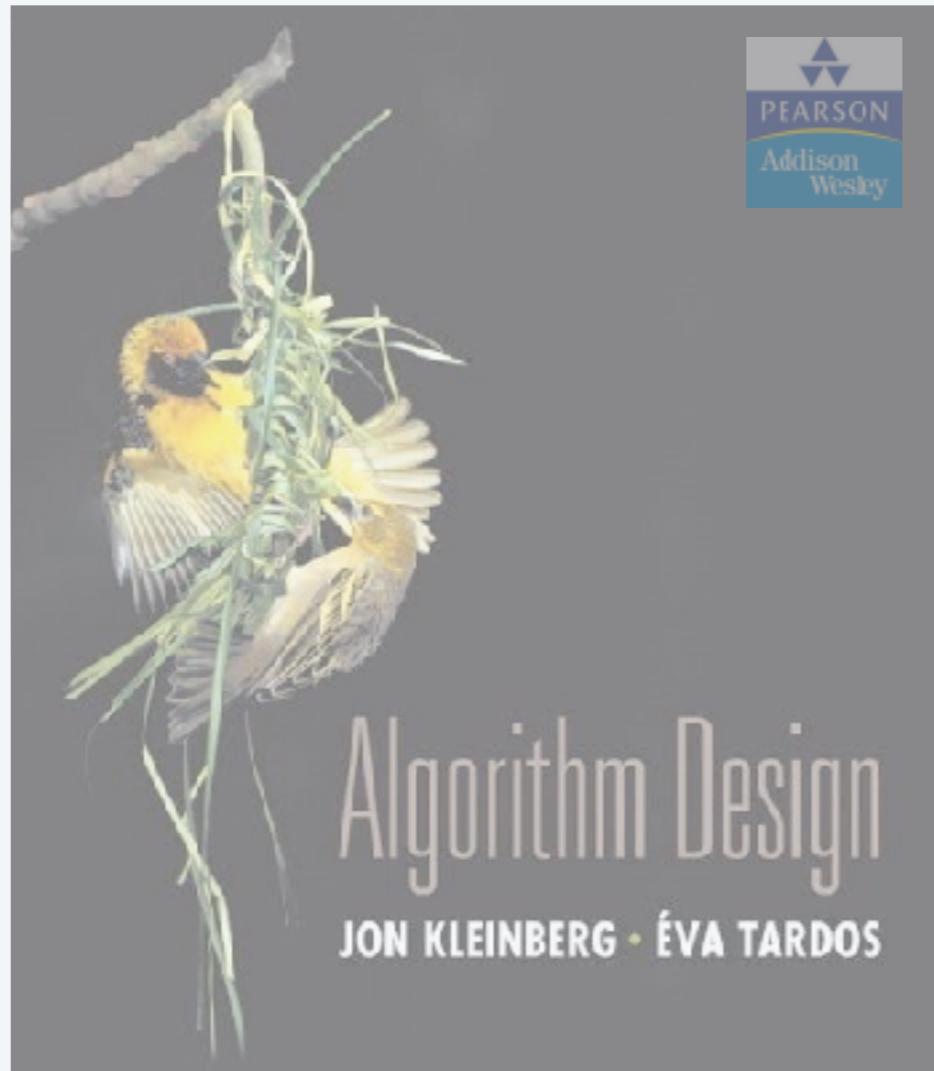
3. GRAPHS

- ▶ *basic definitions and applications*
- ▶ *graph connectivity and graph traversal*
- ▶ *testing bipartiteness*
- ▶ *connectivity in directed graphs*
- ▶ *DAGs and topological ordering*

Lecture slides by Kevin Wayne

Copyright © 2005 Pearson–Addison Wesley

<http://www.cs.princeton.edu/~wayne/kleinberg-tardos>



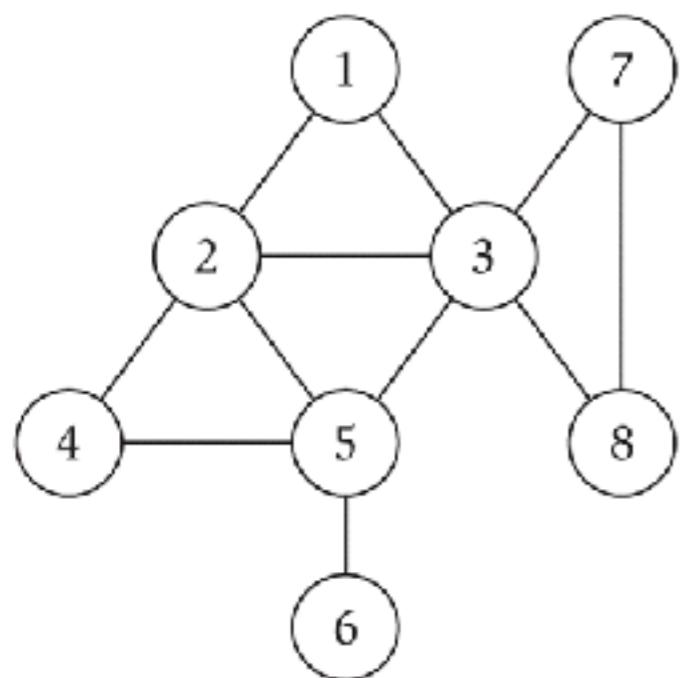
3. GRAPHS

- ▶ *basic definitions and applications*
- ▶ *graph connectivity and graph traversal*
- ▶ *testing bipartiteness*
- ▶ *connectivity in directed graphs*
- ▶ *DAGs and topological ordering*

Undirected graphs

Notation. $G = (V, E)$

- V = nodes (or vertices).
- E = edges (or arcs) between pairs of nodes.
- Captures pairwise relationship between objects.
- Graph size parameters: $n = |V|, m = |E|$.

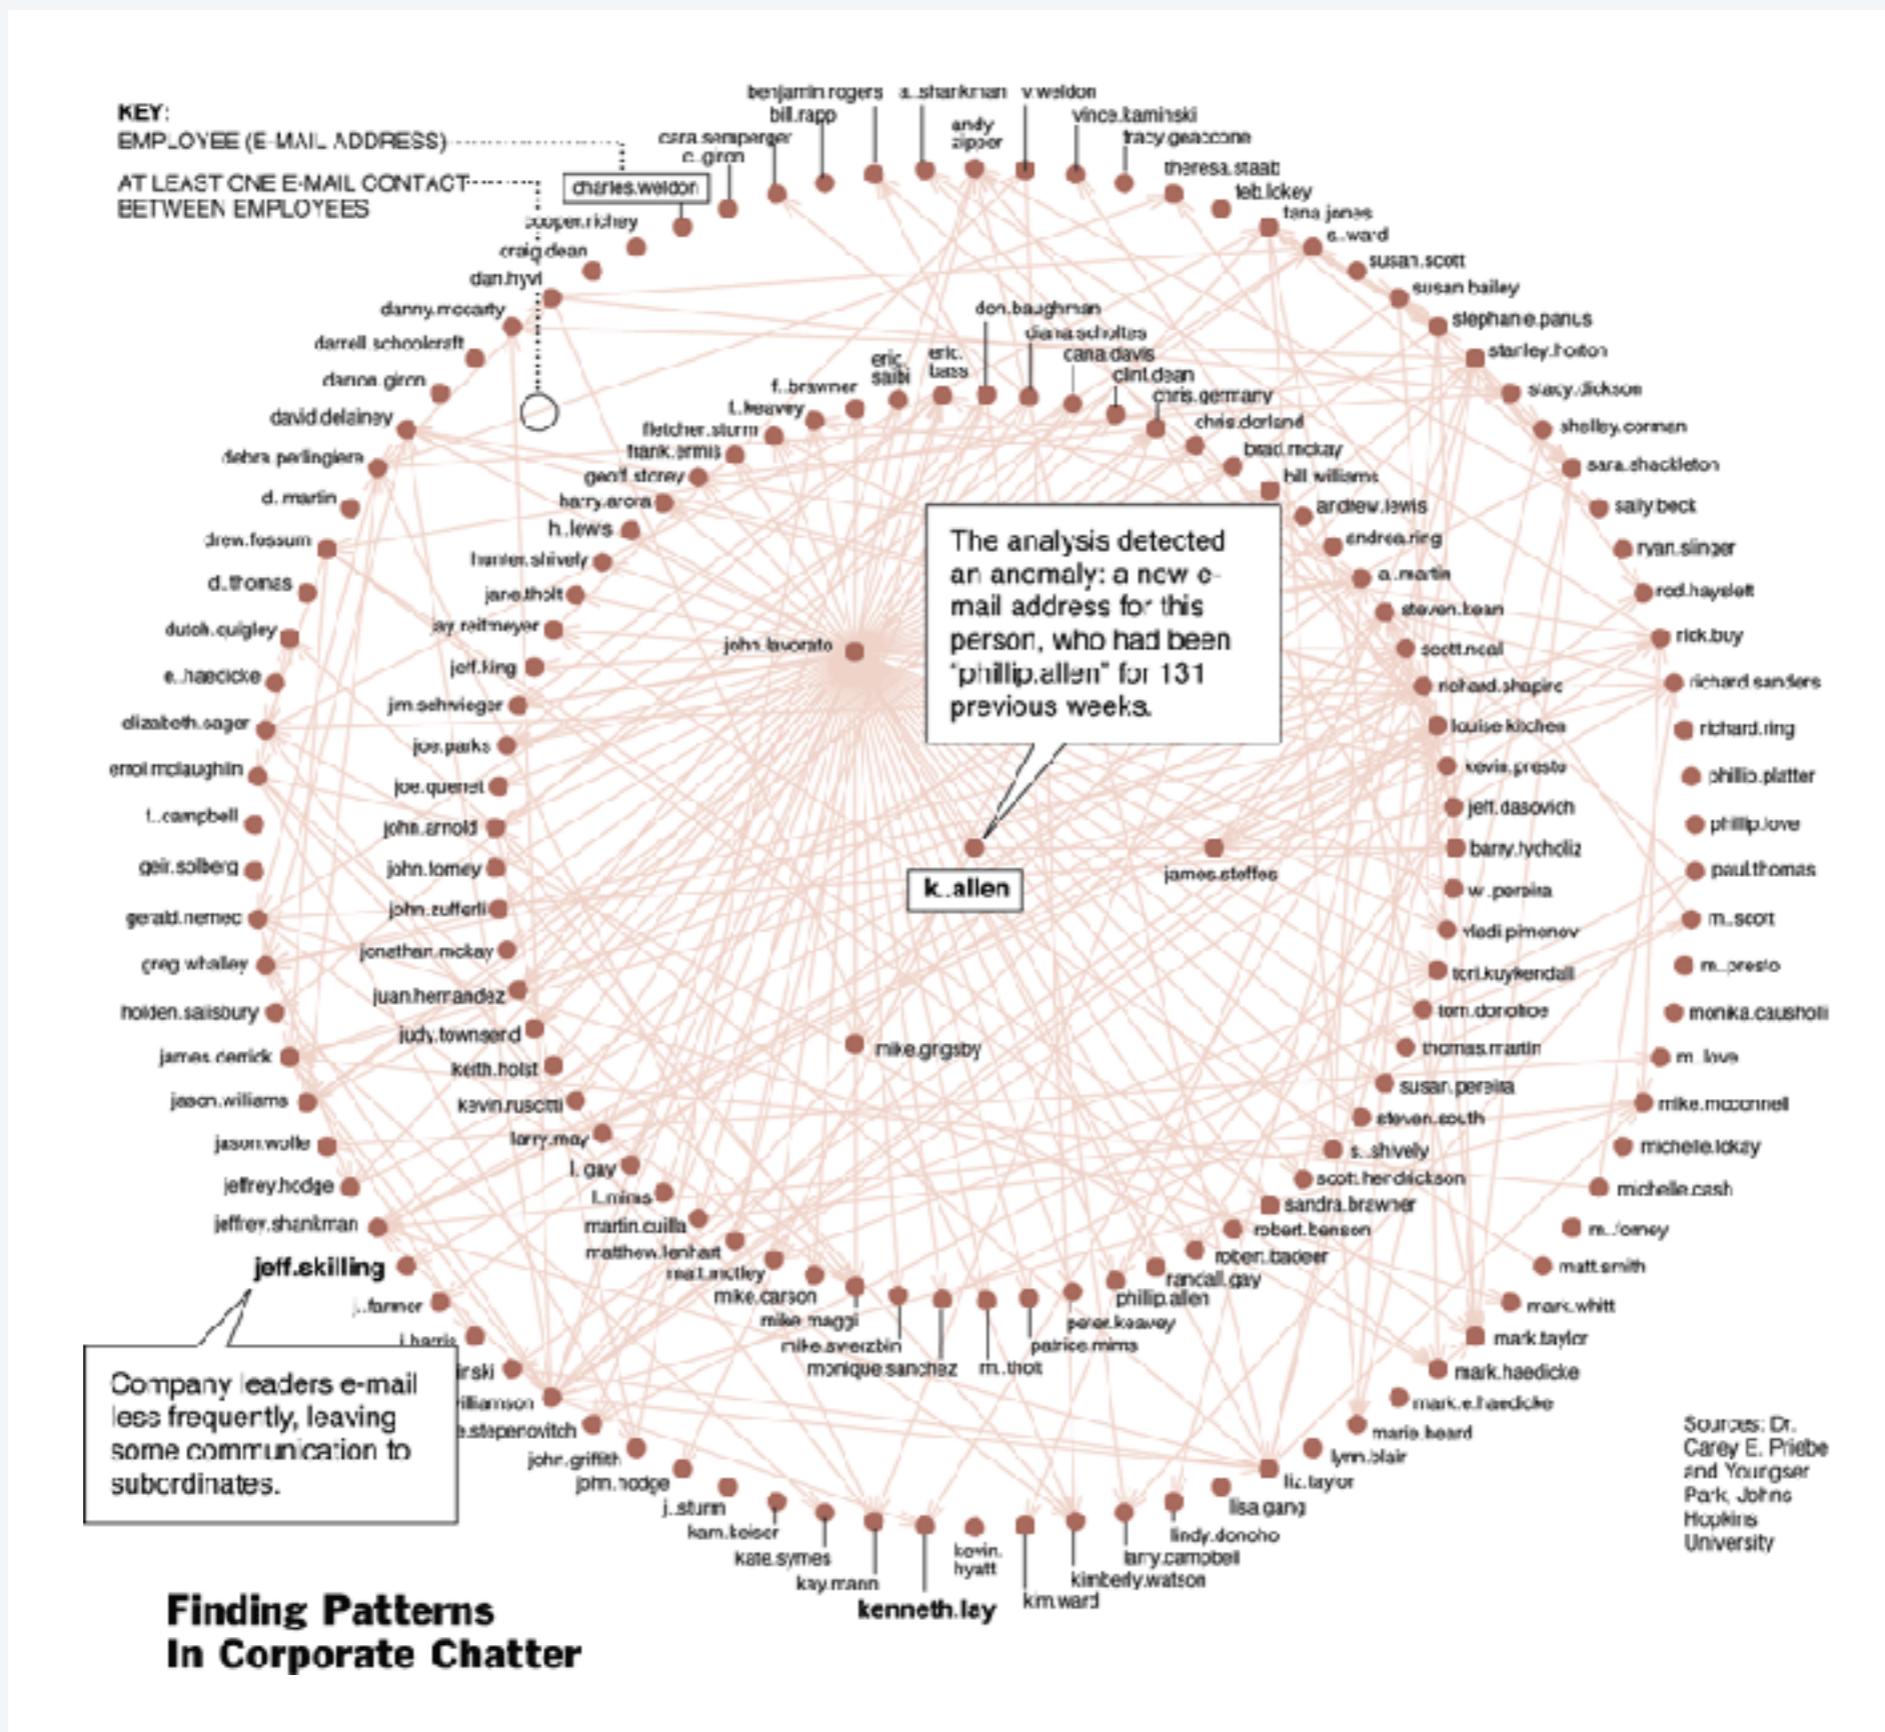


$$V = \{ 1, 2, 3, 4, 5, 6, 7, 8 \}$$

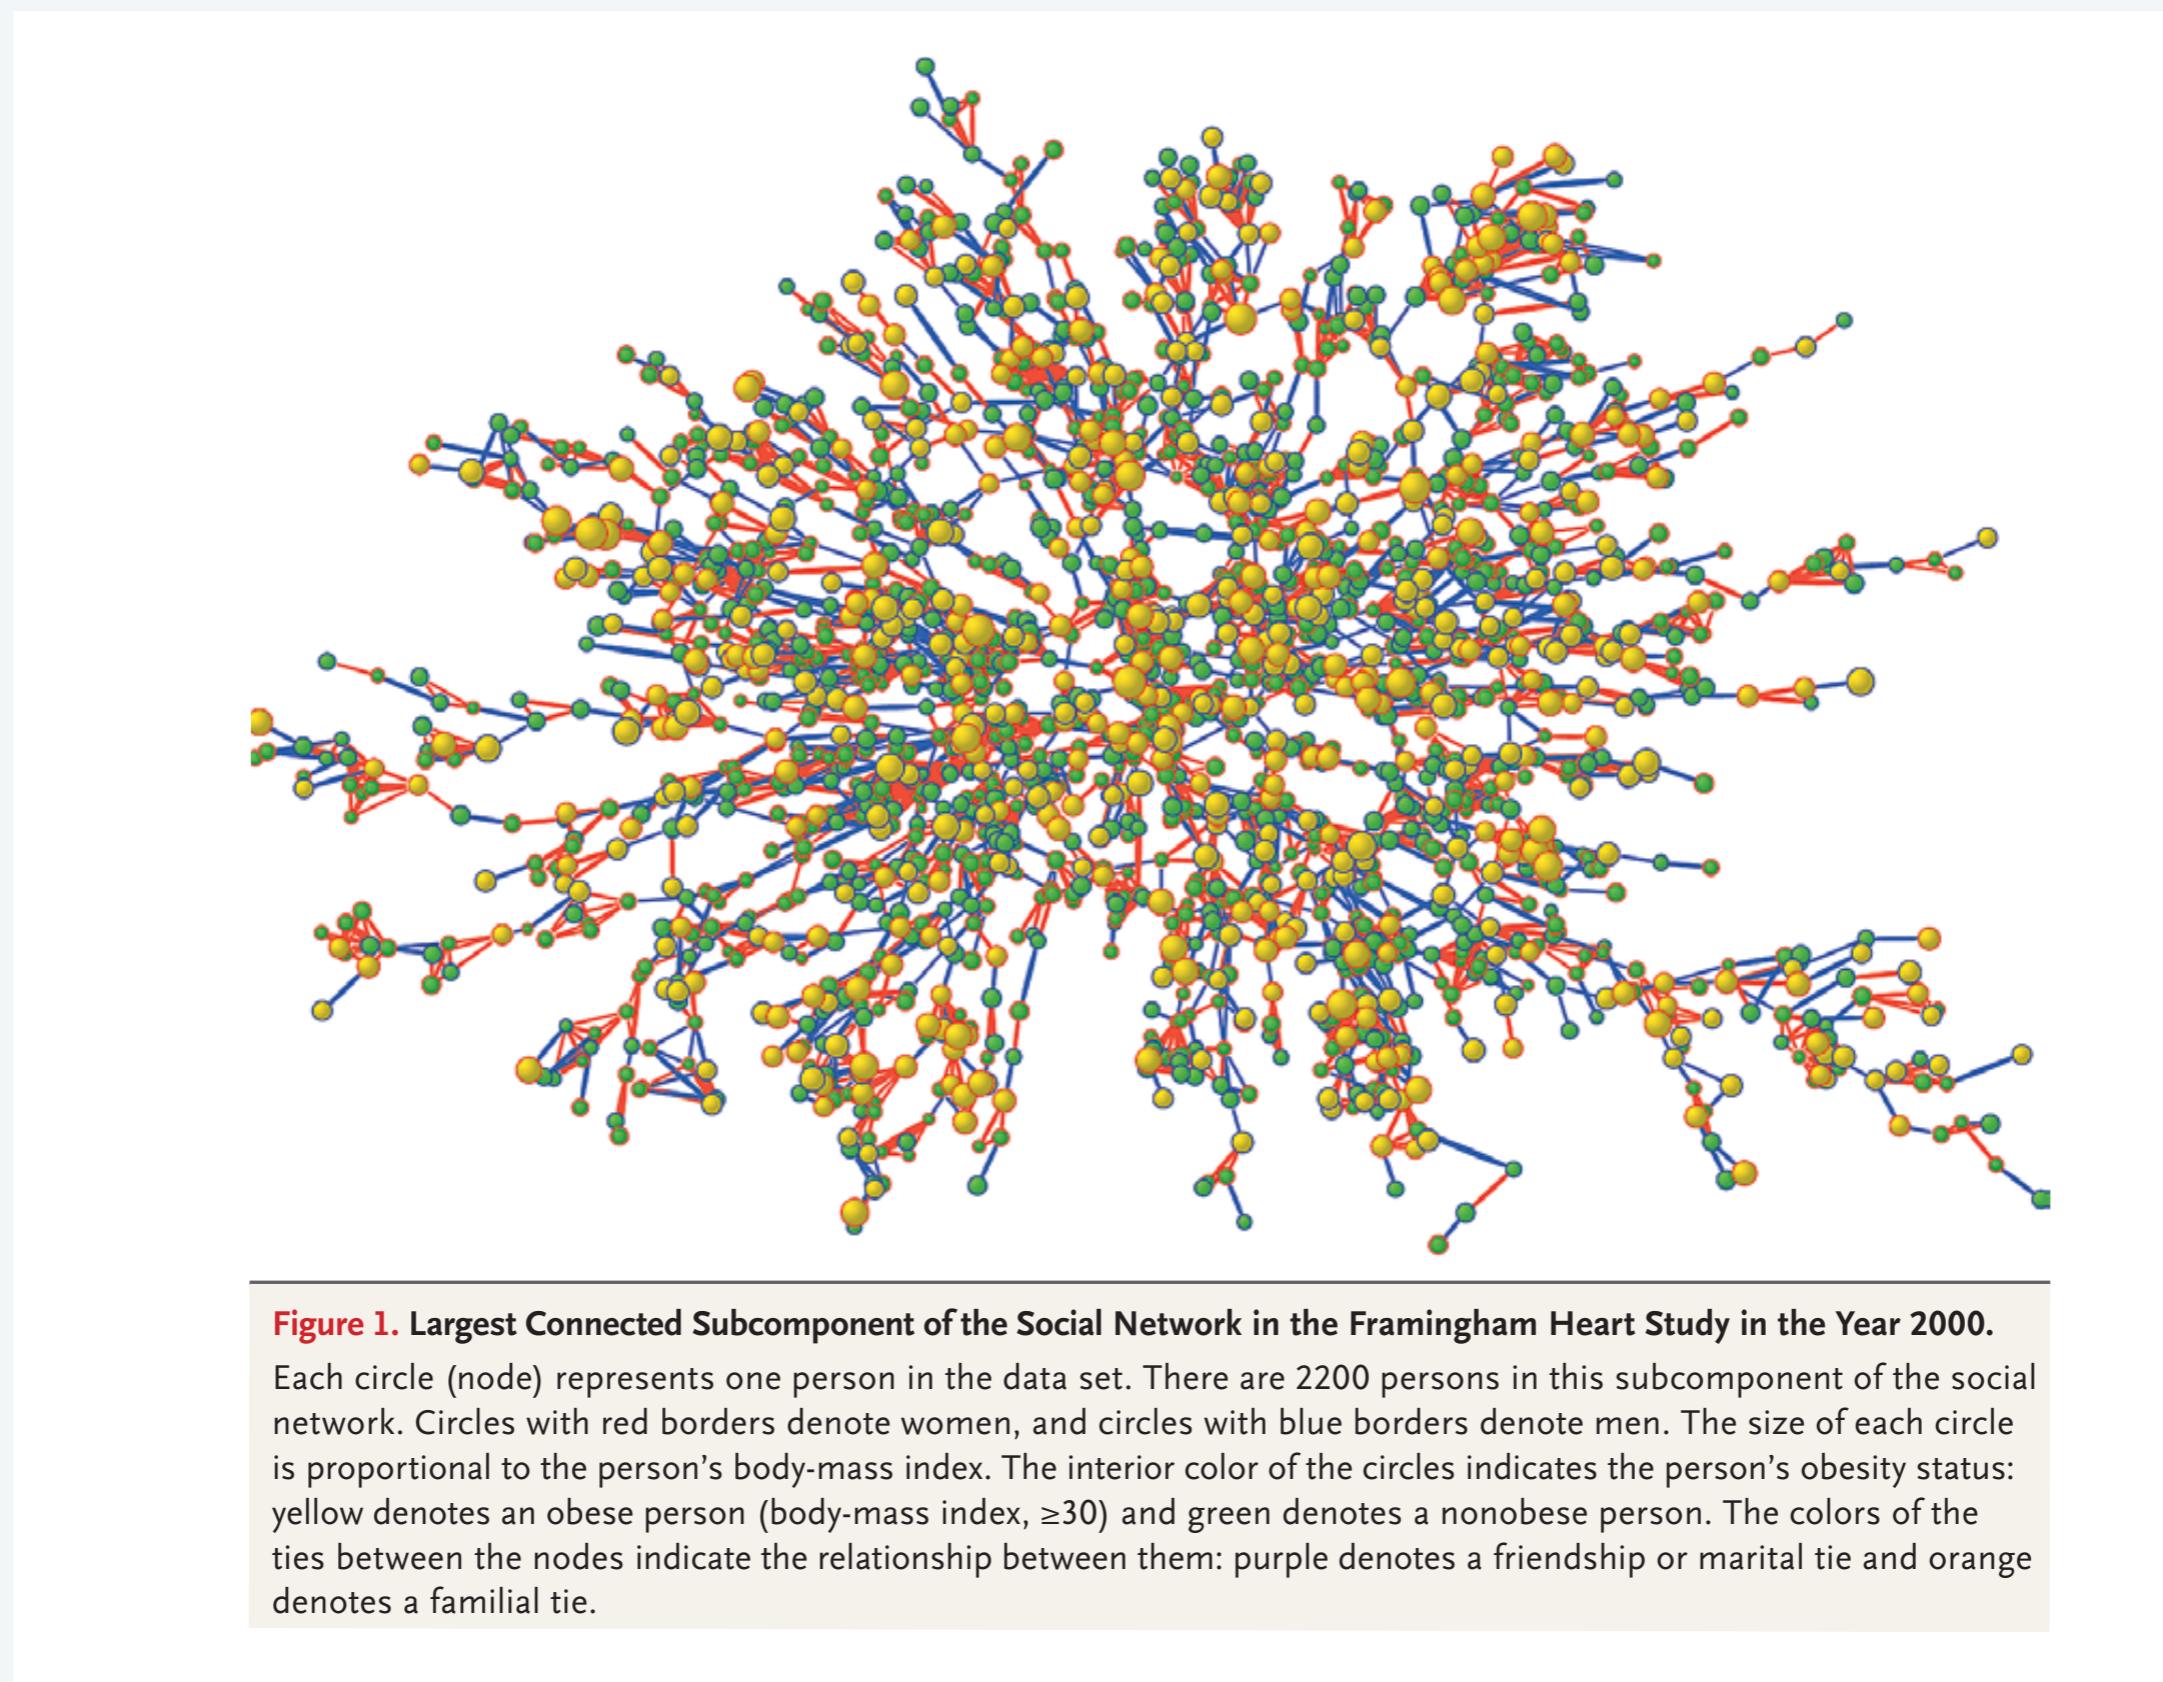
$$E = \{ 1-2, 1-3, 2-3, 2-4, 2-5, 3-5, 3-7, 3-8, 4-5, 5-6, 7-8 \}$$

$$m = 11, n = 8$$

One week of Enron emails



Framingham heart study



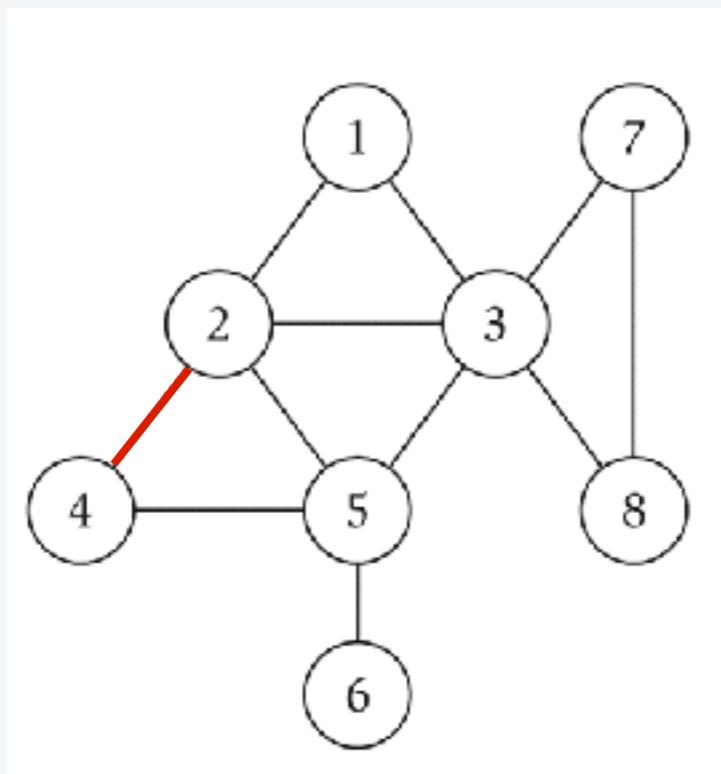
Some graph applications

graph	node	edge
communication	telephone, computer	fiber optic cable
circuit	gate, register, processor	wire
mechanical	joint	rod, beam, spring
financial	stock, currency	transactions
transportation	street intersection, airport	highway, airway route
internet	class C network	connection
game	board position	legal move
social relationship	person, actor	friendship, movie cast
neural network	neuron	synapse
protein network	protein	protein-protein interaction
molecule	atom	bond

Graph representation: adjacency matrix

Adjacency matrix. n -by- n matrix with $A_{uv} = 1$ if (u, v) is an edge.

- Two representations of each edge.
- Space proportional to n^2 .
- Checking if (u, v) is an edge takes $\Theta(1)$ time.
- Identifying all edges takes $\Theta(n^2)$ time.



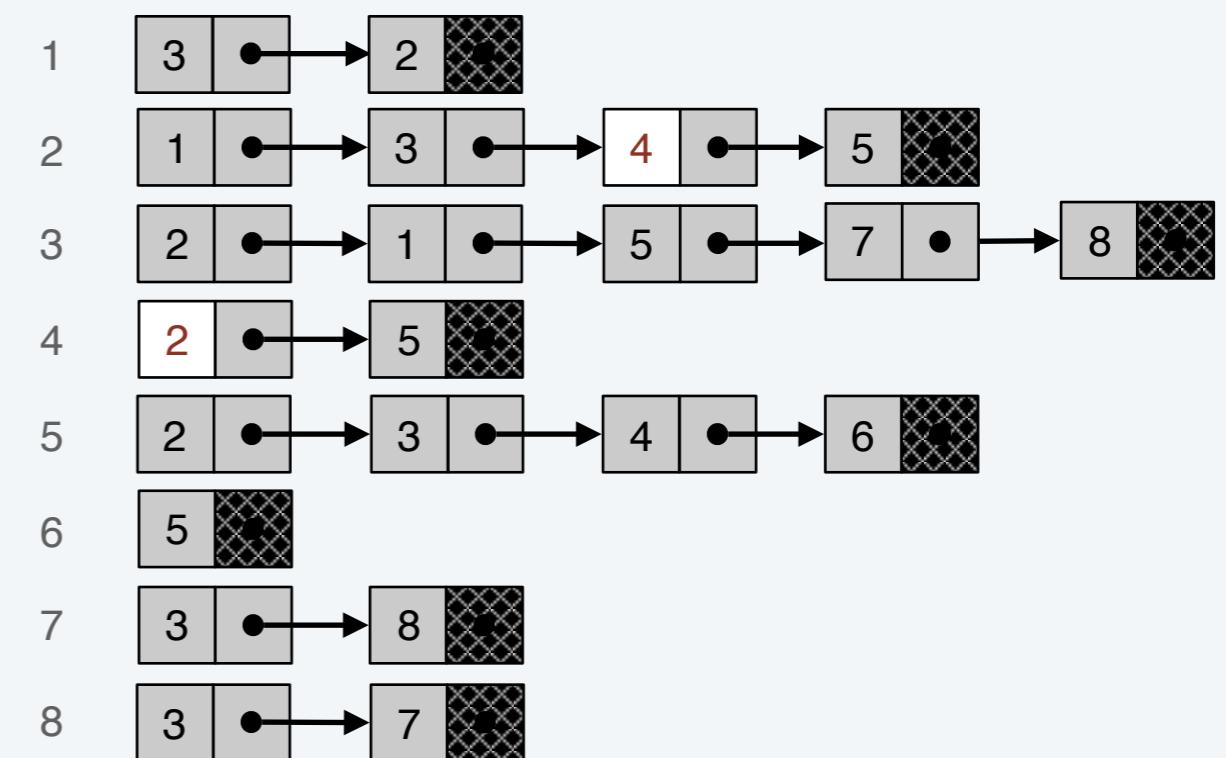
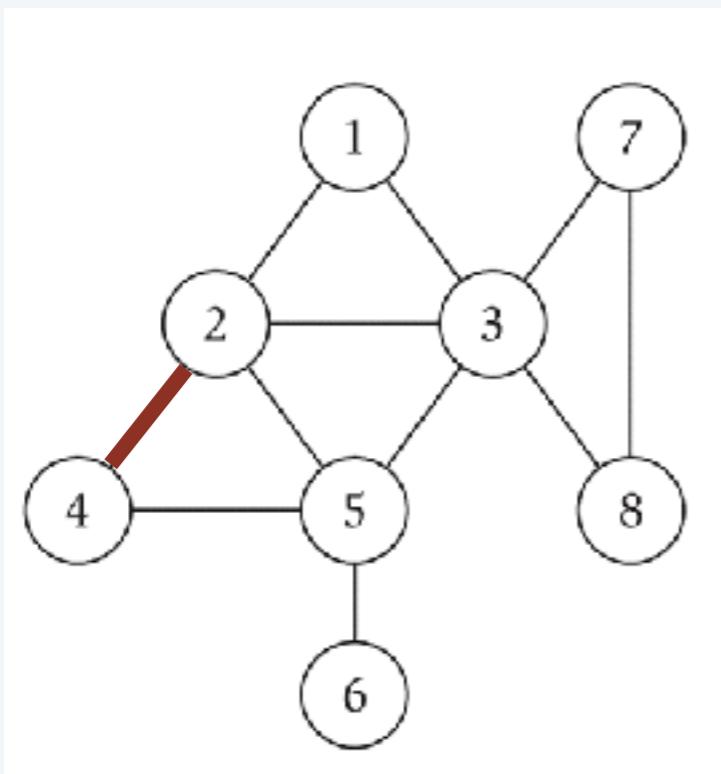
	1	2	3	4	5	6	7	8
1	0	1	1	0	0	0	0	0
2	1	0	1	1	1	0	0	0
3	1	1	0	0	1	0	1	1
4	0	1	0	0	1	0	0	0
5	0	1	1	1	0	1	0	0
6	0	0	0	0	1	0	0	0
7	0	0	1	0	0	0	1	0
8	0	0	1	0	0	0	1	0

Graph representation: adjacency lists

Adjacency lists. Node-indexed array of lists.

- Two representations of each edge.
- Space is $\Theta(m + n)$.
- Checking if (u, v) is an edge takes $O(\text{degree}(u))$ time.
- Identifying all edges takes $\Theta(m + n)$ time.

degree = number of neighbors of u

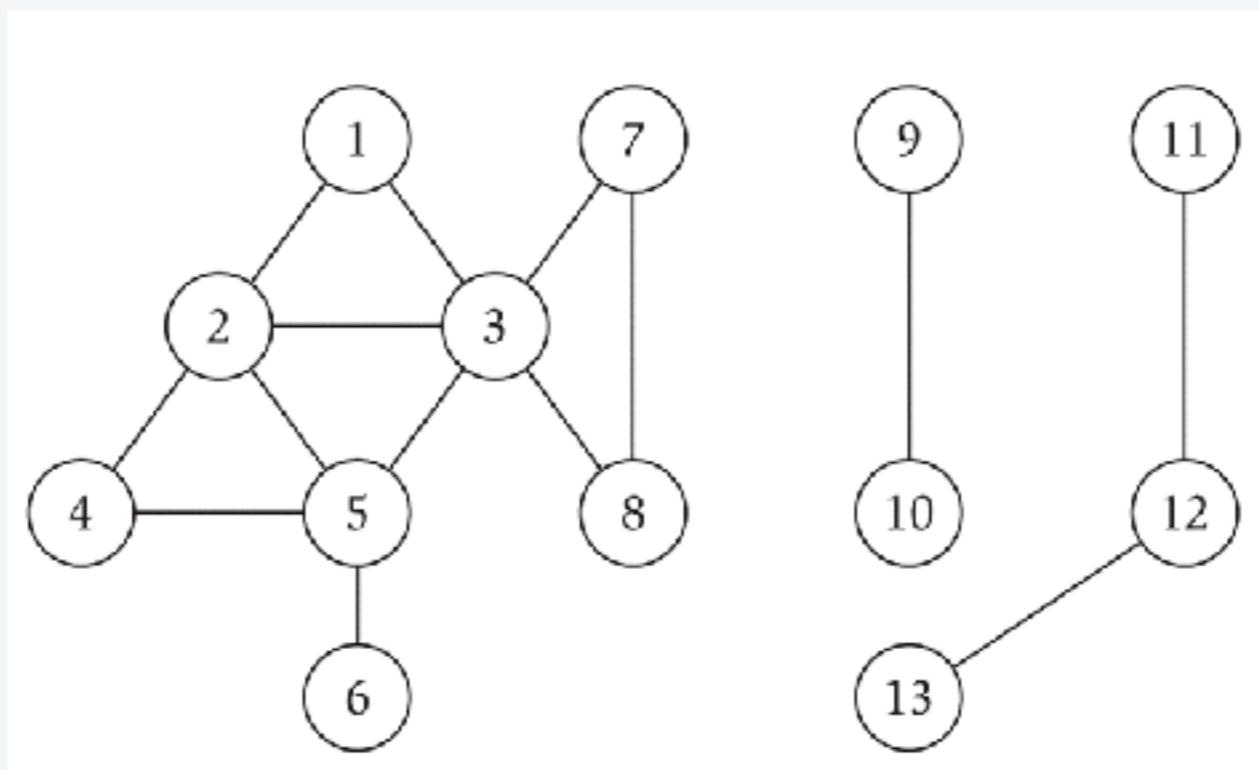


Paths and connectivity

Def. A **path** in an undirected graph $G = (V, E)$ is a sequence of nodes v_1, v_2, \dots, v_k with the property that each consecutive pair v_{i-1}, v_i is joined by an edge in E .

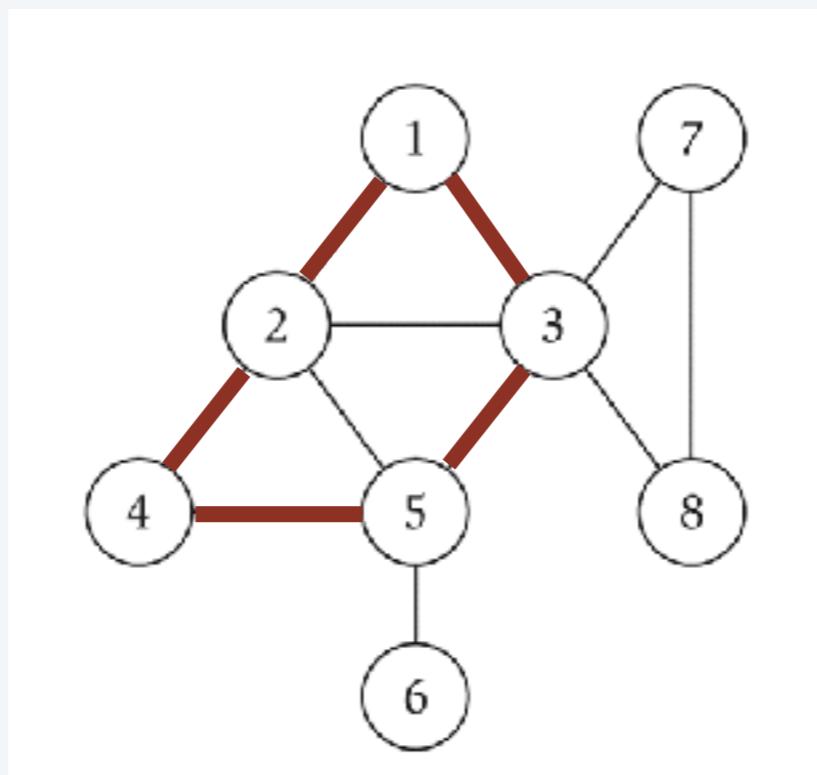
Def. A path is **simple** if all nodes are distinct.

Def. An undirected graph is **connected** if for every pair of nodes u and v , there is a path between u and v .



Cycles

Def. A **cycle** is a path v_1, v_2, \dots, v_k in which $v_1 = v_k$, $k > 2$, and the first $k - 1$ nodes are all distinct.



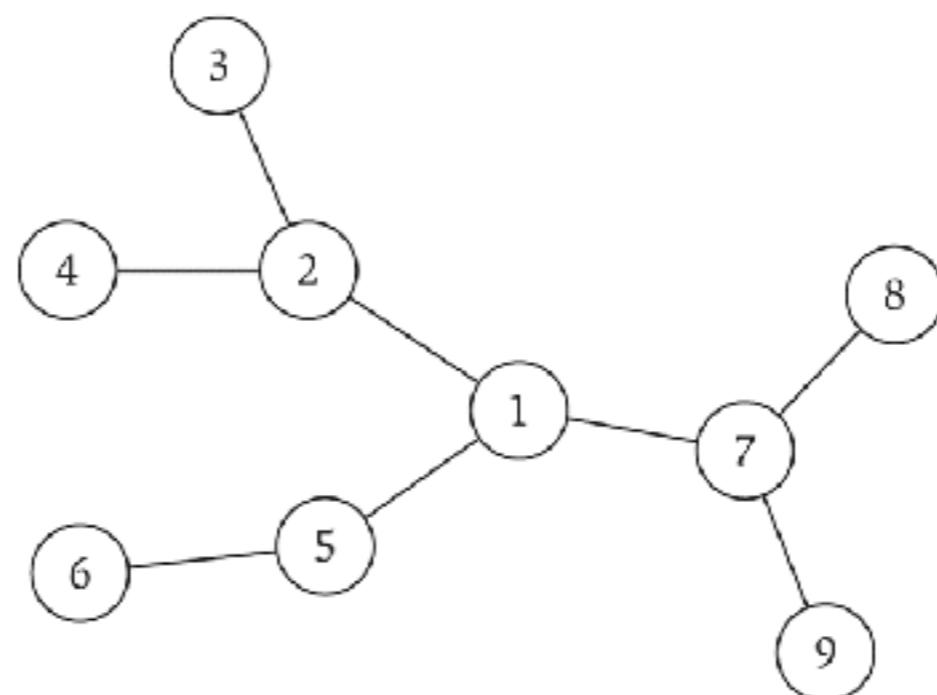
cycle C = 1-2-4-5-3-1

Trees

Def. An undirected graph is a **tree** if it is connected and does not contain a cycle.

Theorem. Let G be an undirected graph on n nodes. Any two of the following statements imply the third:

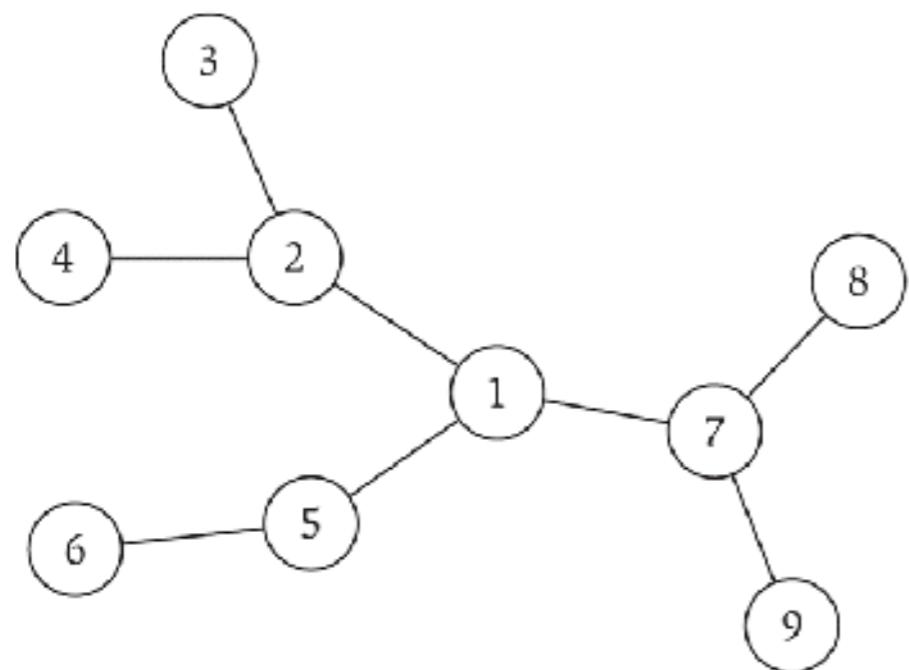
- G is connected.
- G does not contain a cycle.
- G has $n - 1$ edges.



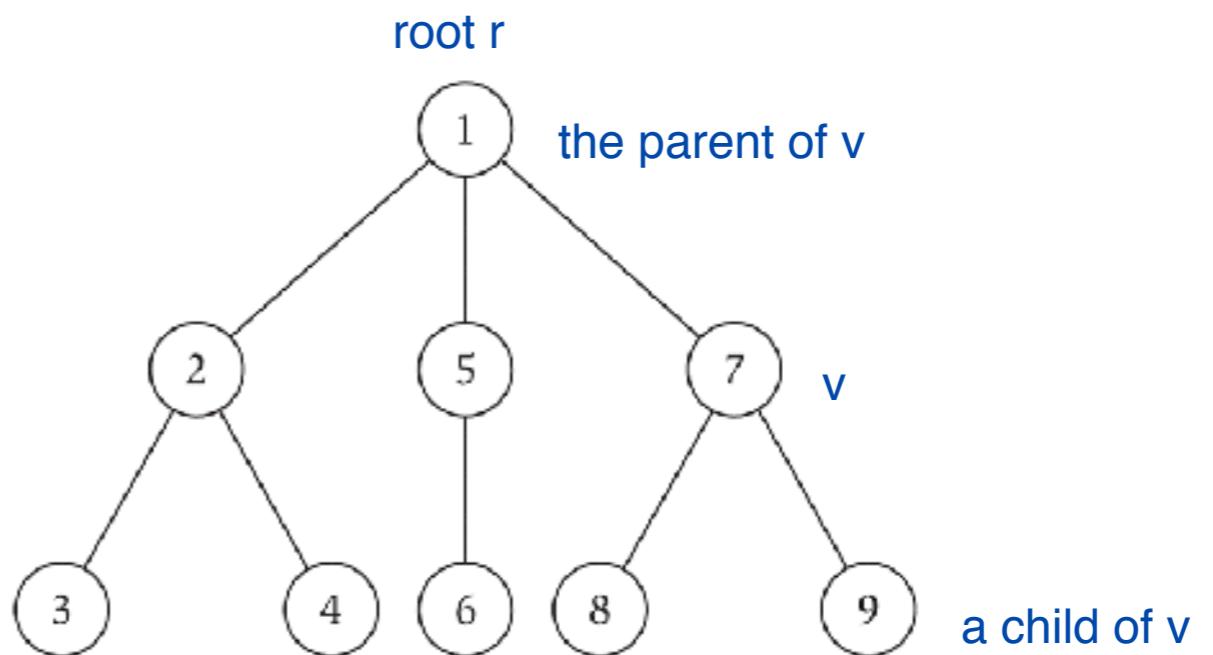
Rooted trees

Given a tree T , choose a root node r and orient each edge away from r .

Importance. Models hierarchical structure.



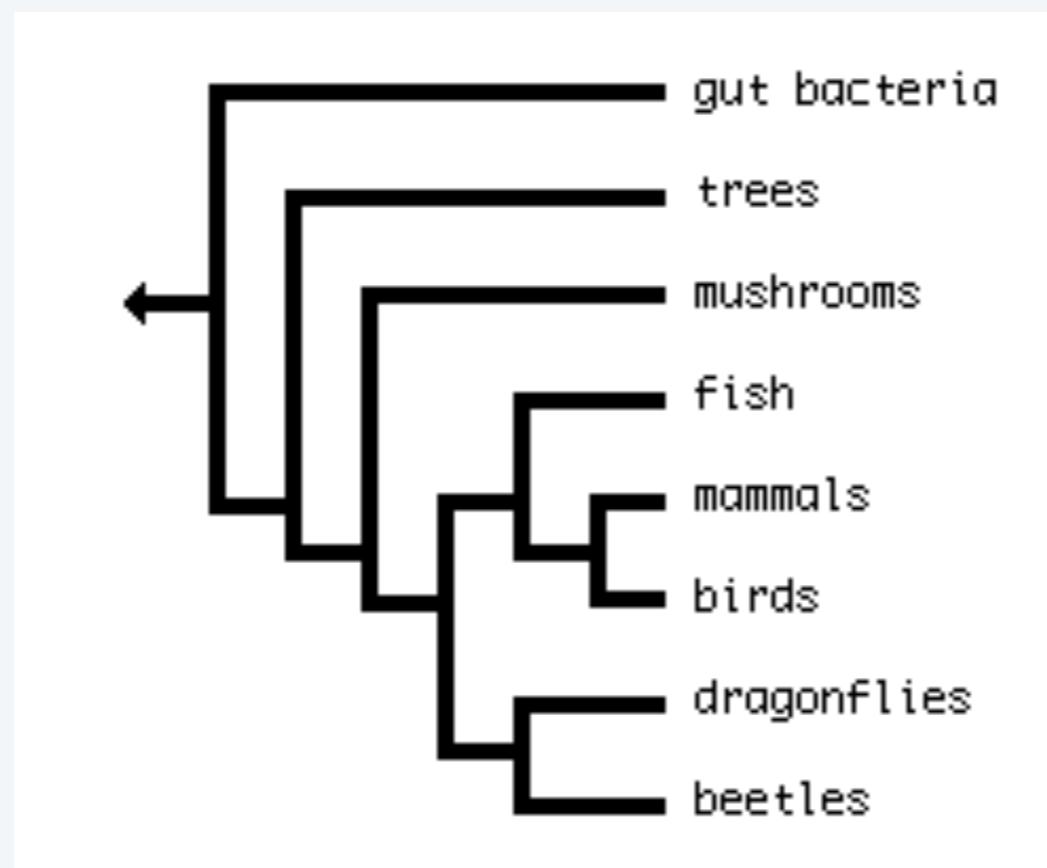
a tree



the same tree, rooted at 1

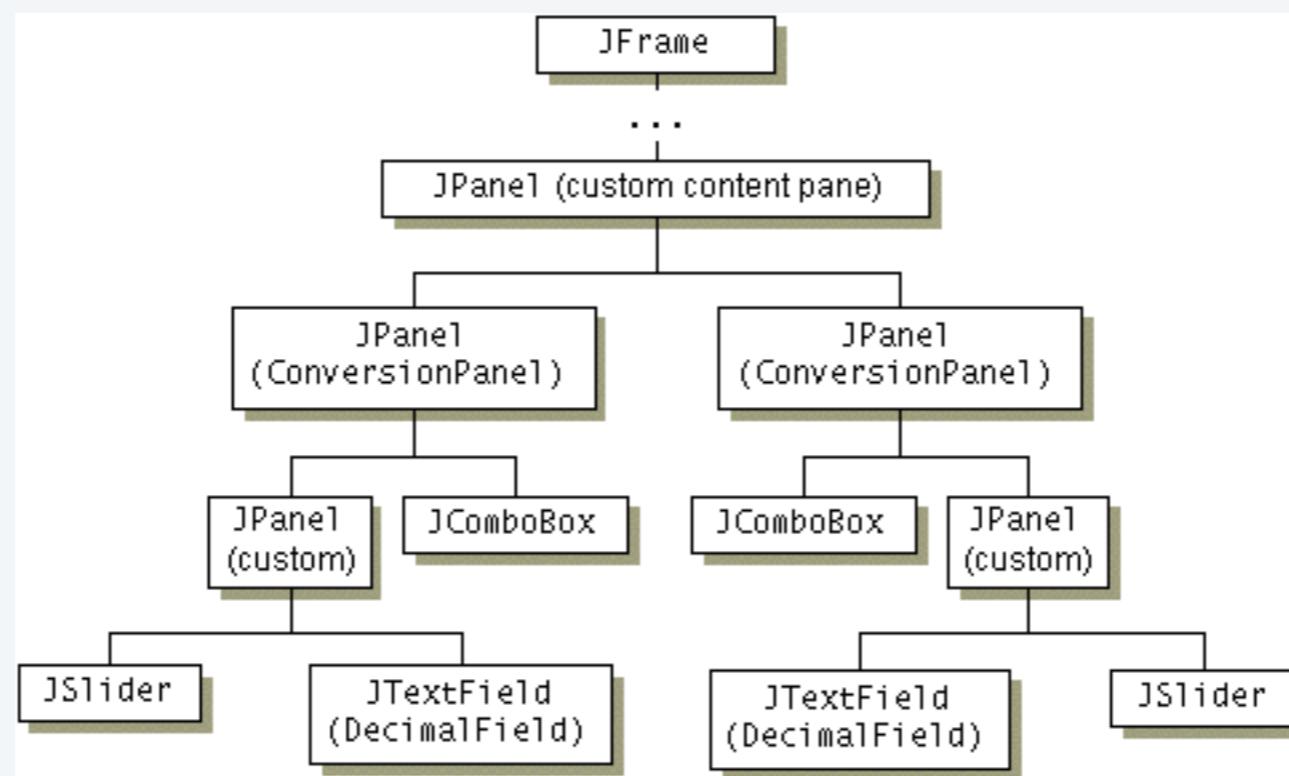
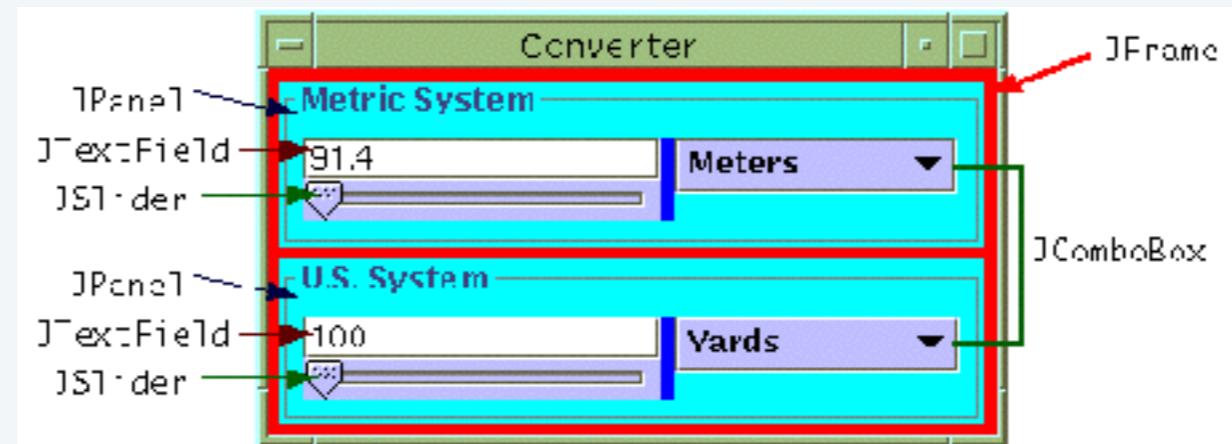
Phylogeny trees

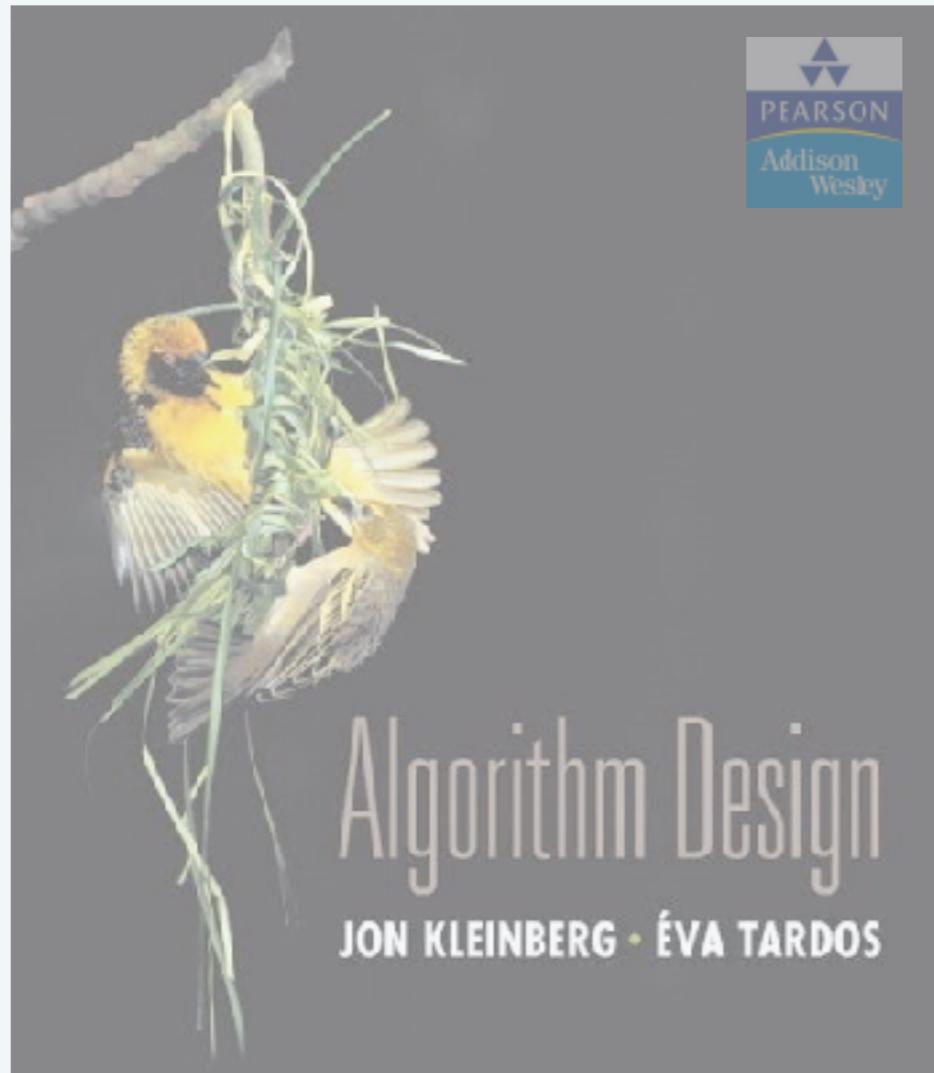
Describe evolutionary history of species.



GUI containment hierarchy

Describe organization of GUI widgets.





3. GRAPHS

- ▶ *basic definitions and applications*
- ▶ *graph connectivity and graph traversal*
- ▶ *testing bipartiteness*
- ▶ *connectivity in directed graphs*
- ▶ *DAGs and topological ordering*

Connectivity

s-t connectivity problem. Given two nodes s and t , is there a path between s and t ?

s-t shortest path problem. Given two nodes s and t , what is the length of a shortest path between s and t ?

Applications.

- Friendster.
- Maze traversal.
- Kevin Bacon number.
- Fewest hops in a communication network.

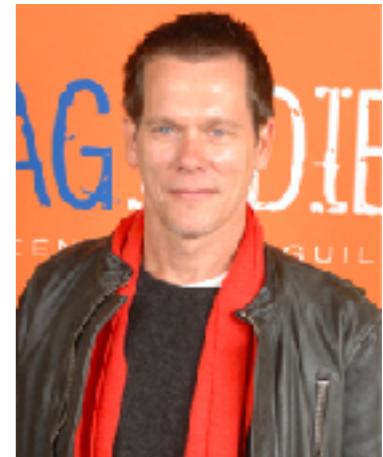
The Small World Experiment

- Milgram 1967
- Picked 300 people at random from Nebraska
- asked them to get a letter to a stockbroker in Boston
they could bypass the letter through friends they know on a first-name basis
- How many steps does it take?
 - Six degrees of separation (play by John Guare)

The Small World Experiment

- 64 chains completed
 - 6.2 average chain length (thus, six degrees of separation)
- Further observations
 - people who owned stocks had shorter paths to the stockbroker than random people
 - people from the Boston area have even closer ties

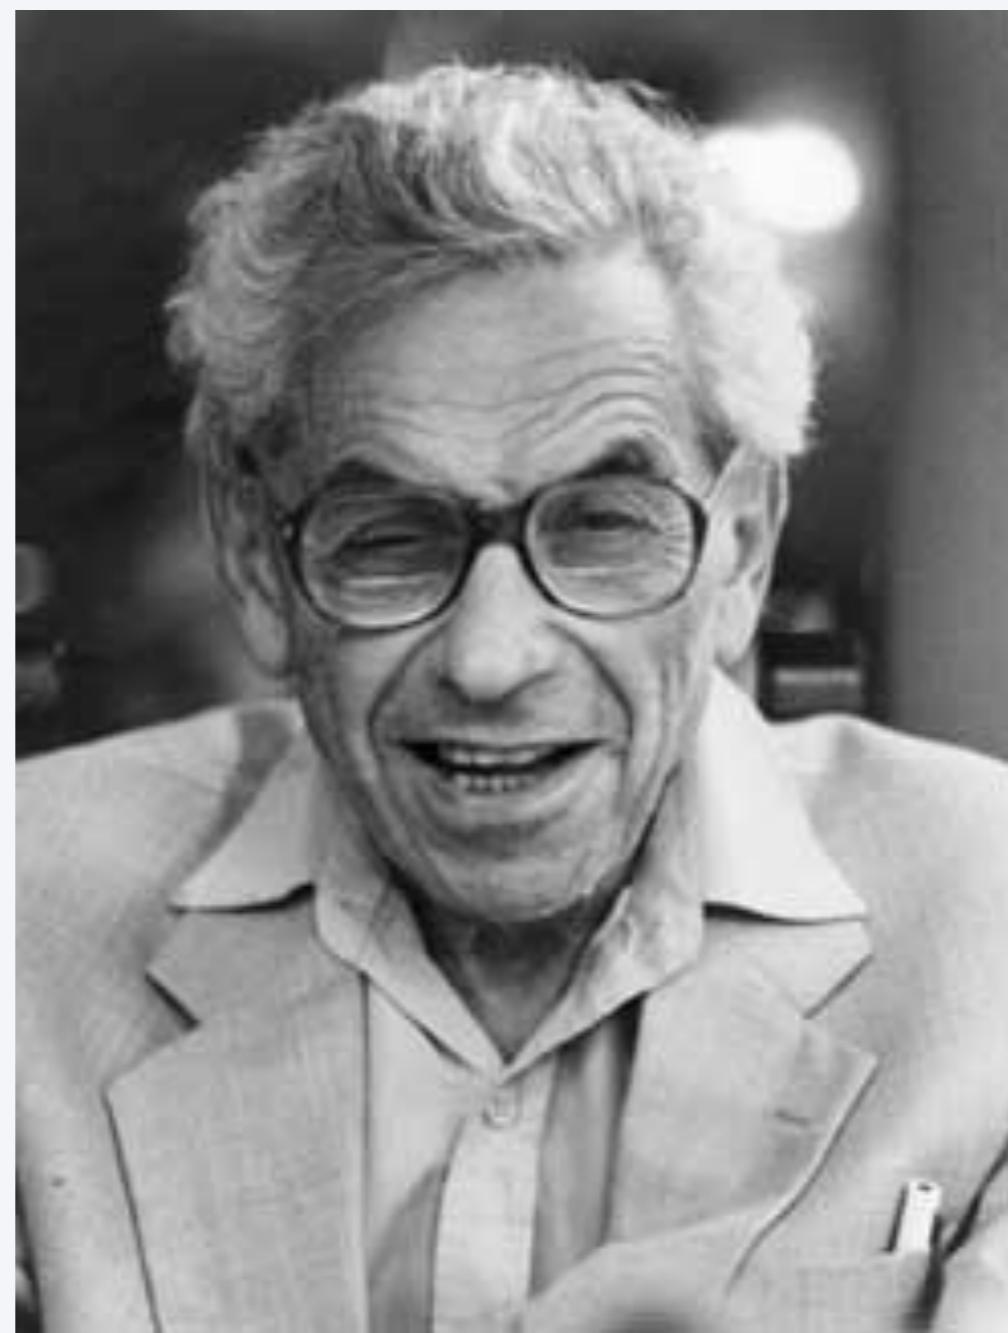
Six Degrees of Kevin Bacon



- Bacon number:
 - Create a network of Hollywood actors
 - Connect two actors if they play together in a movie
 - Bacon number: number of steps to Kevin Bacon
 - As of Dec 2007 the highest (finite) Bacon number is 8
 - Only approx 12% of all actors cannot be linked to him
 - What is the Bacon number of Elvis Presley?
-
- Elvis Presley was in Change of Habit (1969) with Edward Asner
 - Edward Asner was in JFK (1991) with Kevin Bacon

Therefore, Asner has a Bacon number of 1, and Presley (who never appeared in a film with Bacon) has a Bacon number of 2.

Erdős number

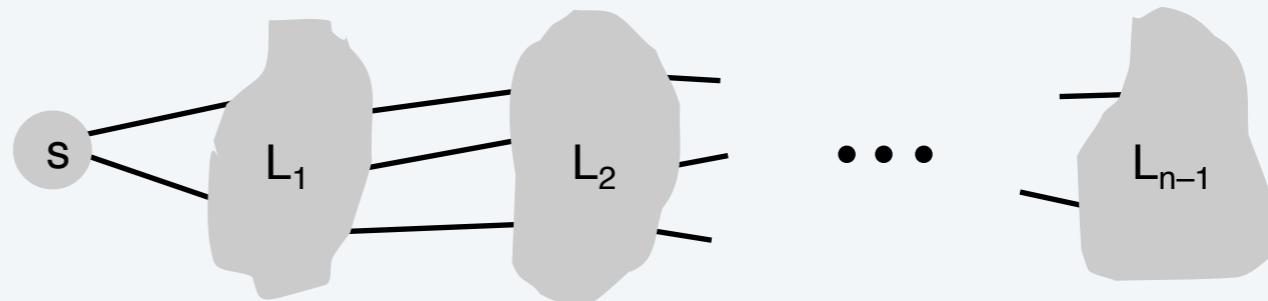


Breadth-first search

BFS intuition. Explore outward from s in all possible directions, adding nodes one “layer” at a time.

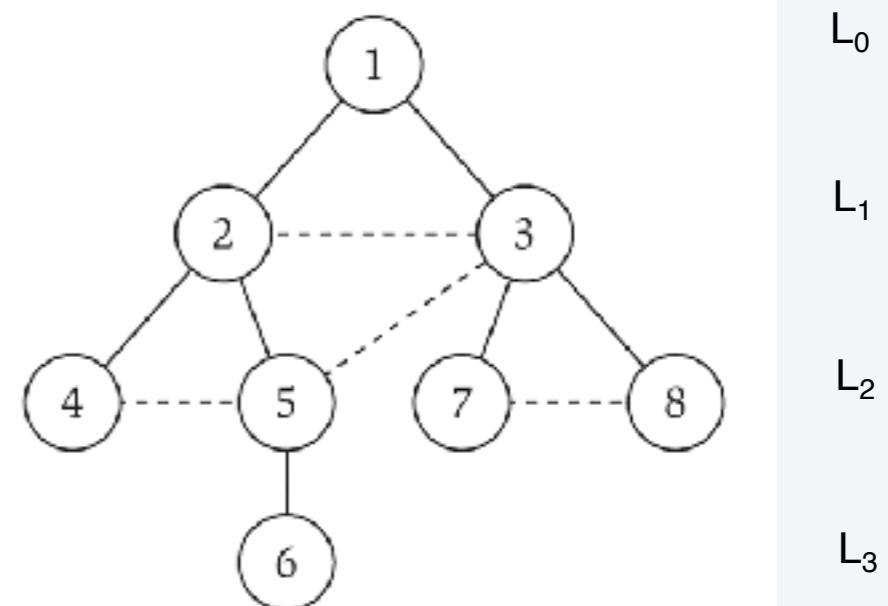
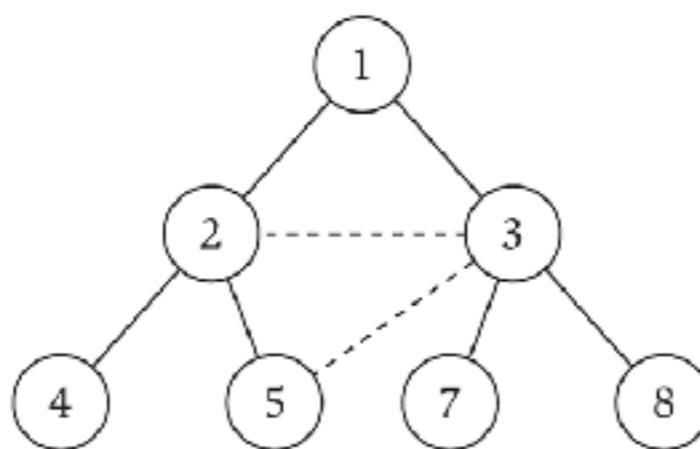
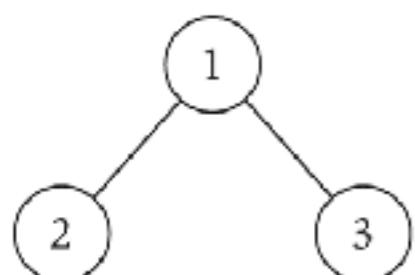
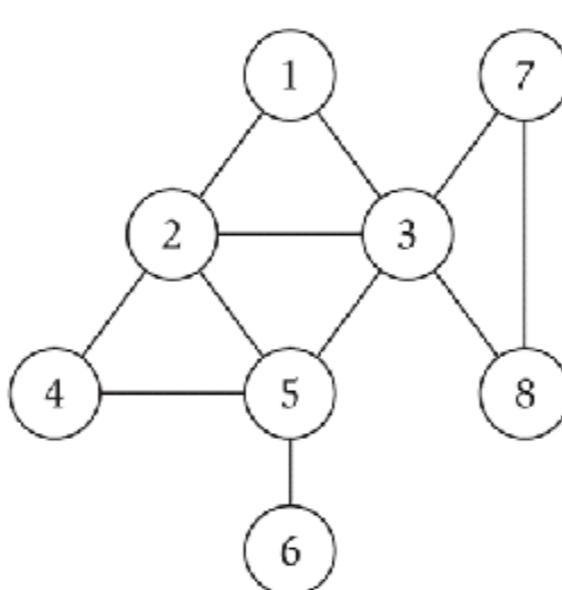
BFS algorithm.

- $L_0 = \{ s \}$.
- $L_1 = \text{all neighbors of } L_0$.
- $L_2 = \text{all nodes that do not belong to } L_0 \text{ or } L_1, \text{ and that have an edge to a node in } L_1$.
- $L_{i+1} = \text{all nodes that do not belong to an earlier layer, and that have an edge to a node in } L_i$.



Breadth-first search

Property. Let T be a BFS tree of $G = (V, E)$, and let (x, y) be an edge of G . Then, the levels of x and y differ by at most 1.



(a)

(b)

(c)

L_0

L_1

L_2

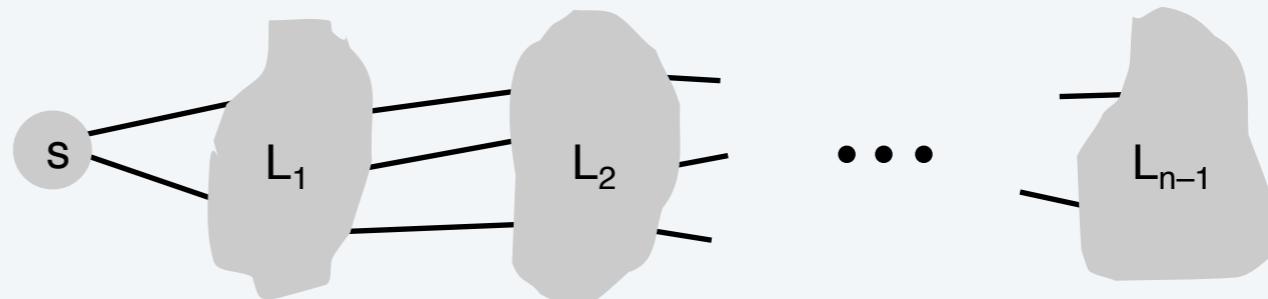
L_3

Breadth-first search

BFS intuition. Explore outward from s in all possible directions, adding nodes one “layer” at a time.

BFS algorithm.

- $L_0 = \{ s \}$.
- $L_1 = \text{all neighbors of } L_0$.
- $L_2 = \text{all nodes that do not belong to } L_0 \text{ or } L_1, \text{ and that have an edge to a node in } L_1$.
- $L_{i+1} = \text{all nodes that do not belong to an earlier layer, and that have an edge to a node in } L_i$.



Theorem. For each i , L_i consists of all nodes at distance exactly i from s . There is a path from s to t iff t appears in some layer.

Breadth-first search: analysis

Theorem. The above implementation of BFS runs in $O(m + n)$ time if the graph is given by its adjacency representation.

Pf.

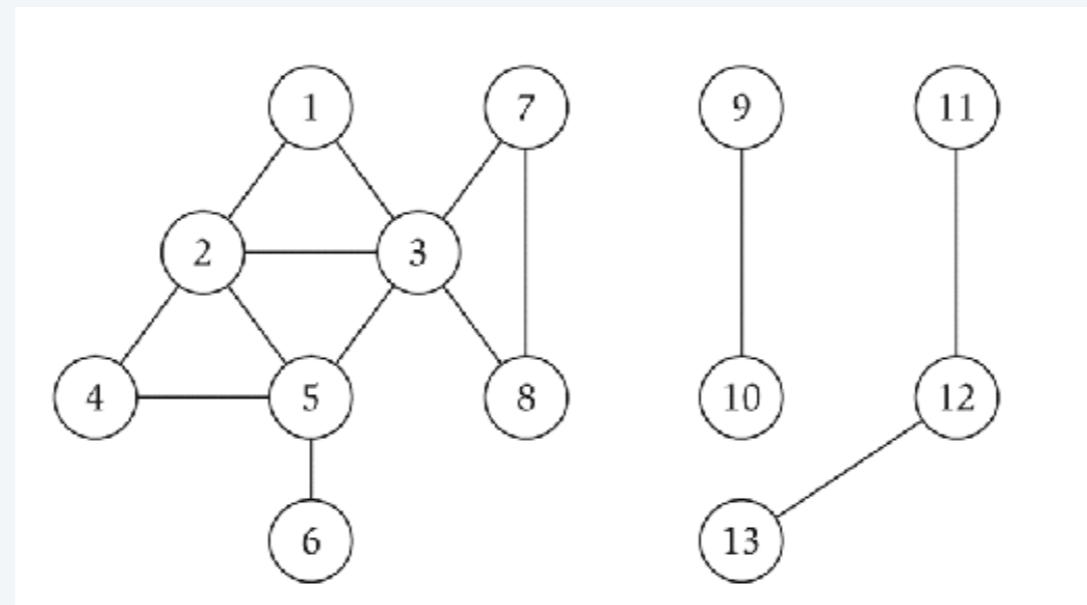
- Easy to prove $O(n^2)$ running time:
 - at most n lists $L[i]$
 - each node occurs on at most one list; for loop runs $\leq n$ times
 - when we consider node u , there are $\leq n$ incident edges (u, v) , and we spend $O(1)$ processing each edge
- Actually runs in $O(m + n)$ time:
 - when we consider node u , there are $\text{degree}(u)$ incident edges (u, v)
 - total time processing edges is $\sum_{u \in V} \text{degree}(u) = 2m$. ■



each edge (u, v) is counted exactly twice
in sum: once in $\text{degree}(u)$ and once in $\text{degree}(v)$

Connected component

Connected component. Find all nodes reachable from s .

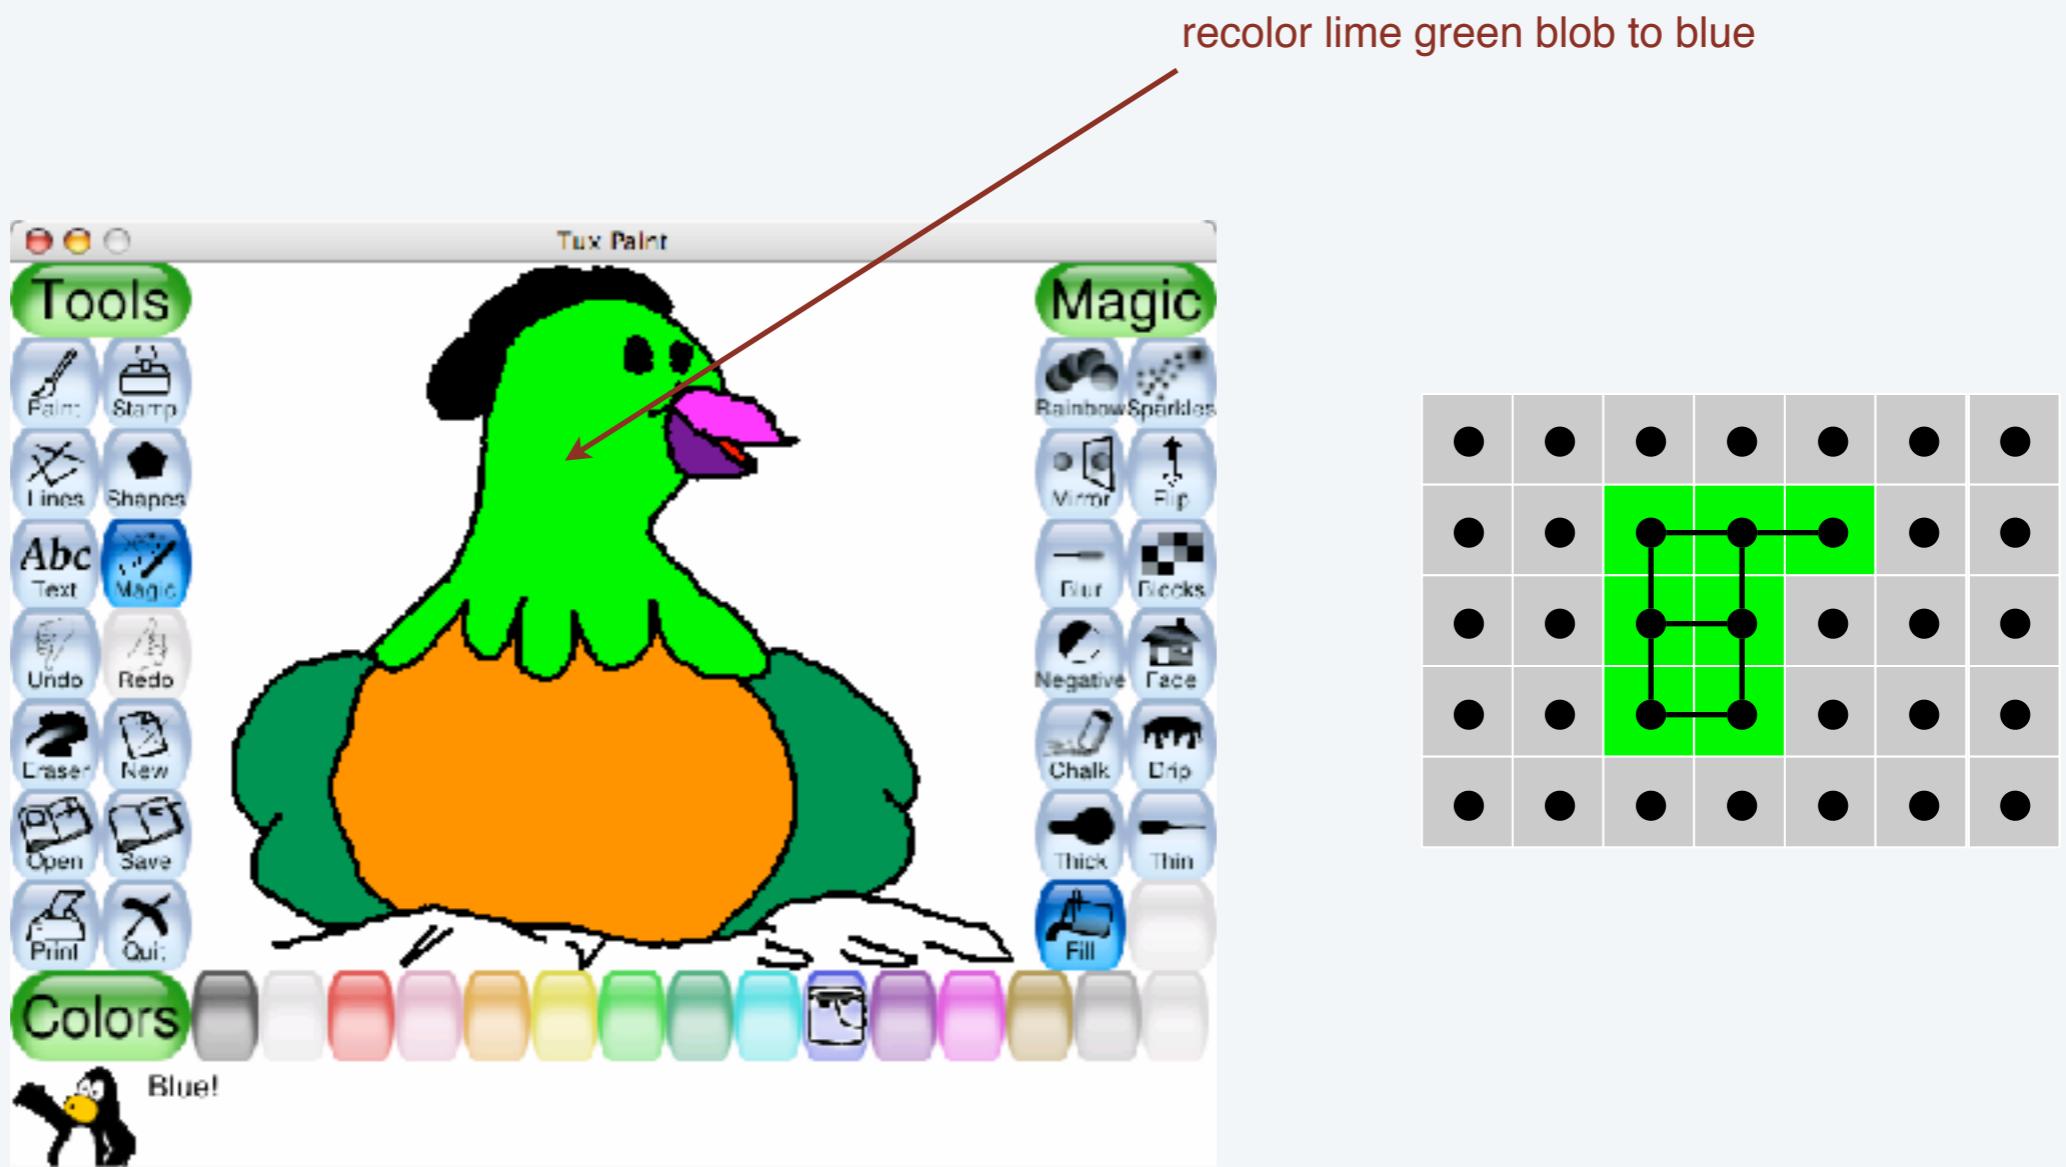


Connected component containing node 1 = { 1, 2, 3, 4, 5, 6, 7, 8 }.

Flood fill

Flood fill. Given lime green pixel in an image, change color of entire blob of neighboring lime pixels to blue.

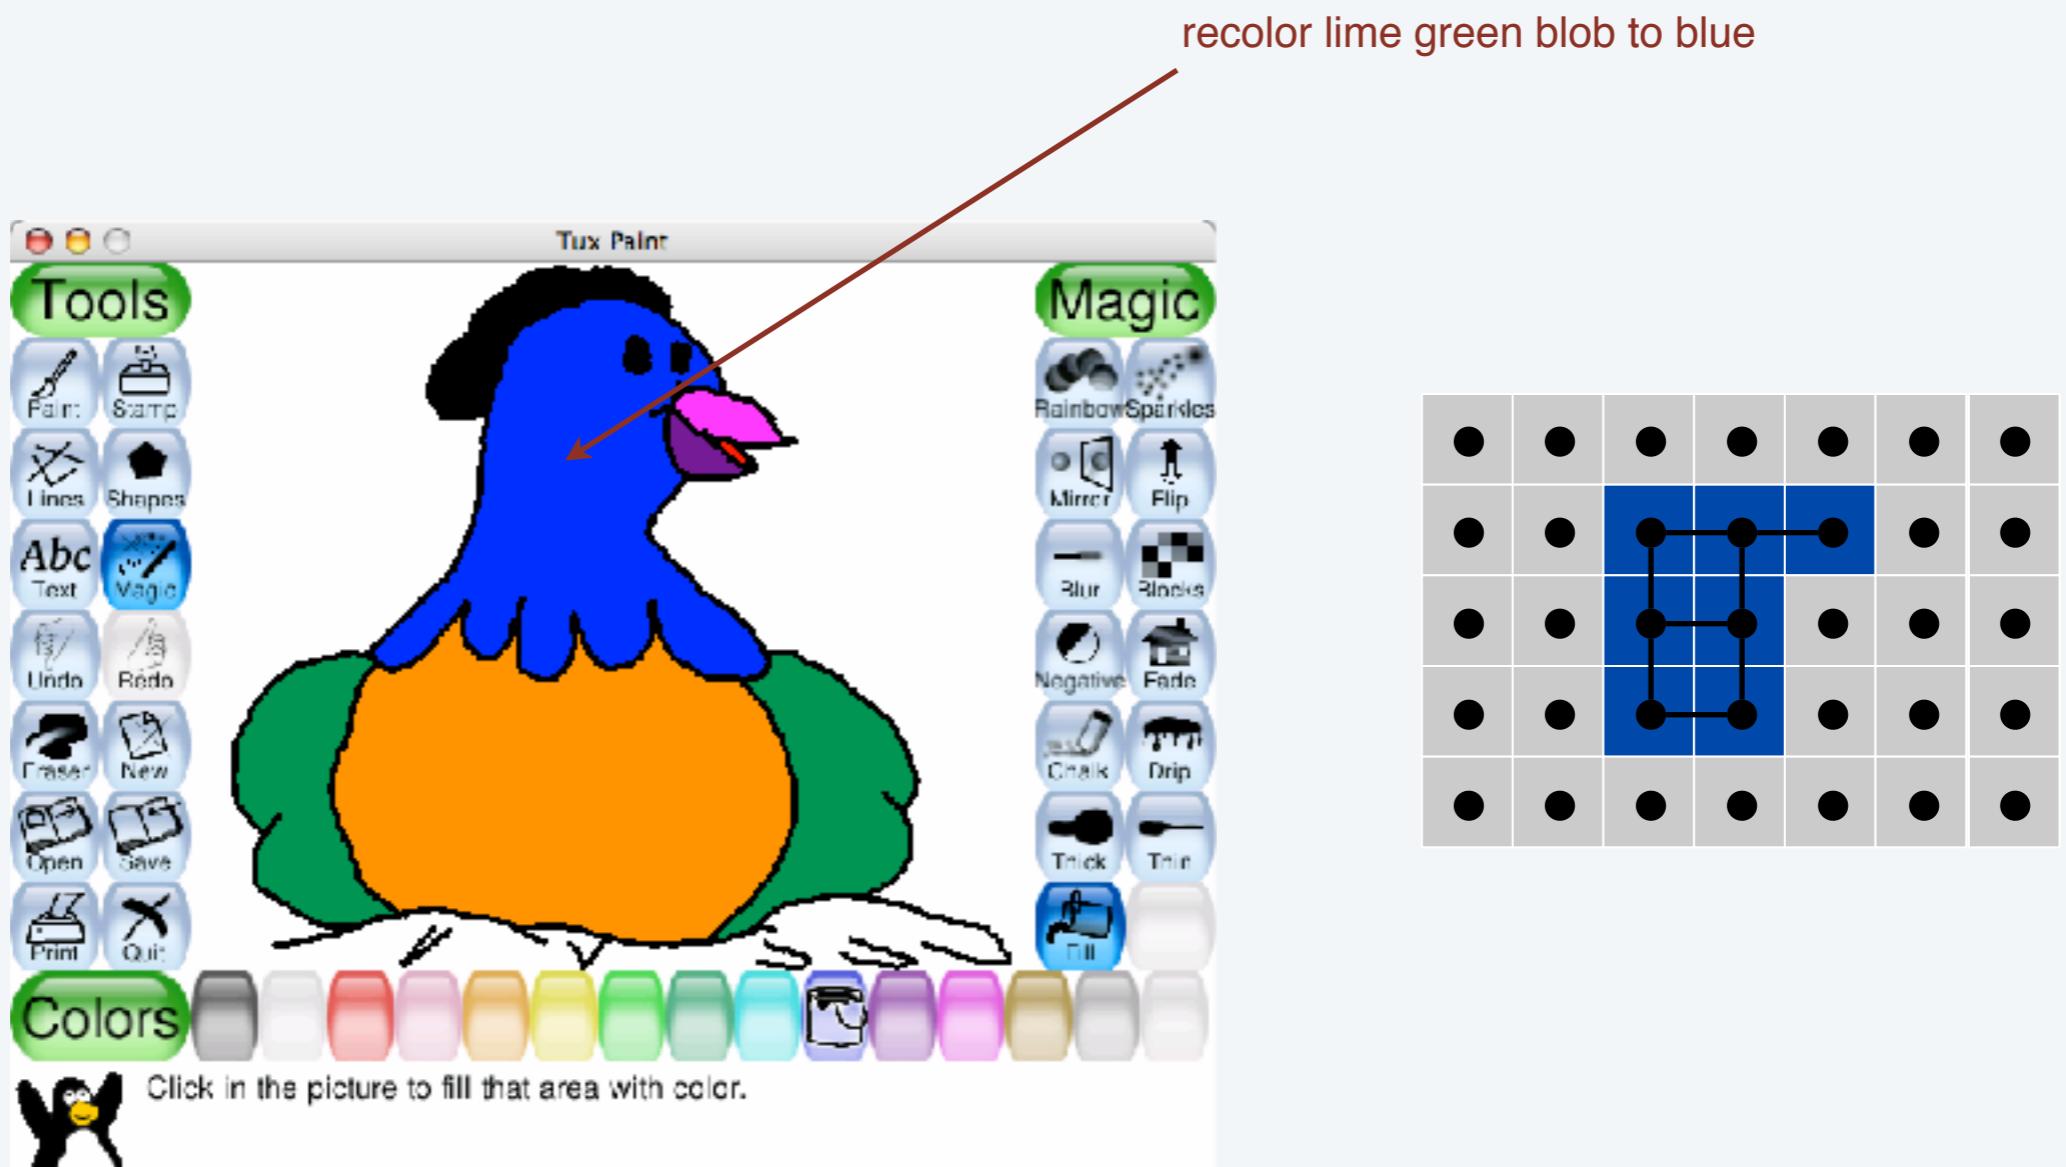
- Node: pixel.
- Edge: two neighboring lime pixels.
- Blob: connected component of lime pixels.



Flood fill

Flood fill. Given lime green pixel in an image, change color of entire blob of neighboring lime pixels to blue.

- Node: pixel.
- Edge: two neighboring lime pixels.
- Blob: connected component of lime pixels.



Connected component

Connected component. Find all nodes reachable from s .

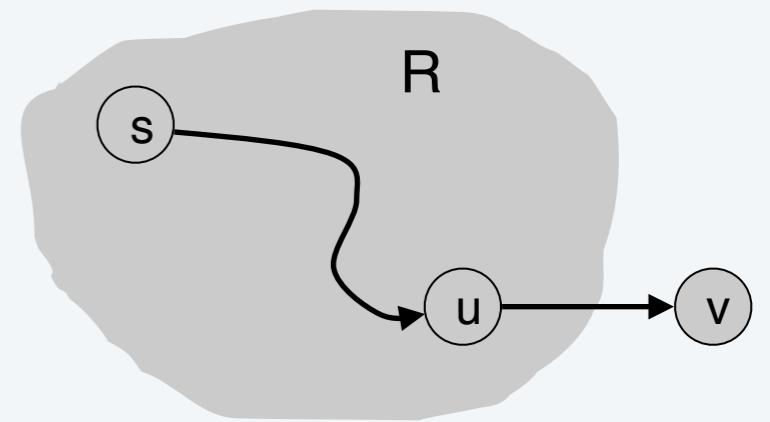
R will consist of nodes to which s has a path

Initially $R = \{s\}$

While there is an edge (u, v) where $u \in R$ and $v \notin R$

Add v to R

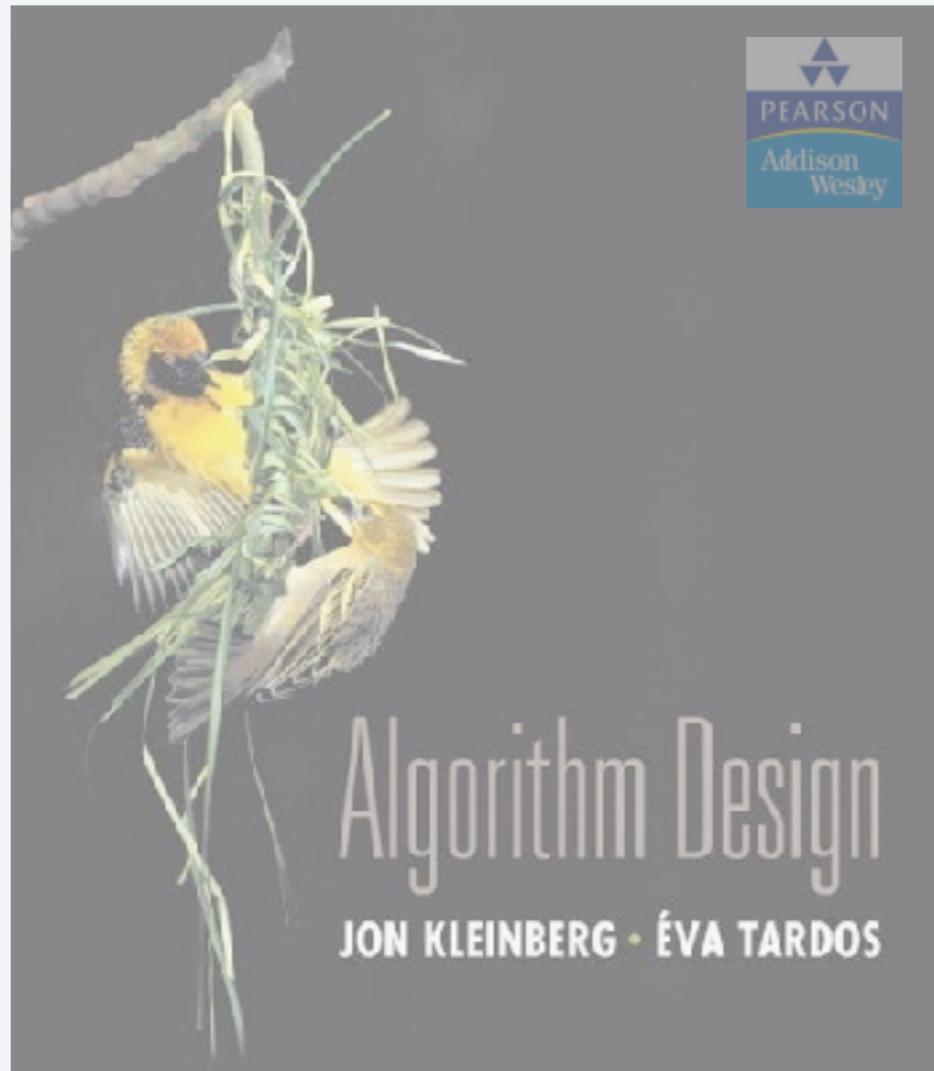
Endwhile



it's safe to add v

Theorem. Upon termination, R is the connected component containing s .

- BFS = explore in order of distance from s .
- DFS = explore in a different way.



3. GRAPHS

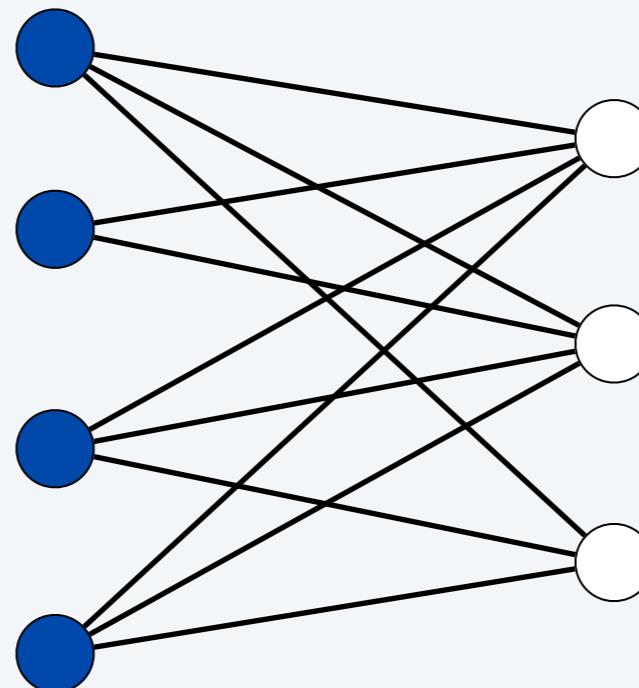
- ▶ *basic definitions and applications*
- ▶ *graph connectivity and graph traversal*
- ▶ ***testing bipartiteness***
- ▶ *connectivity in directed graphs*
- ▶ *DAGs and topological ordering*

Bipartite graphs

Def. An undirected graph $G = (V, E)$ is **bipartite** if the nodes can be colored blue or white such that every edge has one white and one blue end.

Applications.

- Stable matching: med-school residents = blue, hospitals = white.
- Scheduling: machines = blue, jobs = white.



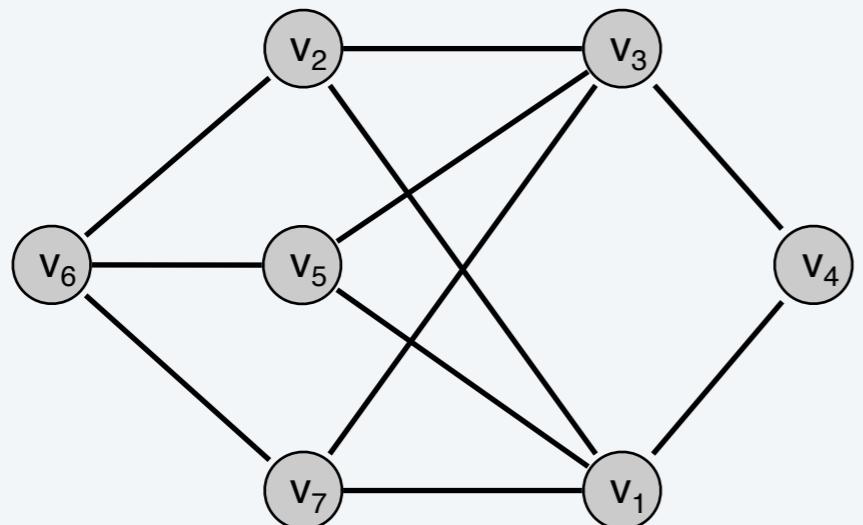
a bipartite graph

Testing bipartiteness

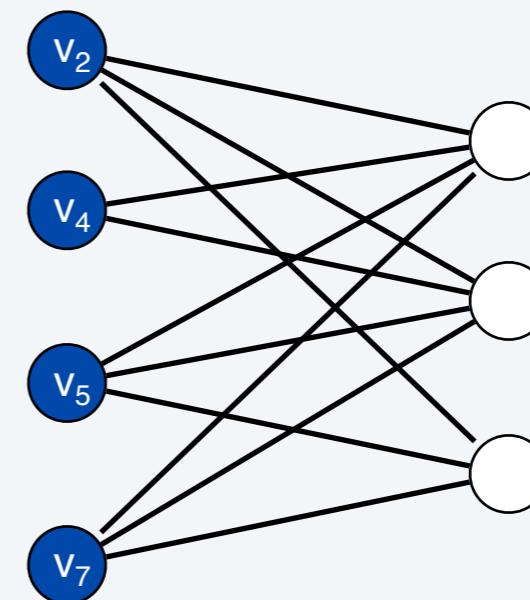
Many graph problems become:

- Easier if the underlying graph is bipartite (matching).
- Tractable if the underlying graph is bipartite (independent set).

Before attempting to design an algorithm, we need to understand structure of bipartite graphs.



a bipartite graph G

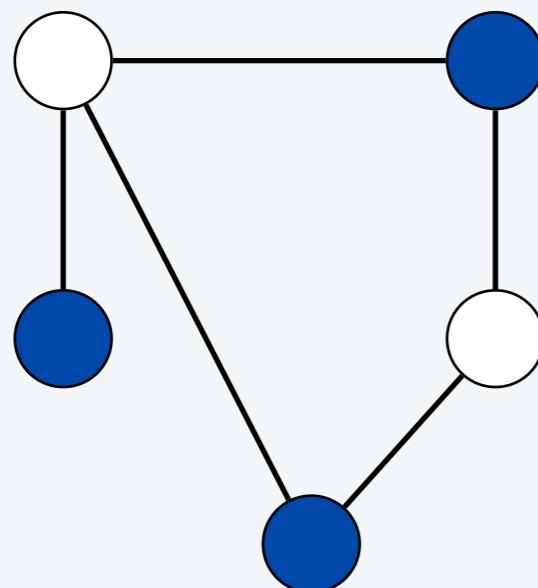


another drawing of G

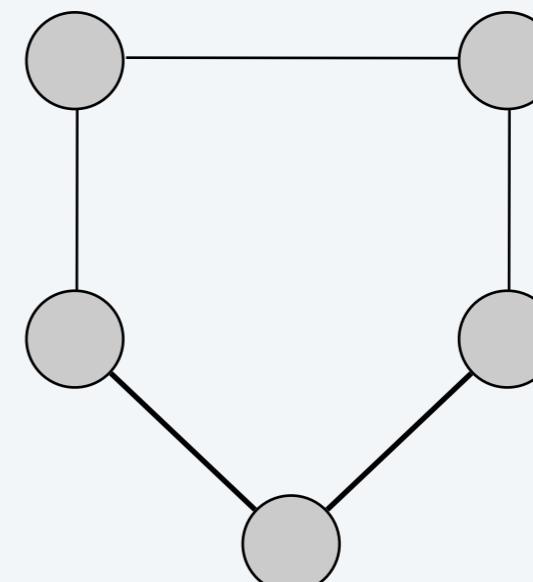
An obstruction to bipartiteness

Lemma. If a graph G is bipartite, it cannot contain an odd-length cycle.

Pf. Not possible to 2-color the odd-length cycle, let alone G .



bipartite
(2-colorable)

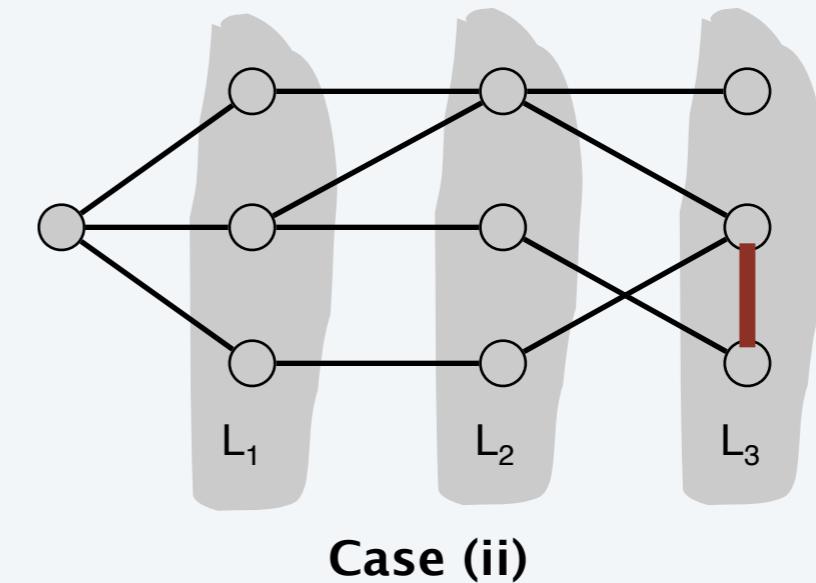
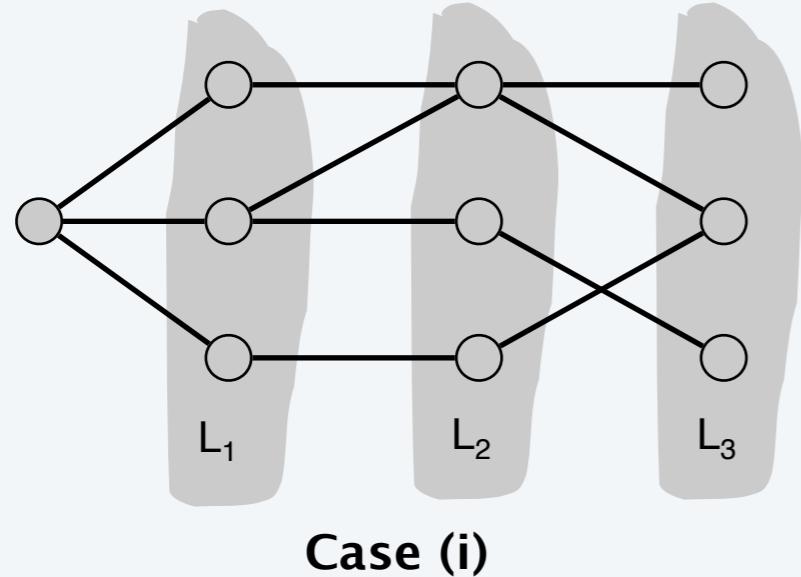


not bipartite
(not 2-colorable)

Bipartite graphs

Lemma. Let G be a connected graph, and let L_0, \dots, L_k be the layers produced by BFS starting at node s . Exactly one of the following holds.

- (i) No edge of G joins two nodes of the same layer, and G is bipartite.
- (ii) An edge of G joins two nodes of the same layer, and G contains an odd-length cycle (and hence is not bipartite).



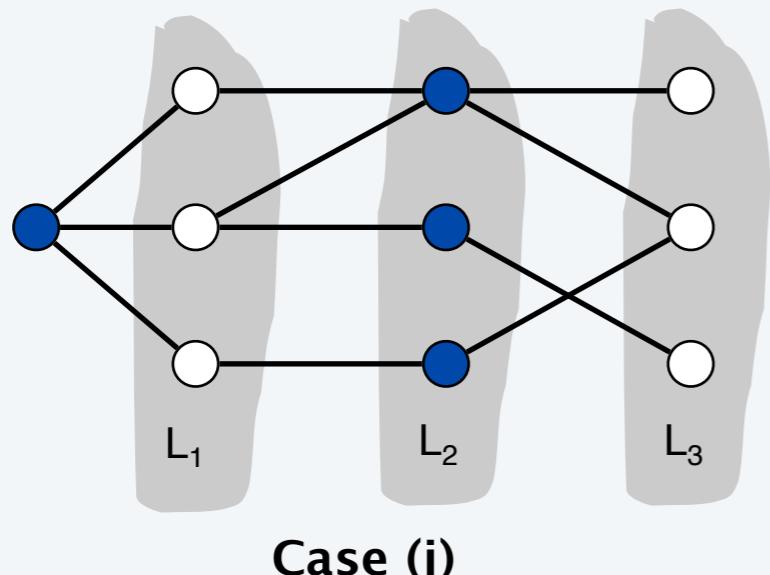
Bipartite graphs

Lemma. Let G be a connected graph, and let L_0, \dots, L_k be the layers produced by BFS starting at node s . Exactly one of the following holds.

- (i) No edge of G joins two nodes of the same layer, and G is bipartite.
- (ii) An edge of G joins two nodes of the same layer, and G contains an odd-length cycle (and hence is not bipartite).

Pf. (i)

- Suppose no edge joins two nodes in same layer.
- By BFS property, each edge joins two nodes in adjacent levels.
- Bipartition: white = nodes on odd levels, blue = nodes on even levels.



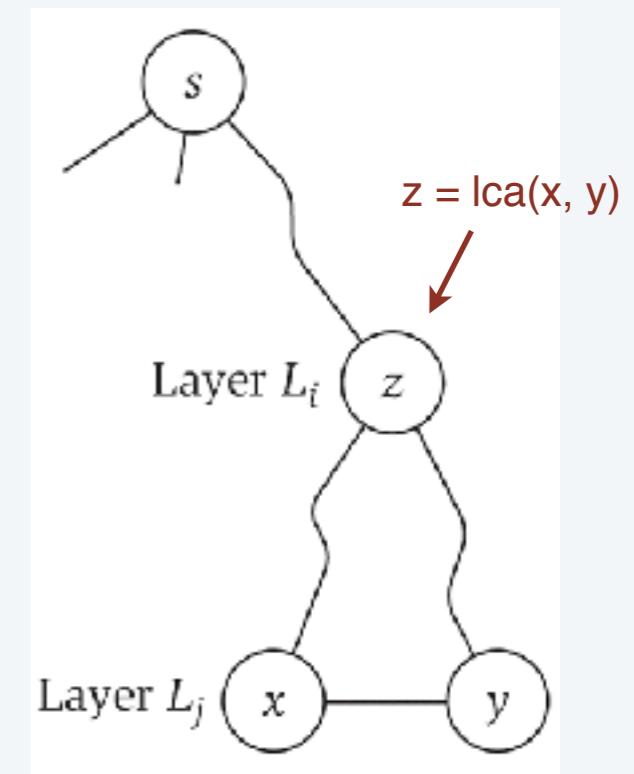
Bipartite graphs

Lemma. Let G be a connected graph, and let L_0, \dots, L_k be the layers produced by BFS starting at node s . Exactly one of the following holds.

- (i) No edge of G joins two nodes of the same layer, and G is bipartite.
- (ii) An edge of G joins two nodes of the same layer, and G contains an odd-length cycle (and hence is not bipartite).

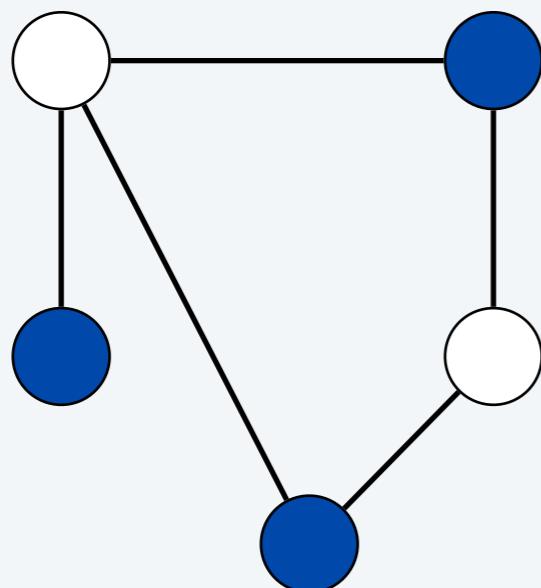
Pf. (ii)

- Suppose (x, y) is an edge with x, y in same level L_j .
- Let $z = lca(x, y) =$ lowest common ancestor.
- Let L_i be level containing z .
- Consider cycle that takes edge from x to y ,
then path from y to z , then path from z to x .
- Its length is $\underbrace{1}_{(x, y)} + \underbrace{(j - i)}_{\text{path from } y \text{ to } z} + \underbrace{(j - i)}_{\text{path from } z \text{ to } x}$, which is odd. ▀

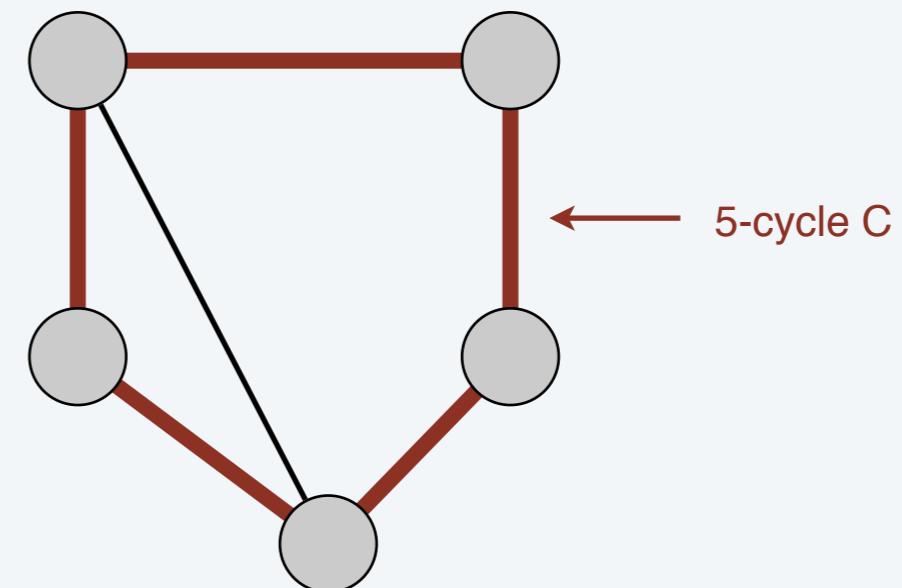


The only obstruction to bipartiteness

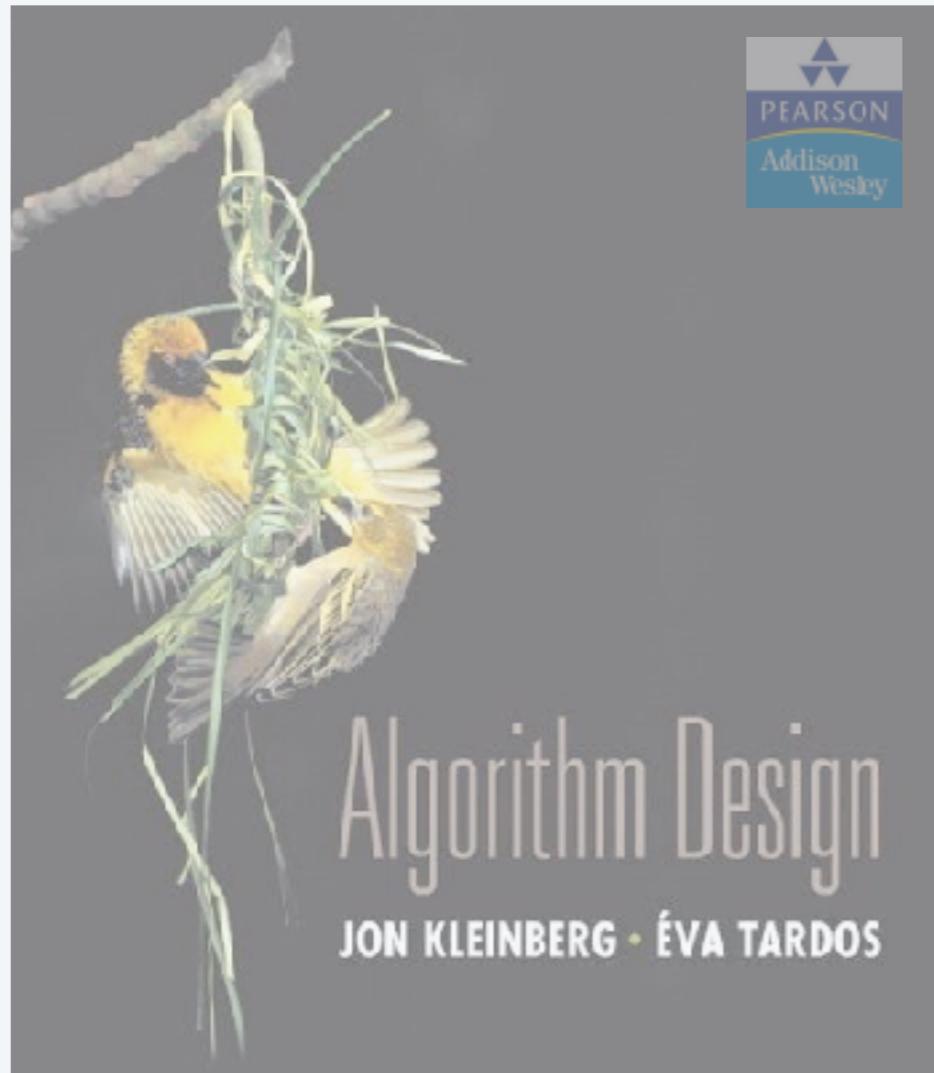
Corollary. A graph G is bipartite iff it contains no odd-length cycle.



bipartite
(2-colorable)



not bipartite
(not 2-colorable)



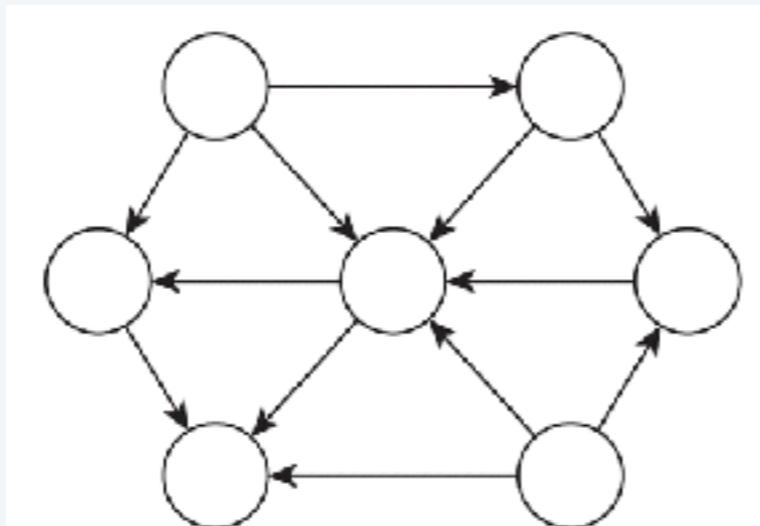
3. GRAPHS

- ▶ *basic definitions and applications*
- ▶ *graph connectivity and graph traversal*
- ▶ *testing bipartiteness*
- ▶ *connectivity in directed graphs*
- ▶ *DAGs and topological ordering*

Directed graphs

Notation. $G = (V, E)$.

- Edge (u, v) leaves node u and enters node v .



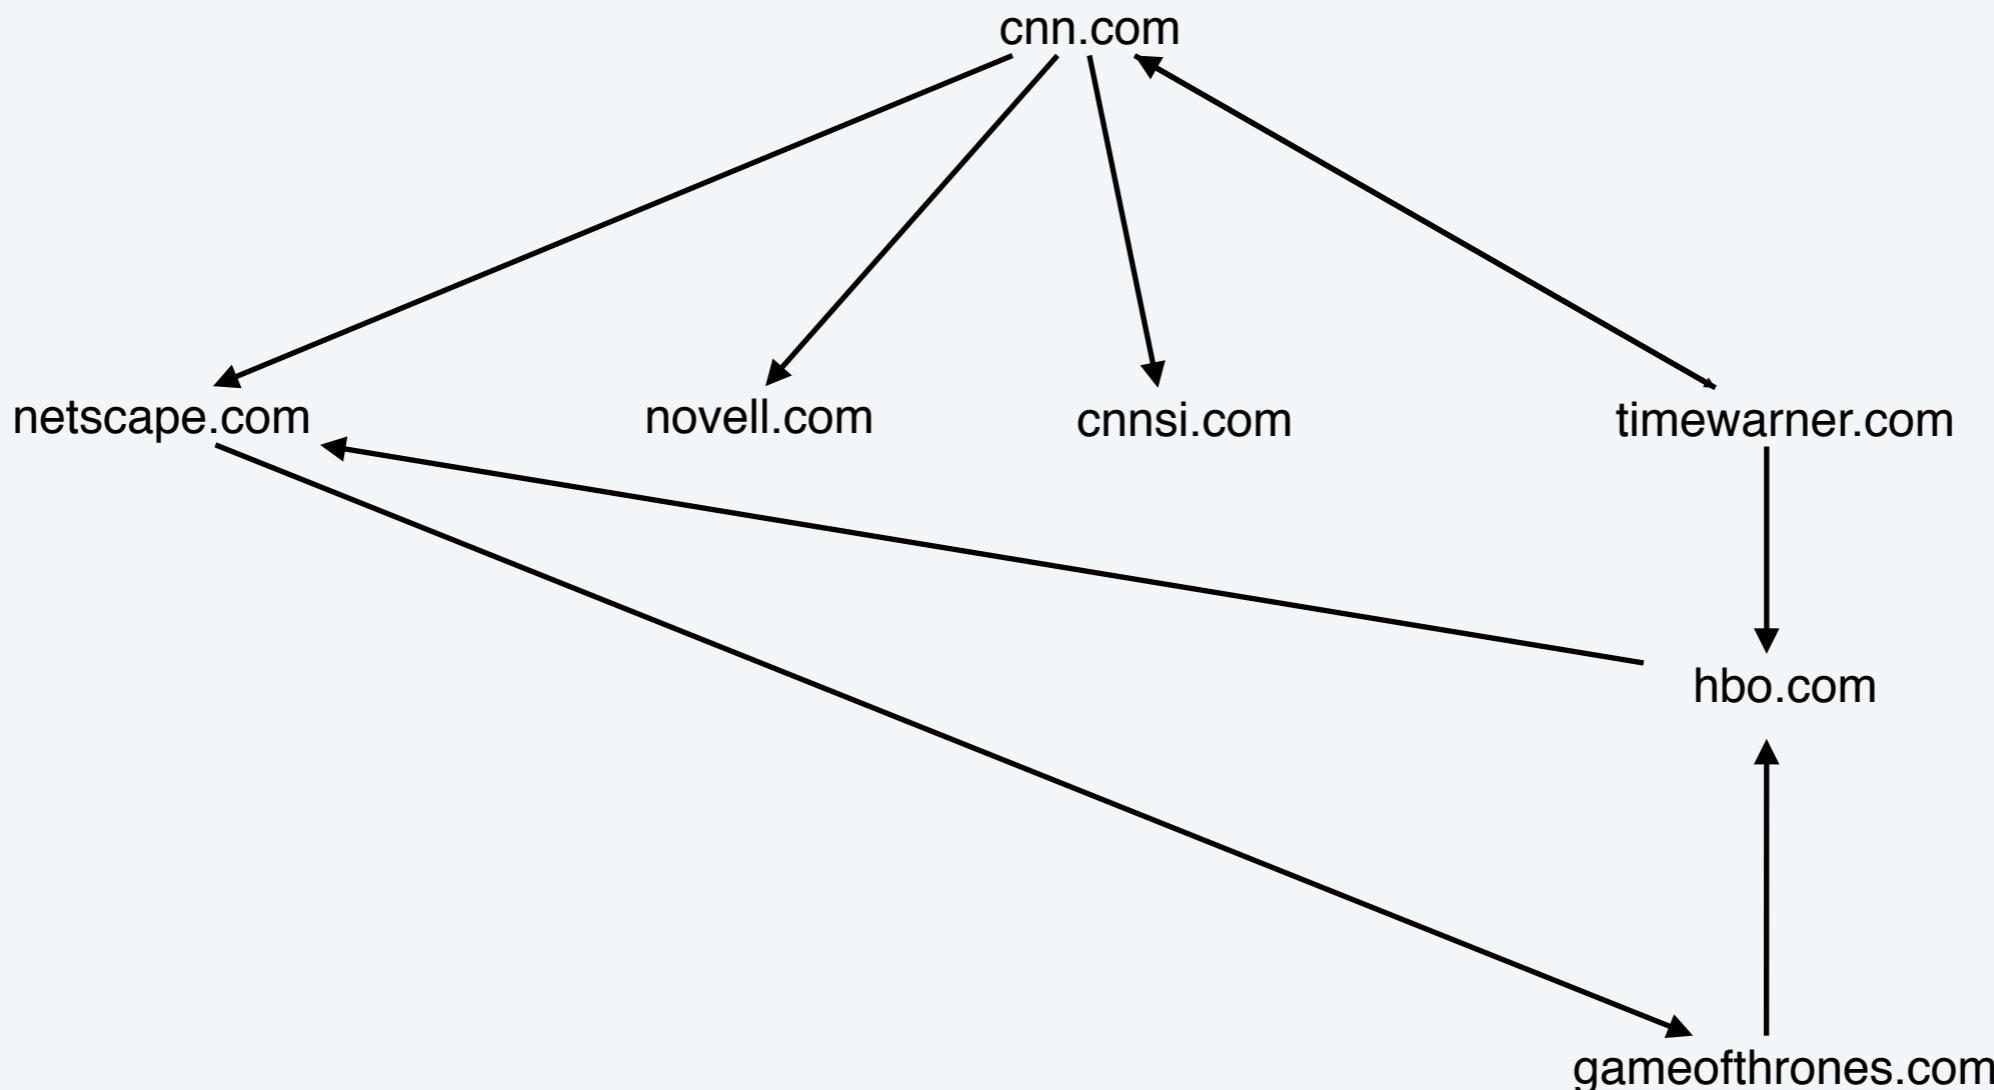
Ex. Web graph: hyperlink points from one web page to another.

- Orientation of edges is crucial.
- Modern web search engines exploit hyperlink structure to rank web pages by importance.

World wide web

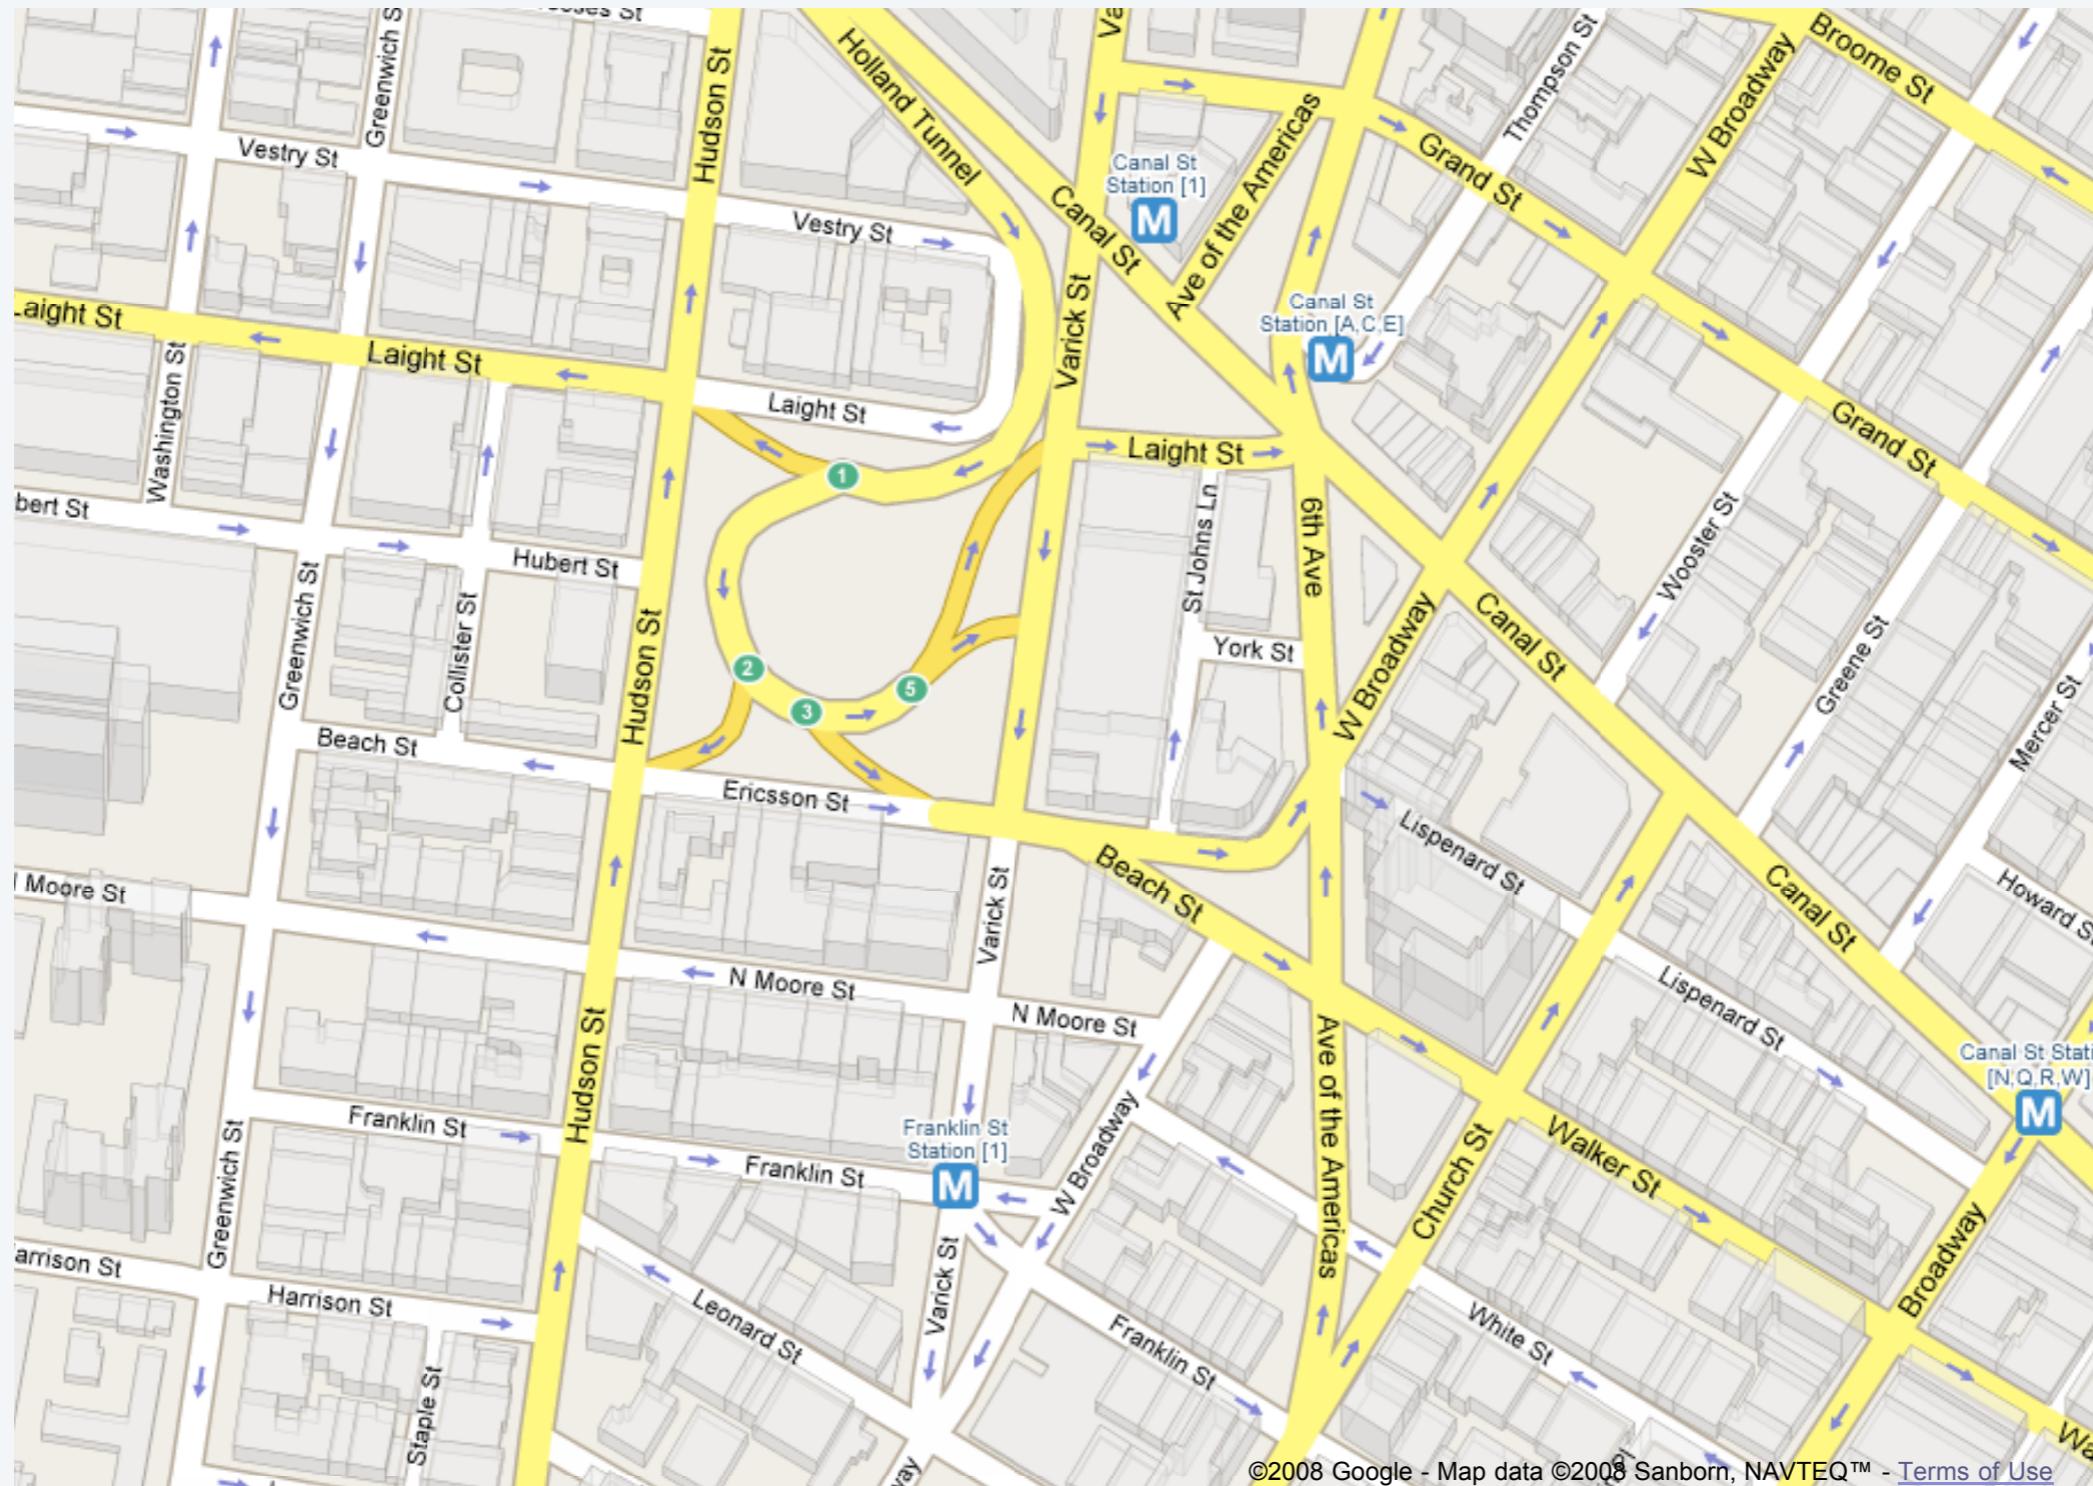
Web graph.

- Node: web page.
- Edge: hyperlink from one page to another (orientation is crucial).
- Modern search engines exploit hyperlink structure to rank web pages by importance.



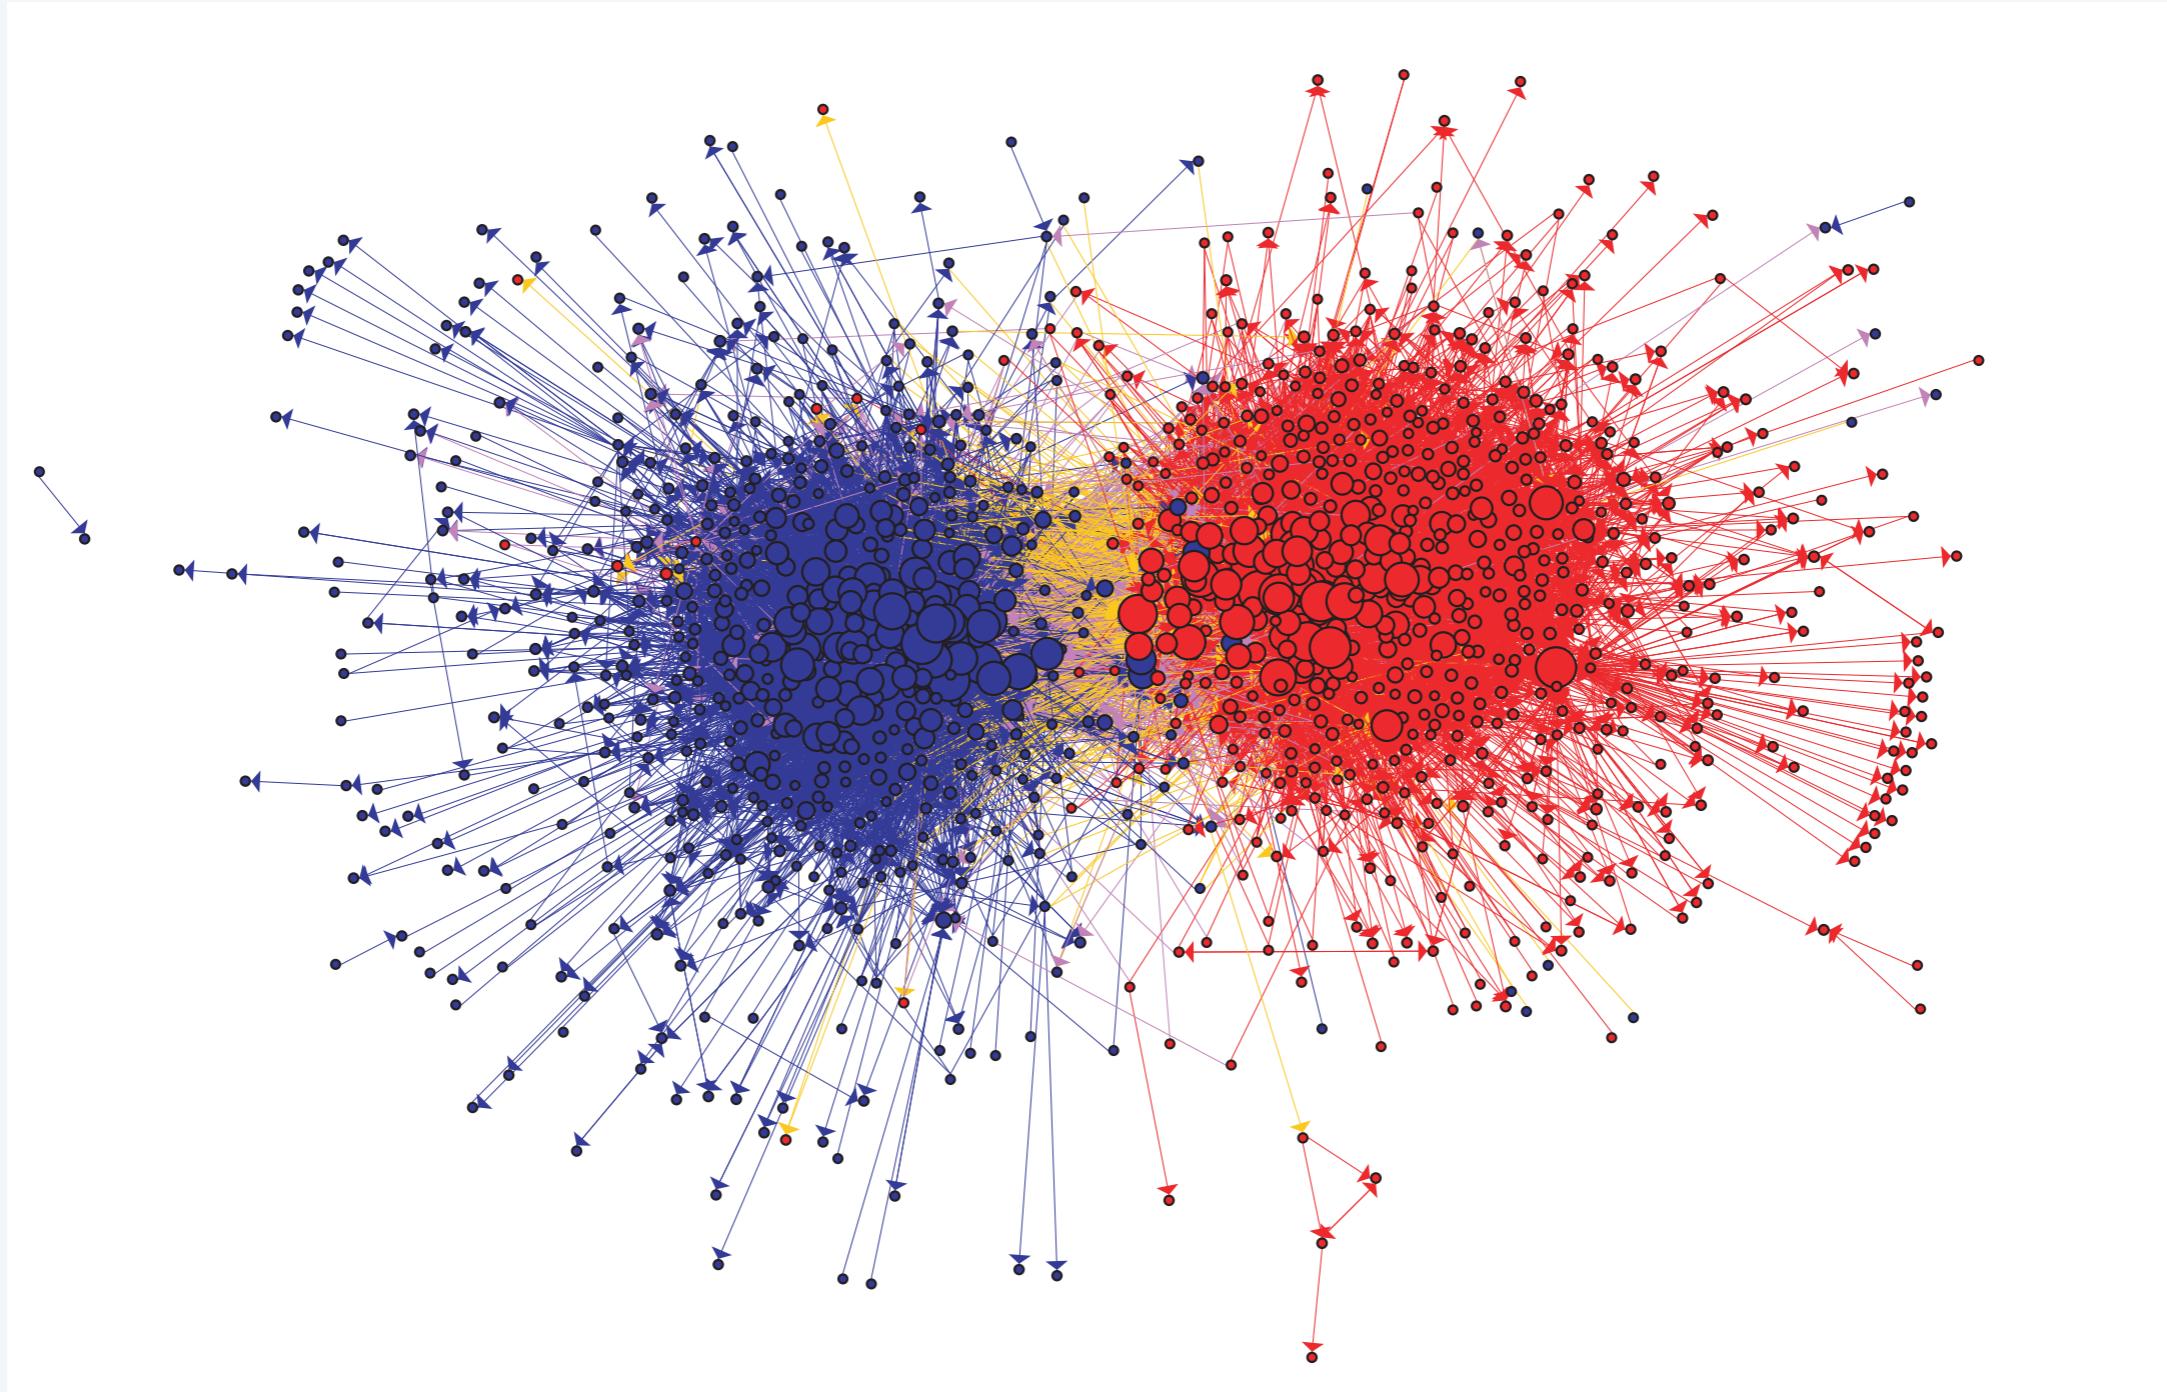
Road network

Node = intersection; edge = one-way street.



Political blogosphere graph

Node = political blog; edge = link.

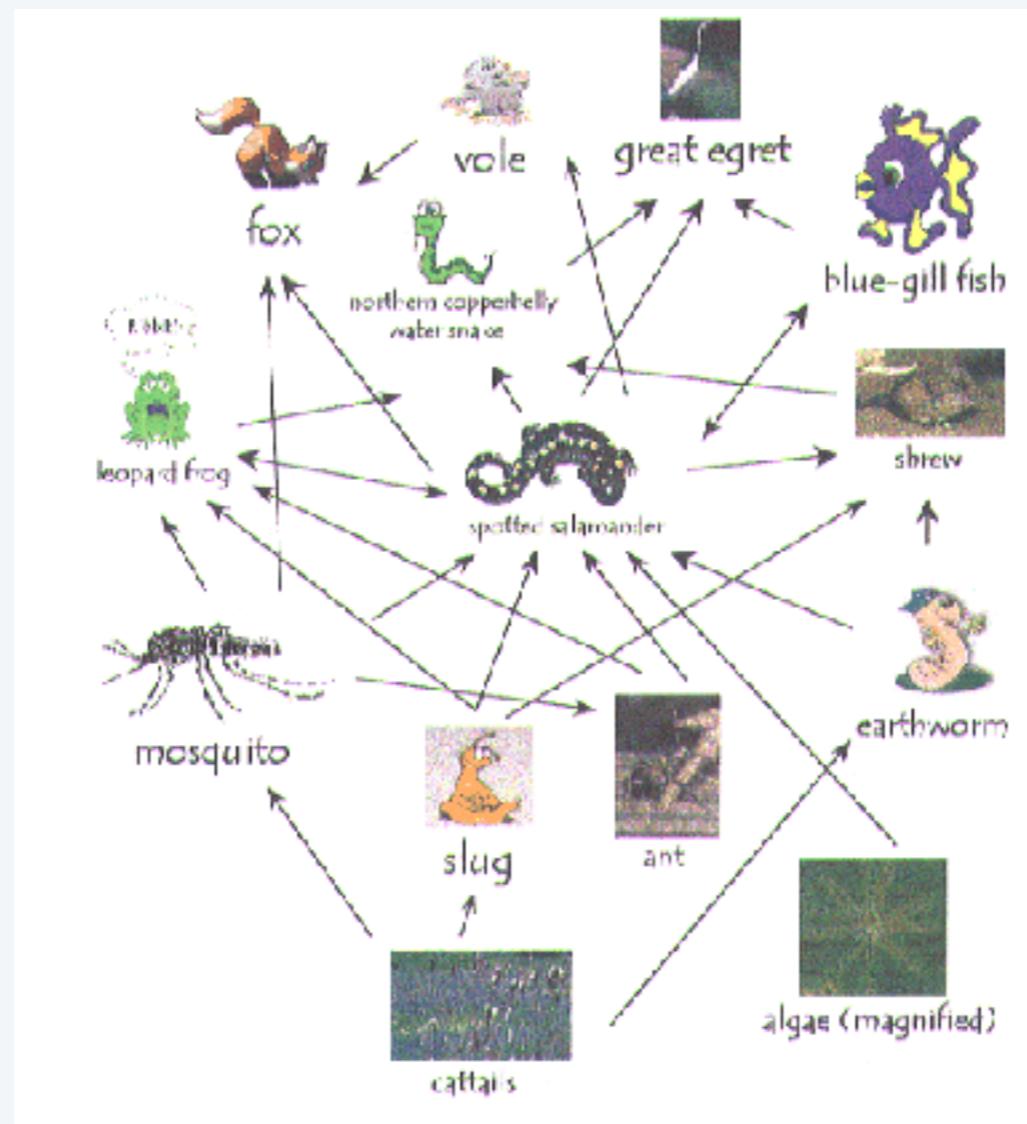


The Political Blogosphere and the 2004 U.S. Election: Divided They Blog, Adamic and Glance, 2005

Ecological food web

Food web graph.

- Node = species.
- Edge = from prey to predator.



Reference: <http://www.twinkl.com/resource/T2-1000-Wetlands-Salamander-SalGraphics/salfoodweb.gif>

Some directed graph applications

directed graph	node	directed edge
transportation	street intersection	one-way street
web	web page	hyperlink
food web	species	predator-prey relationship
WordNet	synset	hypernym
scheduling	task	precedence constraint
financial	bank	transaction
cell phone	person	placed call
infectious disease	person	infection
game	board position	legal move
citation	journal article	citation
object graph	object	pointer
inheritance hierarchy	class	inherits from
control flow	code block	jump

Graph search

Directed reachability. Given a node s , find all nodes reachable from s .

Directed s–t shortest path problem. Given two nodes s and t , what is the length of a shortest path from s to t ?

Graph search. BFS extends naturally to directed graphs.

Web crawler. Start from web page s . Find all web pages linked from s , either directly or indirectly.

Strong connectivity

Def. Nodes u and v are **mutually reachable** if there is both a path from u to v and also a path from v to u .

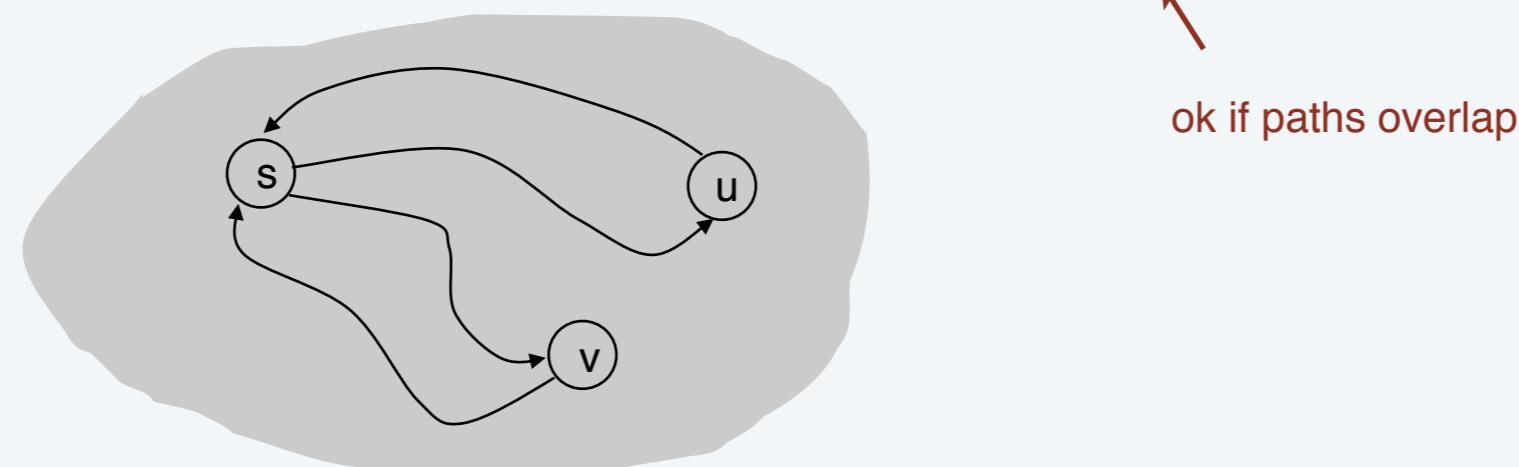
Def. A graph is **strongly connected** if every pair of nodes is mutually reachable.

Lemma. Let s be any node. G is strongly connected iff every node is reachable from s , and s is reachable from every node.

Pf. \Rightarrow Follows from definition.

Pf. \Leftarrow Path from u to v : concatenate $u \rightarrow s$ path with $s \rightarrow v$ path.

Path from v to u : concatenate $v \rightarrow s$ path with $s \rightarrow u$ path. ■

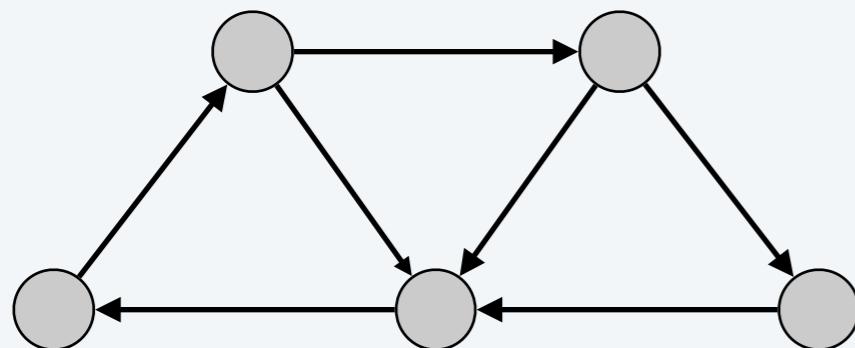


Strong connectivity: algorithm

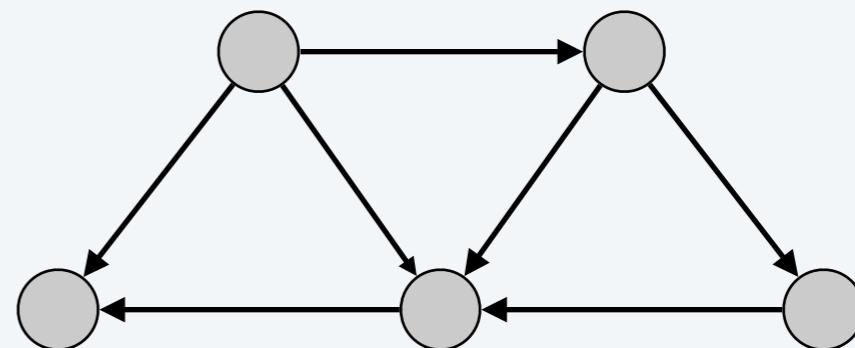
Theorem. Can determine if G is strongly connected in $O(m + n)$ time.

Pf.

- Pick any node s .
- Run BFS from s in G .
- Run BFS from s in G^{reverse} . reverse orientation of every edge in G
- Return true iff all nodes reached in both BFS executions.
- Correctness follows immediately from previous lemma. ■



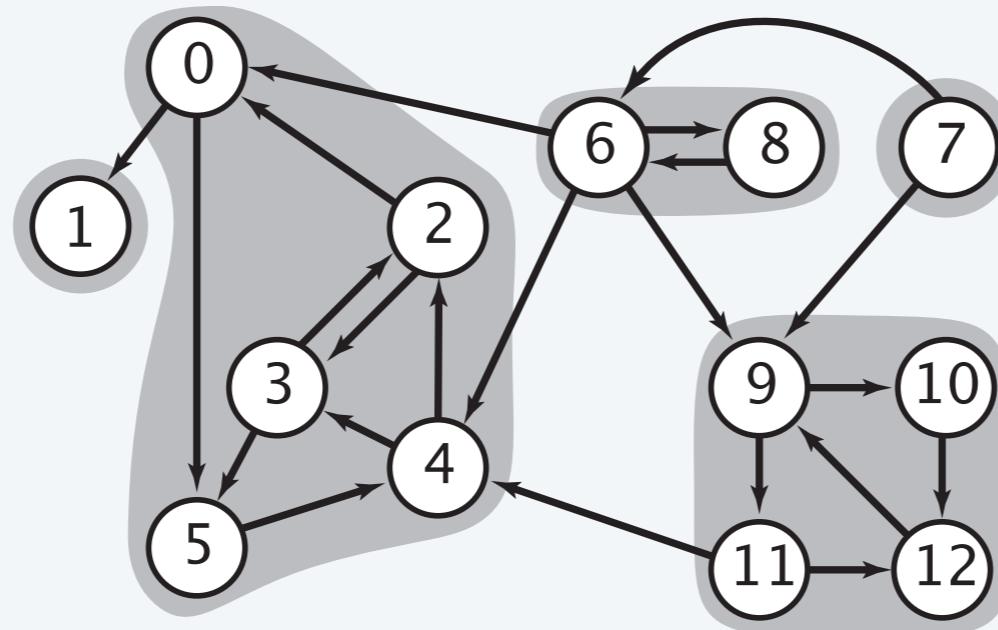
strongly connected



not strongly connected

Strong components

Def. A **strong component** is a maximal subset of mutually reachable nodes.



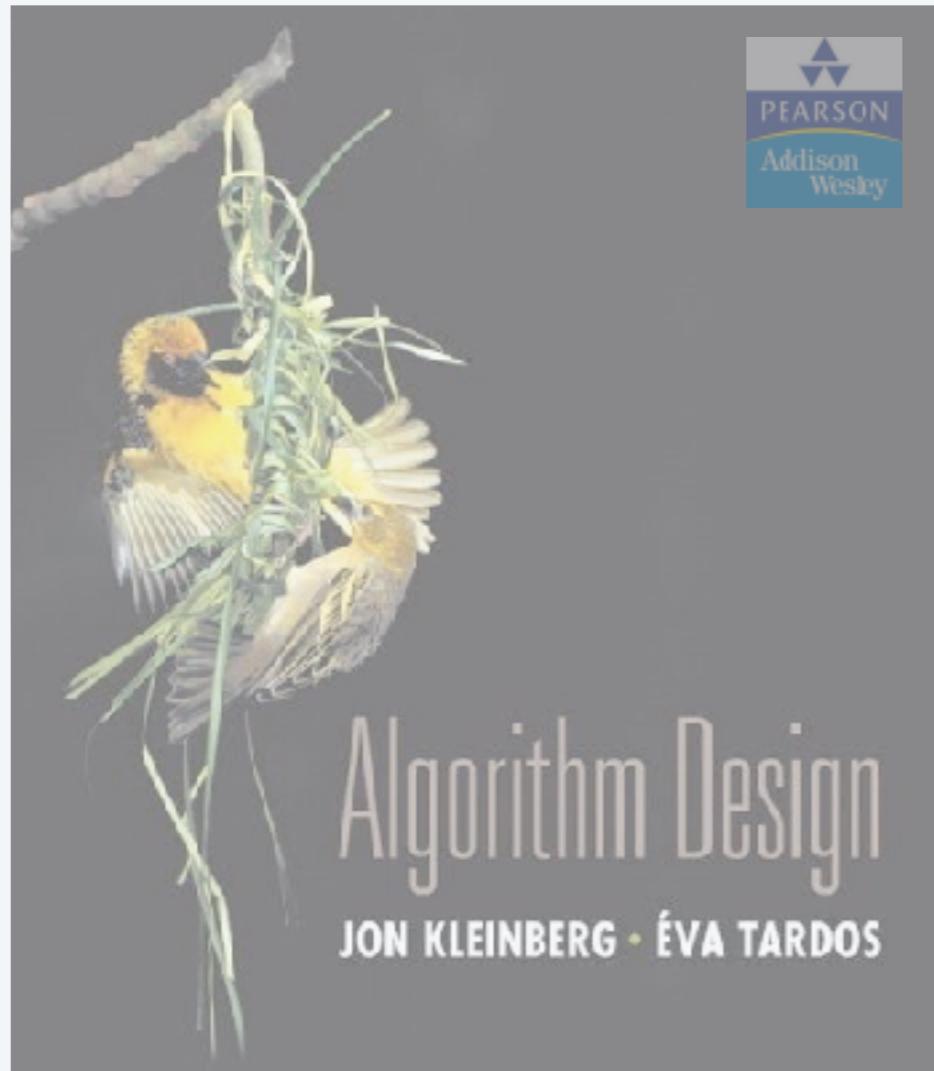
Theorem. [Tarjan 1972] Can find all strong components in $O(m + n)$ time.

SIAM J. COMPUT.
Vol. 1, No. 2, June 1972

DEPTH-FIRST SEARCH AND LINEAR GRAPH ALGORITHMS*

ROBERT TARJAN†

Abstract. The value of depth-first search or “backtracking” as a technique for solving problems is illustrated by two examples. An improved version of an algorithm for finding the strongly connected components of a directed graph and an algorithm for finding the biconnected components of an undirected graph are presented. The space and time requirements of both algorithms are bounded by $k_1 V + k_2 E + k_3$ for some constants k_1, k_2 , and k_3 , where V is the number of vertices and E is the number of edges of the graph being examined.



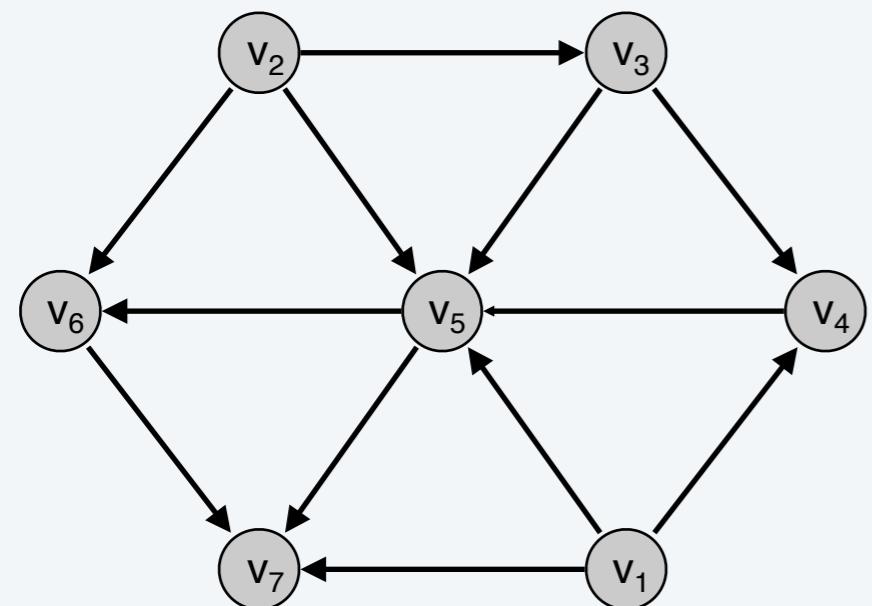
3. GRAPHS

- ▶ *basic definitions and applications*
- ▶ *graph connectivity and graph traversal*
- ▶ *testing bipartiteness*
- ▶ *connectivity in directed graphs*
- ▶ *DAGs and topological ordering*

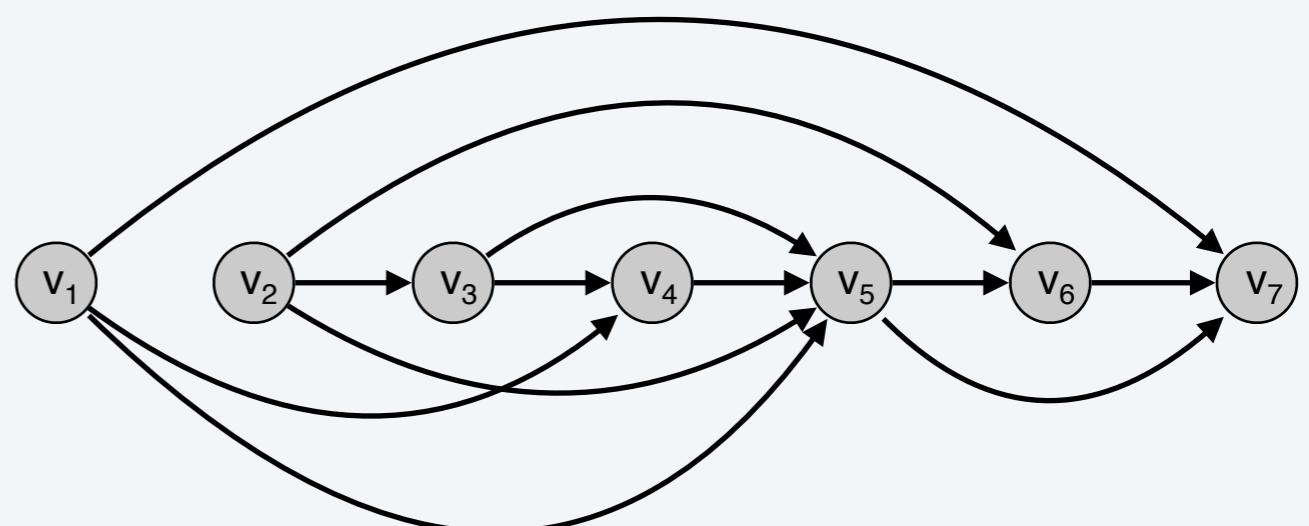
Directed acyclic graphs

Def. A **DAG** is a directed graph that contains no directed cycles.

Def. A **topological order** of a directed graph $G = (V, E)$ is an ordering of its nodes as v_1, v_2, \dots, v_n so that for every edge (v_i, v_j) we have $i < j$.



a DAG



a topological ordering

Precedence constraints

Precedence constraints. Edge (v_i, v_j) means task v_i must occur before v_j .

Applications.

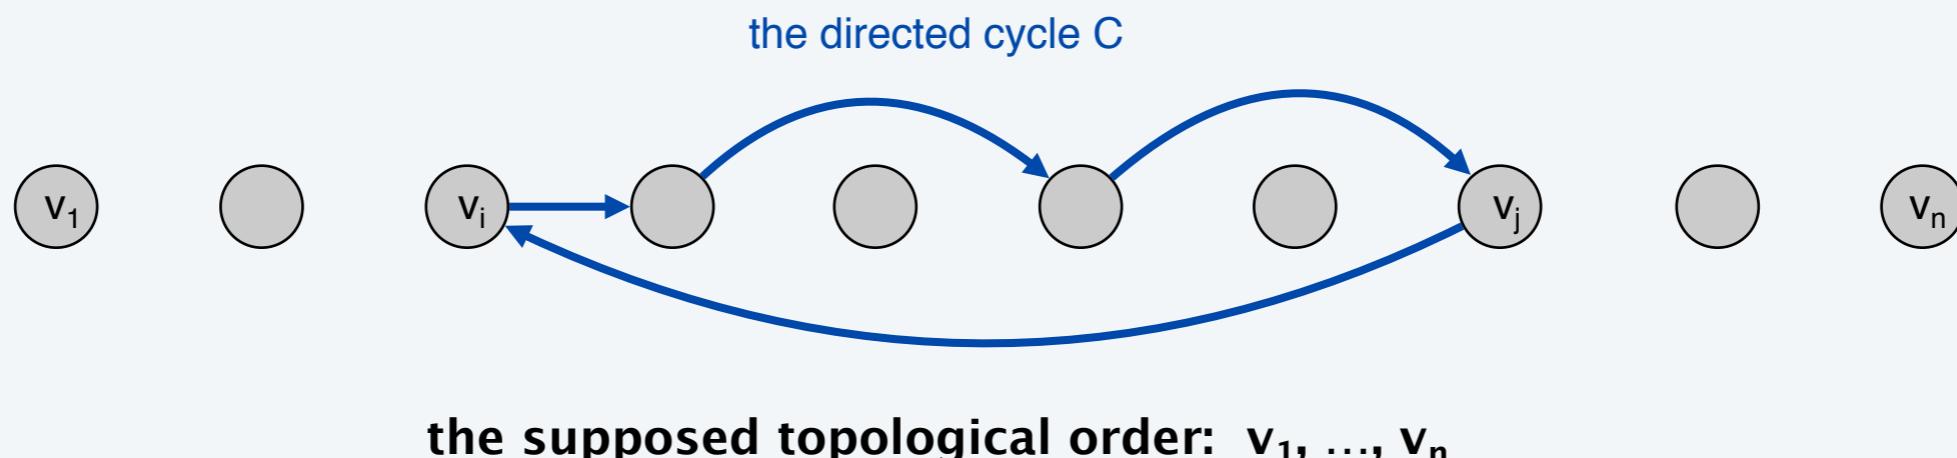
- Course prerequisite graph: course v_i must be taken before v_j .
- Compilation: module v_i must be compiled before v_j .
- Pipeline of computing jobs: output of job v_i needed to determine input of job v_j .

Directed acyclic graphs

Lemma. If G has a topological order, then G is a DAG.

Pf. [by contradiction]

- Suppose that G has a topological order v_1, v_2, \dots, v_n and that G also has a directed cycle C . Let's see what happens.
- Let v_i be the lowest-indexed node in C , and let v_j be the node just before v_i ; thus (v_j, v_i) is an edge.
- By our choice of i , we have $i < j$.
- On the other hand, since (v_j, v_i) is an edge and v_1, v_2, \dots, v_n is a topological order, we must have $j < i$, a contradiction. ■



Directed acyclic graphs

Lemma. If G has a topological order, then G is a DAG.

Q. Does every DAG have a topological ordering?

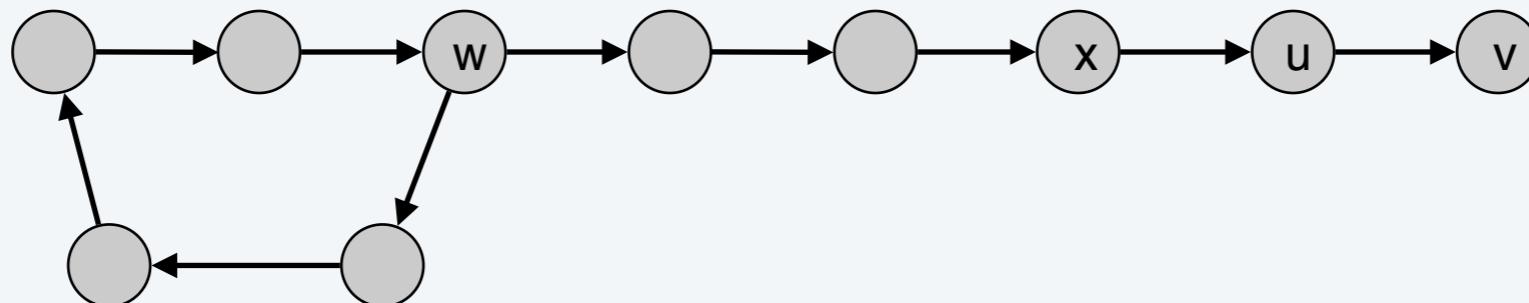
Q. If so, how do we compute one?

Directed acyclic graphs

Lemma. If G is a DAG, then G has a node with no entering edges.

Pf. [by contradiction]

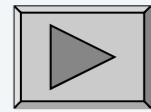
- Suppose that G is a DAG and every node has at least one entering edge. Let's see what happens.
- Pick any node v , and begin following edges backward from v . Since v has at least one entering edge (u, v) we can walk backward to u .
- Then, since u has at least one entering edge (x, u) , we can walk backward to x .
- Repeat until we visit a node, say w , twice.
- Let C denote the sequence of nodes encountered between successive visits to w . C is a cycle. ■



Directed acyclic graphs

Lemma. If G is a DAG, then G has a topological ordering.

Pf. [by induction on n]



- Base case: true if $n = 1$.
- Given DAG on $n > 1$ nodes, find a node v with no entering edges.
- $G - \{v\}$ is a DAG, since deleting v cannot create cycles.
- By inductive hypothesis, $G - \{v\}$ has a topological ordering.
- Place v first in topological ordering; then append nodes of $G - \{v\}$ in topological order. This is valid since v has no entering edges. ▀

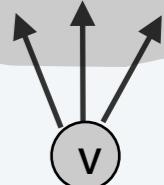
To compute a topological ordering of G :

Find a node v with no incoming edges and order it first

Delete v from G

Recursively compute a topological ordering of $G - \{v\}$
and append this order after v

DAG

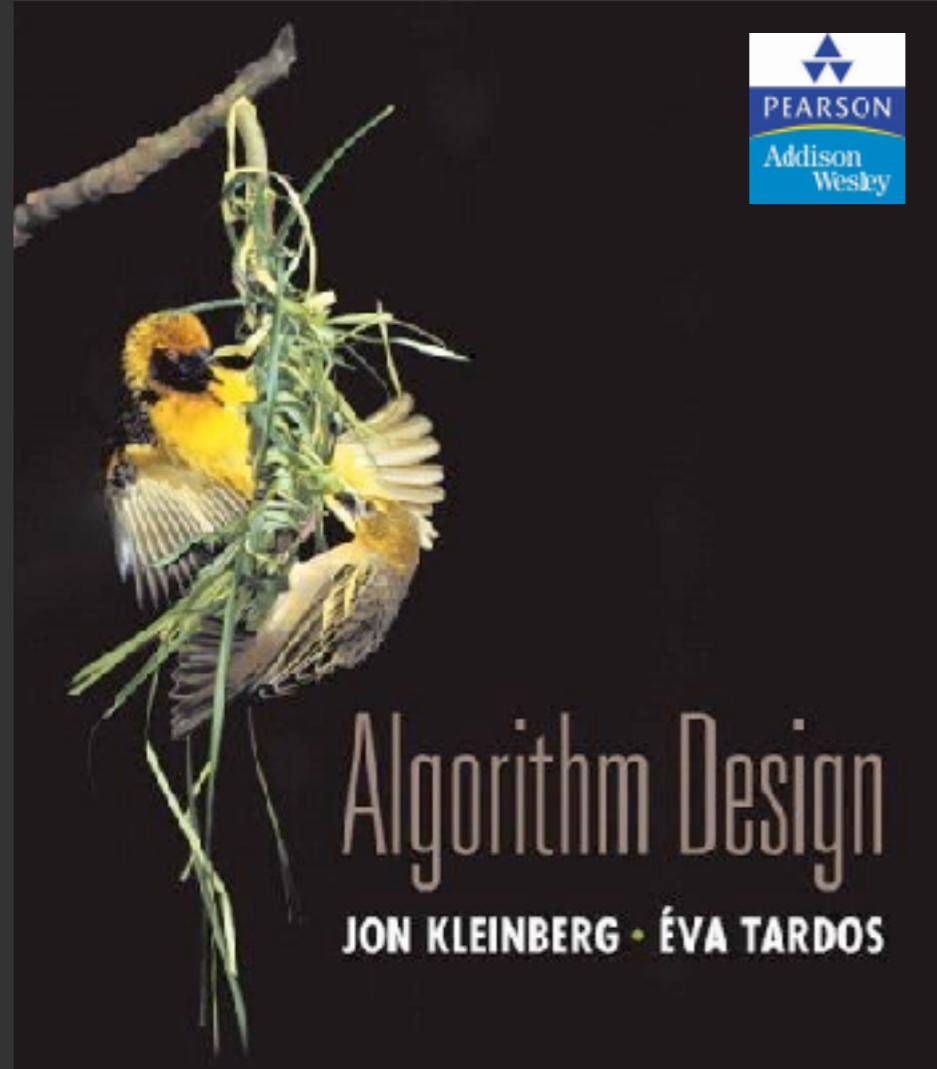


Topological sorting algorithm: running time

Theorem. Algorithm finds a topological order in $O(m + n)$ time.

Pf.

- Maintain the following information:
 - $count(w)$ = remaining number of incoming edges
 - S = set of remaining nodes with no incoming edges
- Initialization: $O(m + n)$ via single scan through graph.
- Update: to delete v
 - remove v from S
 - decrement $count(w)$ for all edges from v to w ;
and add w to S if $count(w)$ hits 0
 - this is $O(1)$ per edge ■



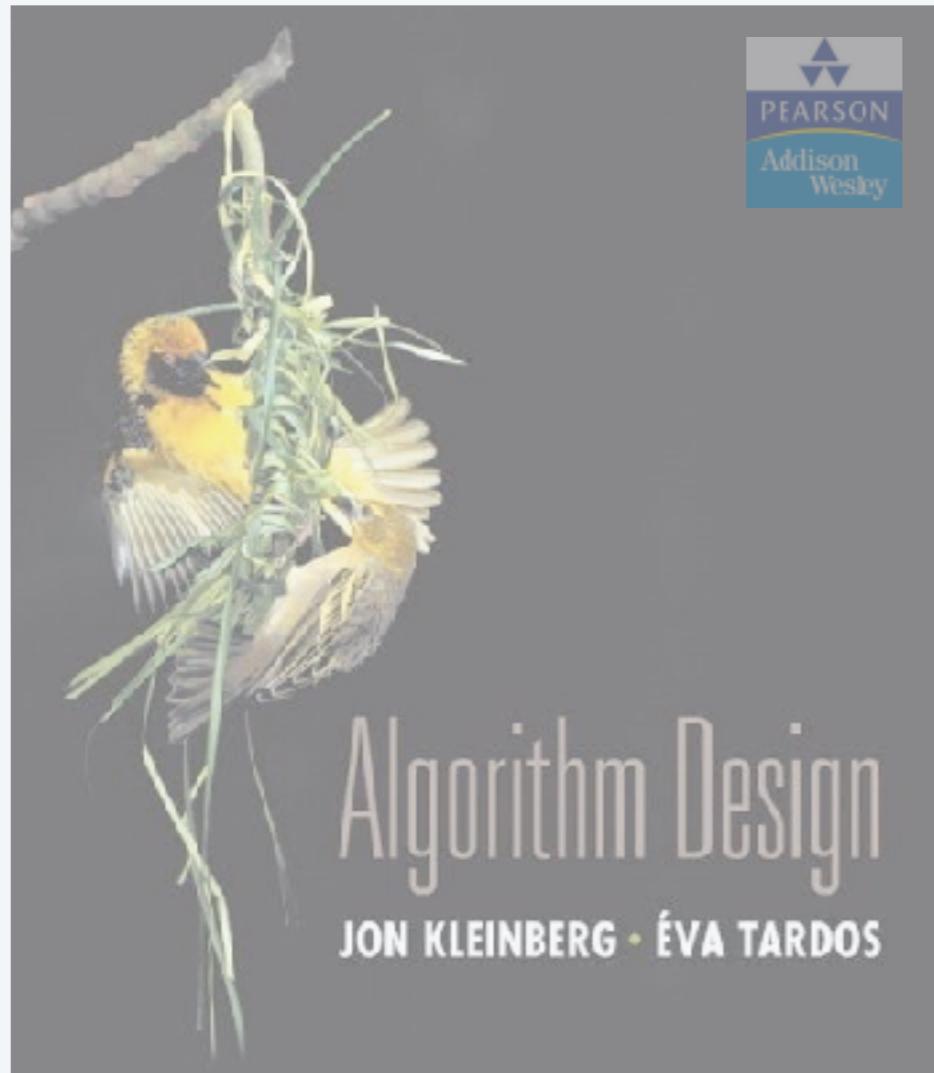
4. GREEDY ALGORITHMS I

- ▶ *coin changing*
- ▶ *interval scheduling and partitioning*
- ▶ *scheduling to minimize lateness*
- ▶ *optimal caching*

Lecture slides by Kevin Wayne

Copyright © 2005 Pearson–Addison Wesley

<http://www.cs.princeton.edu/~wayne/kleinberg-tardos>

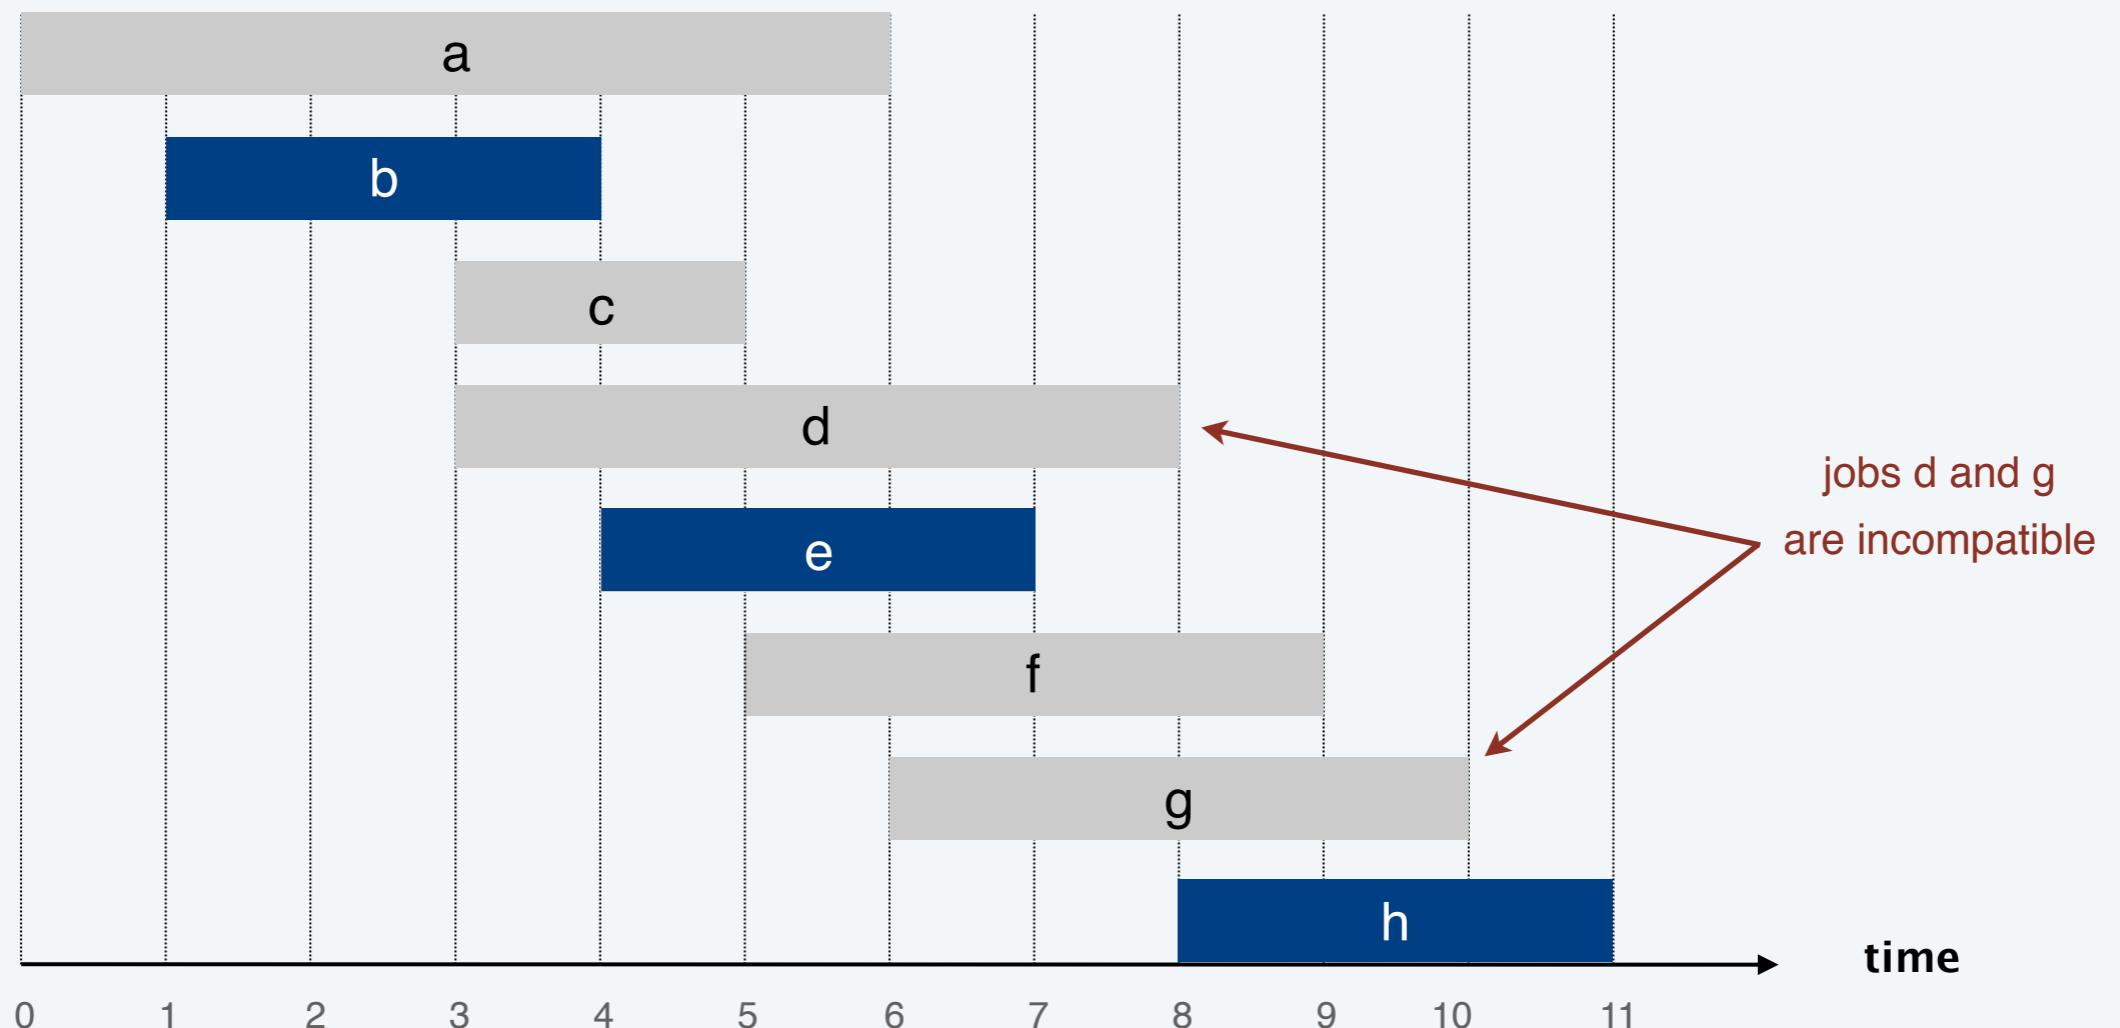


4. GREEDY ALGORITHMS I

- ▶ *coin changing*
- ▶ *interval scheduling and partitioning*
- ▶ *scheduling to minimize lateness*
- ▶ *optimal caching*

Interval scheduling

- Job j starts at s_j and finishes at f_j .
- Two jobs **compatible** if they don't overlap.
- Goal: find maximum subset of mutually compatible jobs.



Greedy algorithms

What makes an algorithm greedy?

Think of what describes the most simple algorithms you know.

Interval scheduling: greedy algorithms

Greedy template. Consider jobs in some natural order.

Take each job provided it's compatible with the ones already taken.

- [Earliest start time] Consider jobs in ascending order of s_j .
- [Earliest finish time] Consider jobs in ascending order of f_j .
- [Shortest interval] Consider jobs in ascending order of $f_j - s_j$.
- [Fewest conflicts] For each job j , count the number of conflicting jobs c_j . Schedule in ascending order of c_j .

Greedy algorithms

What makes an algorithm greedy?

- simple and fast computations
- decision is based on local/few information
- once you make a decision you cannot go back to change it

Interval scheduling: greedy algorithms

Greedy template. Consider jobs in some natural order.

Take each job provided it's compatible with the ones already taken.

- [Earliest start time] Consider jobs in ascending order of s_j .
- [Earliest finish time] Consider jobs in ascending order of f_j .
- [Shortest interval] Consider jobs in ascending order of $f_j - s_j$.
- [Fewest conflicts] For each job j , count the number of conflicting jobs c_j . Schedule in ascending order of c_j .

Interval scheduling: greedy algorithms

Greedy template. Consider jobs in some natural order.

Take each job provided it's compatible with the ones already taken.

counterexample for earliest start time



counterexample for shortest interval



counterexample for fewest conflicts



Interval scheduling: earliest-finish-time-first algorithm



EARLIEST-FINISH-TIME-FIRST ($n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$)

SORT jobs by finish time so that $f_1 \leq f_2 \leq \dots \leq f_n$

$A \leftarrow \emptyset$ ← set of jobs selected

FOR $j = 1$ TO n

IF job j is compatible with A

$A \leftarrow A \cup \{ j \}$

RETURN A

Proposition. Can implement earliest-finish-time first in $O(n \log n)$ time.

- Keep track of job j^* that was added last to A .
- Job j is compatible with A iff $s_j \geq f_{j^*}$.
- Sorting by finish time takes $O(n \log n)$ time.

Interval scheduling: analysis of earliest-finish-time-first algorithm

Theorem. The earliest-finish-time-first algorithm is optimal.

Pf. What do we need to prove?

What is the output of the greedy algorithm?

What does it mean that it is optimal?

Pf.

- our solution S may or may not be optimal, so let's pick a solution S^* that we know **is optimal**. (what does that mean again?)
- there may be multiple optimal solutions, so let's choose S^* to be the one **most similar** to S . (what makes it most similar?)
- by **swapping two jobs** we can make S^* even **more similar** to S in **contradiction** with the choice of S^* being most similar.

Interval scheduling: analysis of earliest-finish-time-first algorithm

Theorem. The earliest-finish-time-first algorithm is optimal

Pf.

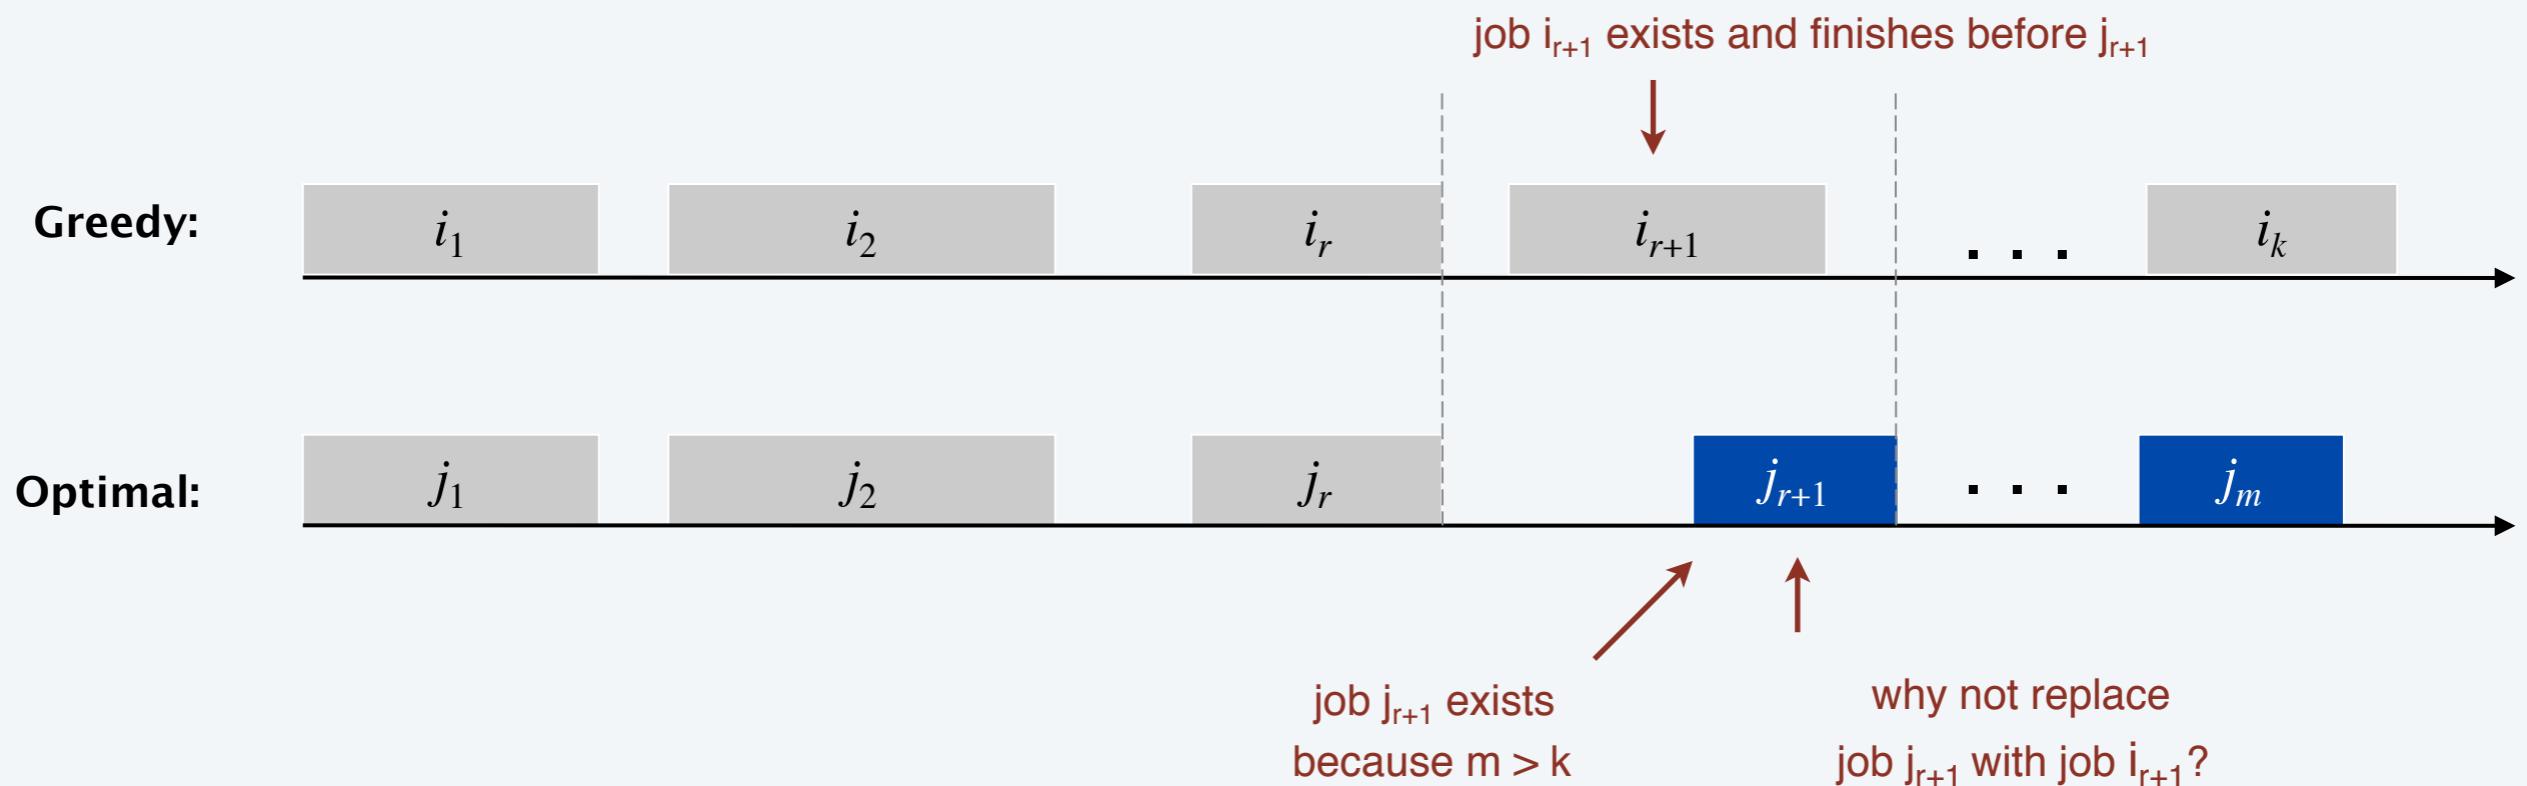
- our solution S may or may not be optimal, so let's pick a solution S^* that we know is optimal. (what does that mean again?)
- there may be multiple optimal solutions, so let's choose S^* to be the one most similar to S . (what makes it most similar?)
- by swapping two jobs we can make S^* even more similar to S in contradiction with the choice of S^* being most similar.

Interval scheduling: analysis of earliest-finish-time-first algorithm

Theorem. The earliest-finish-time-first algorithm is optimal

Pf.

- our solution S may or may not be optimal, so let's pick a solution S^* that we know is optimal. (what does that mean again?)
- there may be multiple optimal solutions, so let's choose S^* to be the one most similar to S . (what makes it most similar?)
- by swapping two jobs we can make S^* even more similar to S in contradiction with the choice of S^* being most similar.

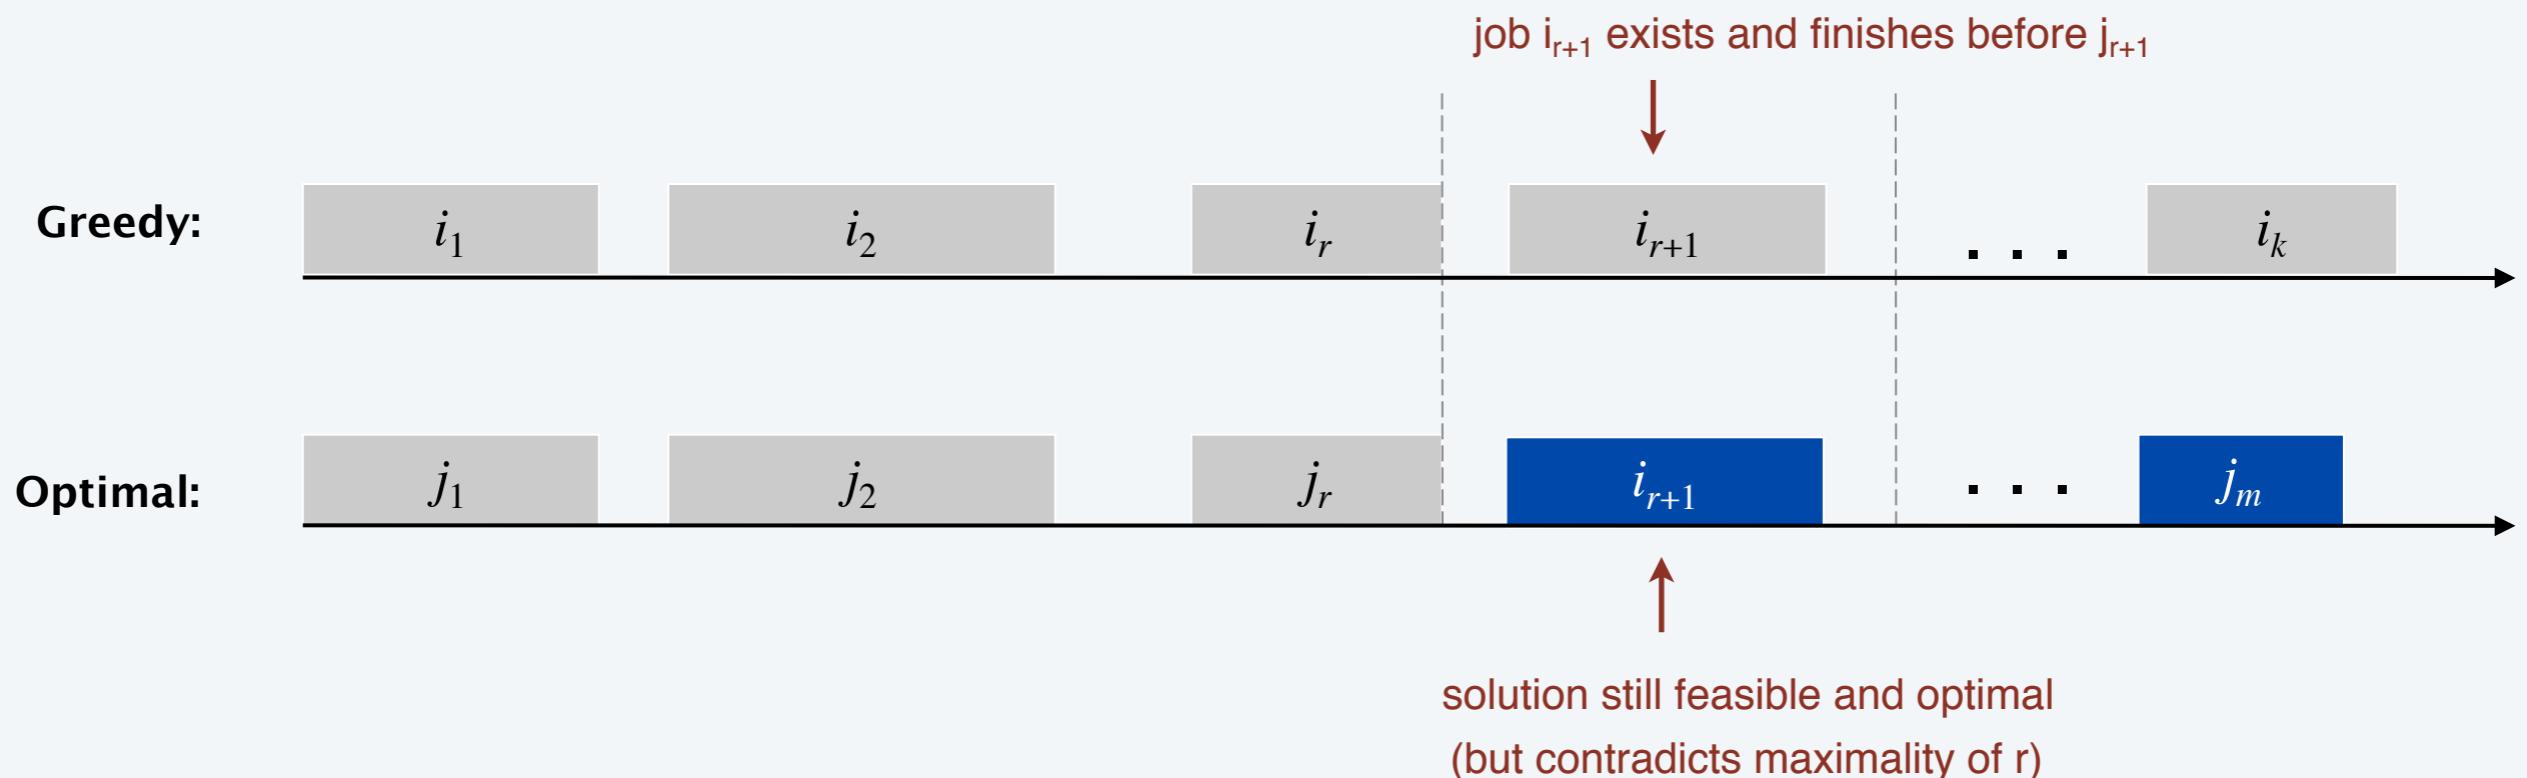


Interval scheduling: analysis of earliest-finish-time-first algorithm

Theorem. The earliest-finish-time-first algorithm is optimal

Pf.

- our solution S may or may not be optimal, so let's pick a solution S^* that we know is optimal. (what does that mean again?)
- there may be multiple optimal solutions, so let's choose S^* to be the one most similar to S . (what makes it most similar?)
- by swapping two jobs we can make S^* even more similar to S in contradiction with the choice of S^* being most similar.

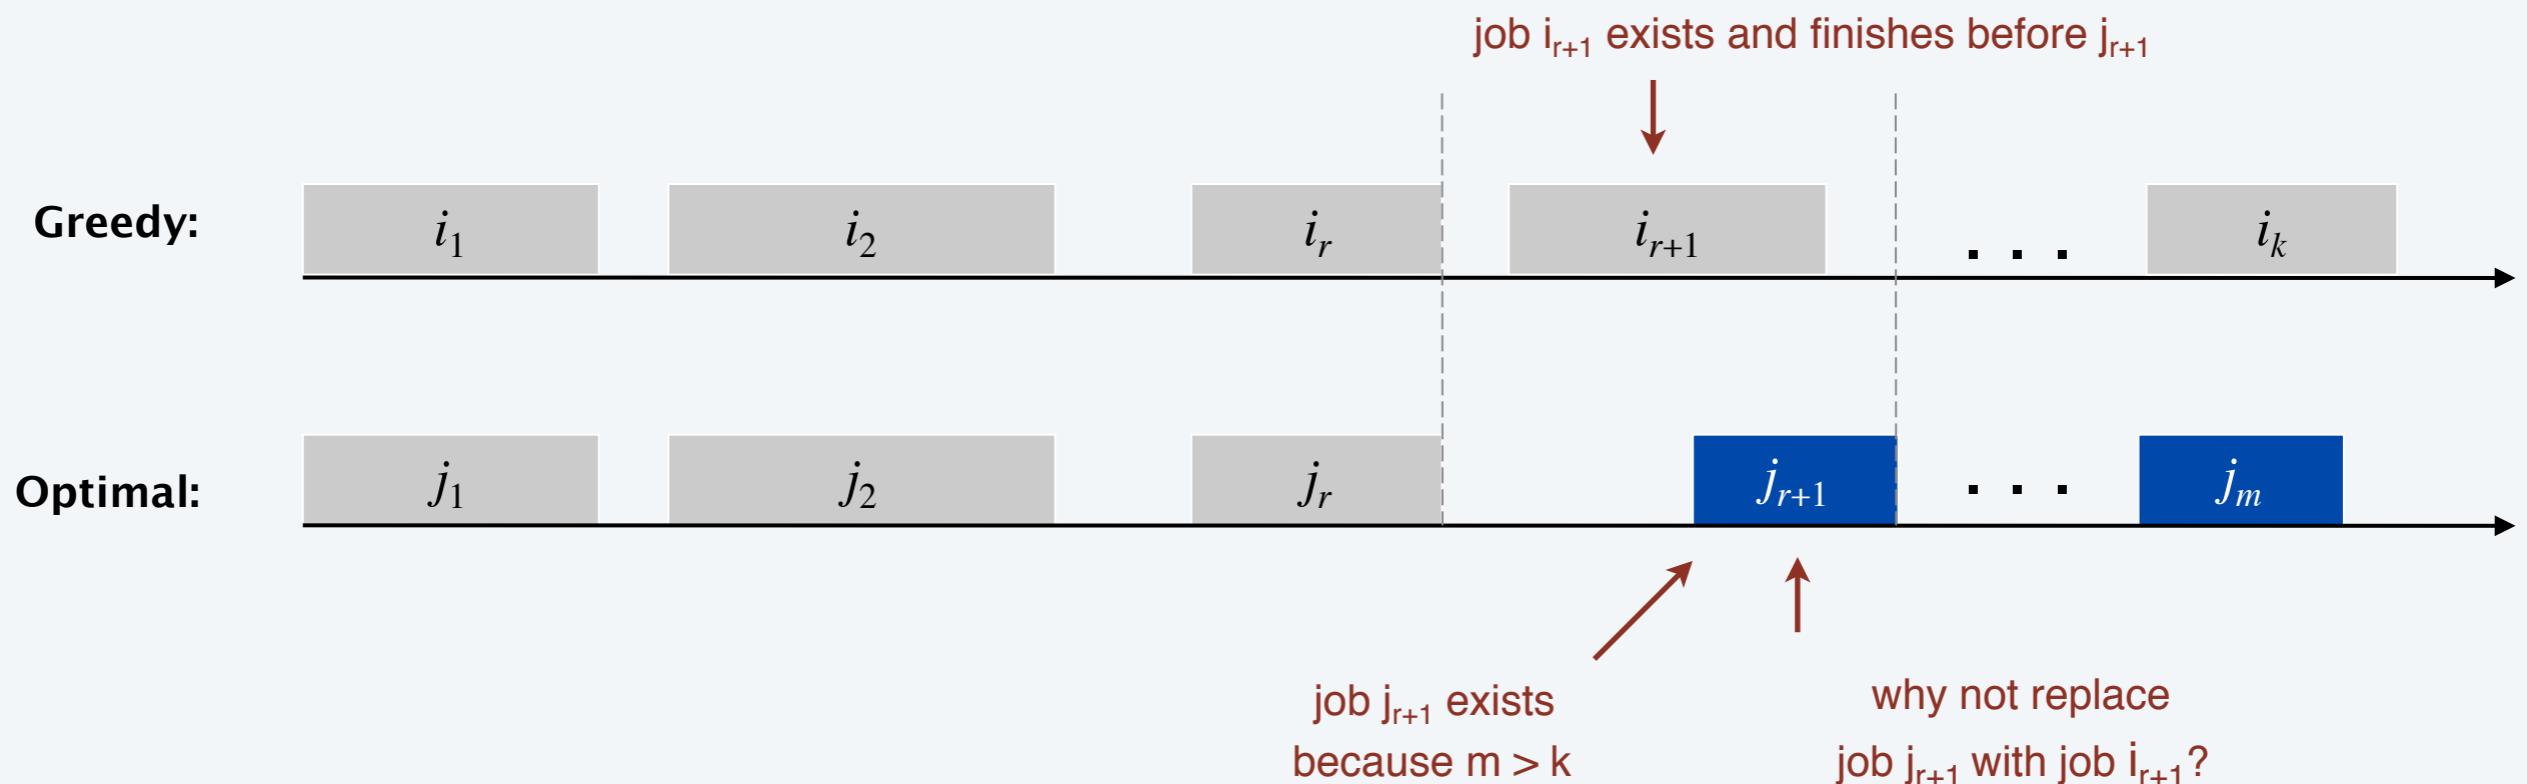


Interval scheduling: analysis of earliest-finish-time-first algorithm

Theorem. The earliest-finish-time-first algorithm is optimal.

Pf. [by contradiction]

- Assume greedy is not optimal, and let's see what happens.
- Let i_1, i_2, \dots, i_k denote set of jobs selected by greedy.
- Let j_1, j_2, \dots, j_m denote set of jobs in an optimal solution with $i_1 = j_1, i_2 = j_2, \dots, i_r = j_r$ for the largest possible value of r .

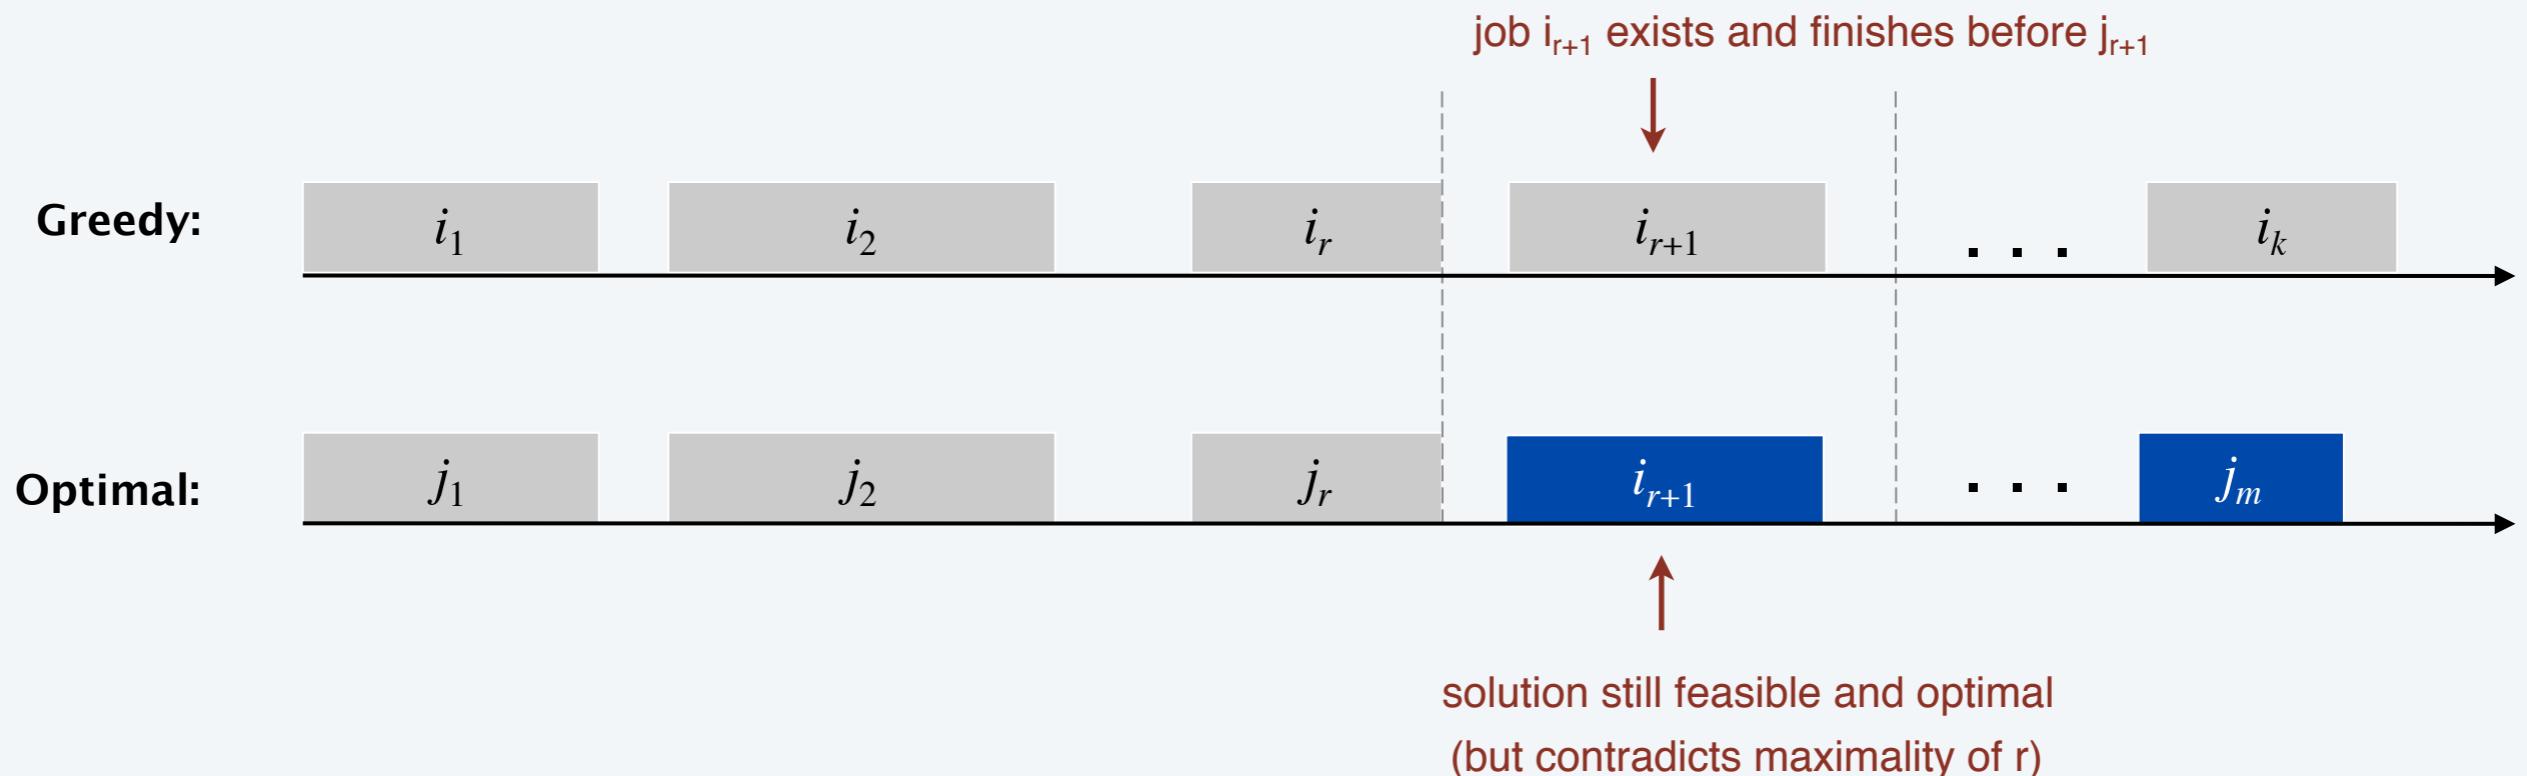


Interval scheduling: analysis of earliest-finish-time-first algorithm

Theorem. The earliest-finish-time-first algorithm is optimal.

Pf. [by contradiction]

- Assume greedy is not optimal, and let's see what happens.
- Let i_1, i_2, \dots, i_k denote set of jobs selected by greedy.
- Let j_1, j_2, \dots, j_m denote set of jobs in an optimal solution with $i_1 = j_1, i_2 = j_2, \dots, i_r = j_r$ for the largest possible value of r .

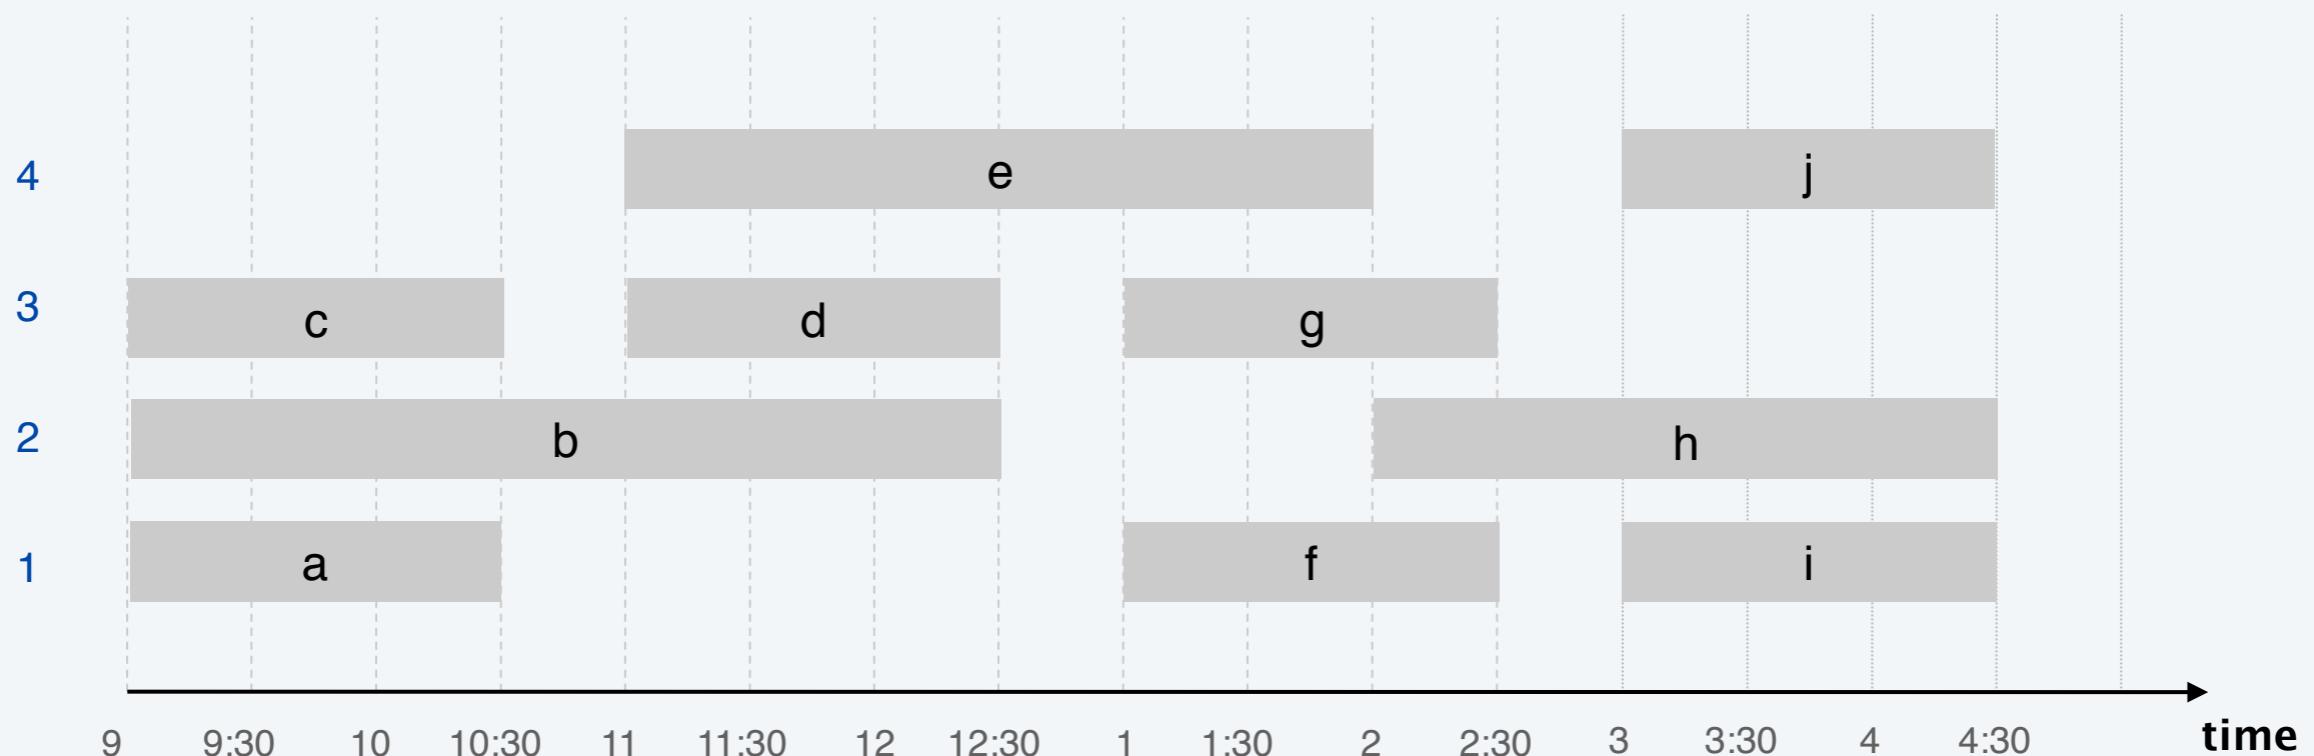


Interval partitioning

Interval partitioning.

- Lecture j starts at s_j and finishes at f_j .
- Goal: find minimum number of classrooms to schedule all lectures so that no two lectures occur at the same time in the same room.

Ex. This schedule uses 4 classrooms to schedule 10 lectures.

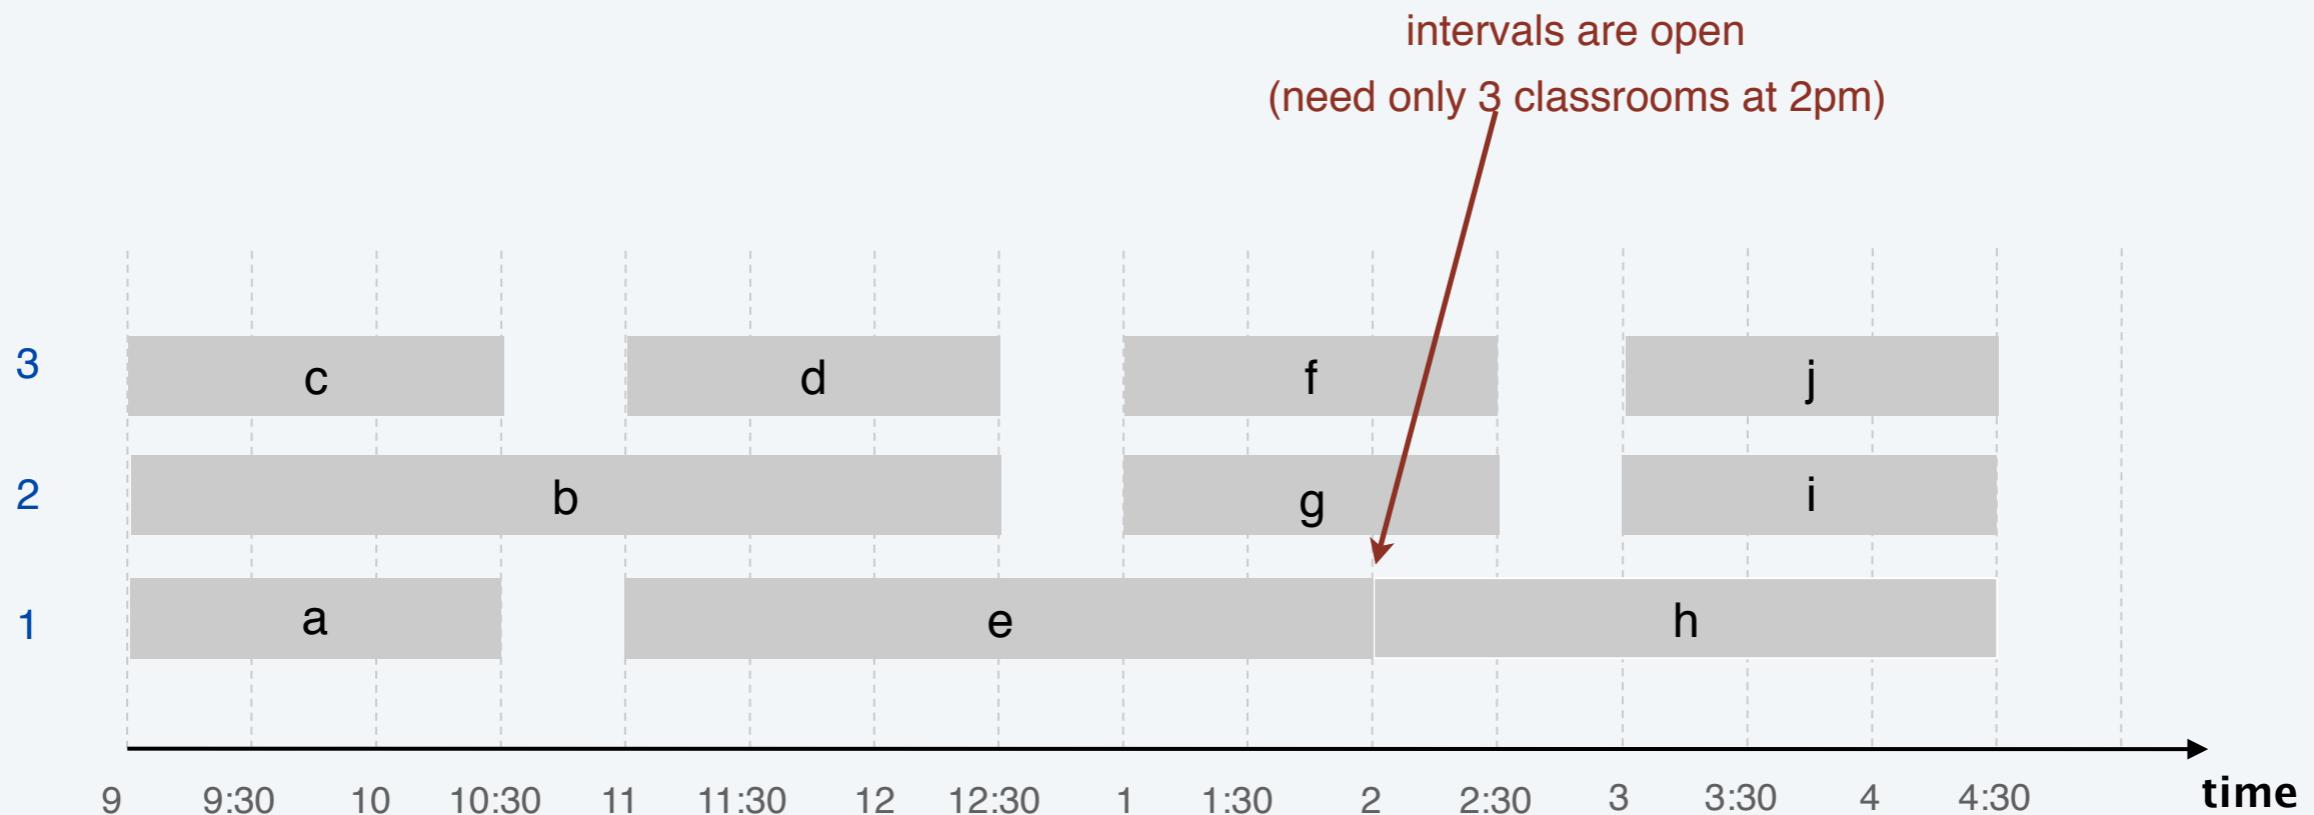


Interval partitioning

Interval partitioning.

- Lecture j starts at s_j and finishes at f_j .
- Goal: find minimum number of classrooms to schedule all lectures so that no two lectures occur at the same time in the same room.

Ex. This schedule uses **3** classrooms to schedule 10 lectures.



Greedy algorithms

What makes an algorithm greedy?

- simple and fast computations
- decision is based on local/few information
- once you make a decision you cannot go back to change it

Interval partitioning: greedy algorithms

Greedy template. Consider lectures in some natural order.

Assign each lecture to an available classroom (which one?);
allocate a new classroom if none are available.

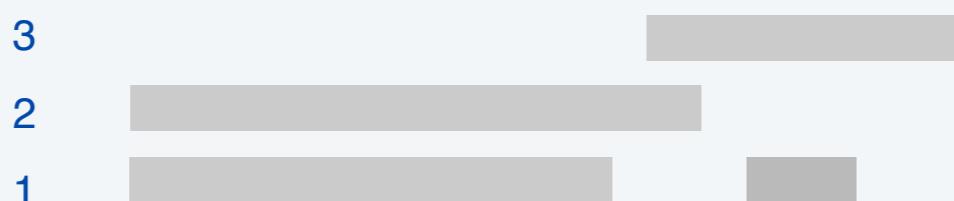
- [Earliest start time] Consider lectures in ascending order of s_j .
- [Earliest finish time] Consider lectures in ascending order of f_j .
- [Shortest interval] Consider lectures in ascending order of $f_j - s_j$.
- [Fewest conflicts] For each lecture j , count the number of
conflicting lectures c_j . Schedule in ascending order of c_j .

Interval partitioning: greedy algorithms

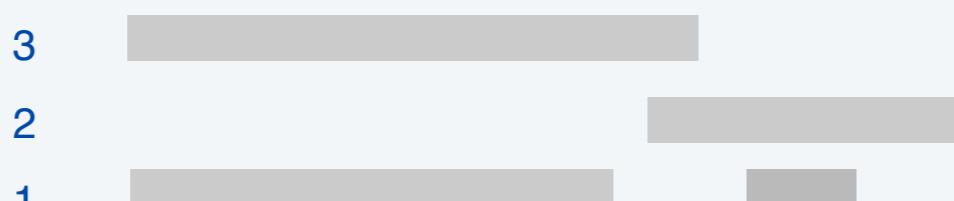
Greedy template. Consider lectures in some natural order.

Assign each lecture to an available classroom (which one?);
allocate a new classroom if none are available.

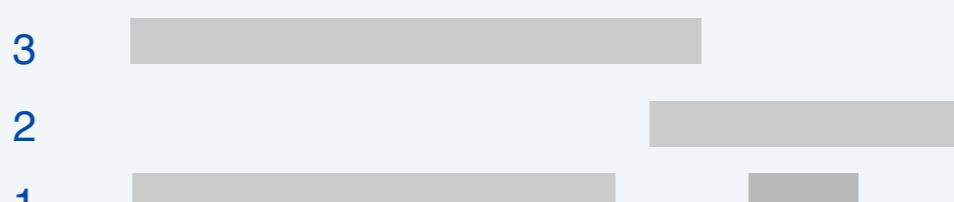
counterexample for earliest finish time



counterexample for shortest interval



counterexample for fewest conflicts



Interval partitioning: earliest-start-time-first algorithm



EARLIEST-START-TIME-FIRST ($n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$)

SORT lectures by start time so that $s_1 \leq s_2 \leq \dots \leq s_n$.

$d \leftarrow 0$ ← number of allocated classrooms

FOR $j = 1$ *TO* n

IF lecture j is compatible with some classroom

Schedule lecture j in any such classroom k.

ELSE

Allocate a new classroom $d + 1$.

Schedule lecture j in classroom $d + 1$.

$d \leftarrow d + 1$

RETURN schedule.

Interval partitioning: earliest-start-time-first algorithm

Proposition. The earliest-start-time-first algorithm can be implemented in $O(n \log n)$ time.

Pf. Store classrooms in a **priority queue** (key = finish time of its last lecture).

- To determine whether lecture j is compatible with some classroom, compare s_j to key of min classroom k in priority queue.
- To add lecture j to classroom k , increase key of classroom k to f_j .
- Total number of priority queue operations is $O(n)$.
- Sorting by start time takes $O(n \log n)$ time. ▀

Remark. This implementation chooses a classroom k whose finish time of its last lecture is the **earliest**.

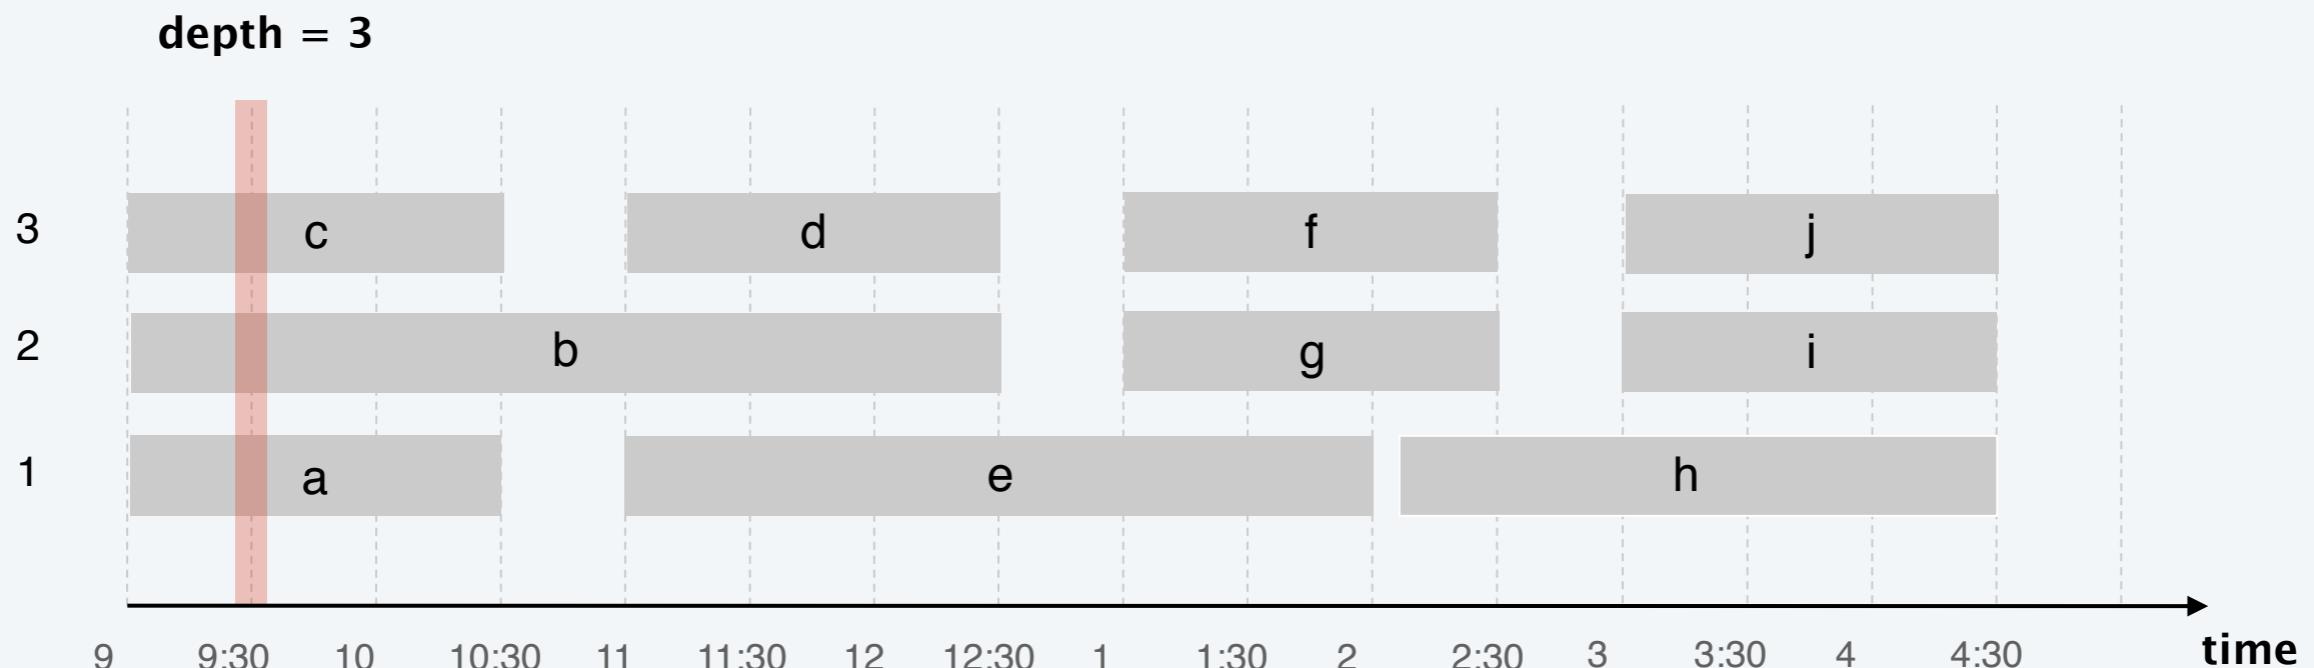
Interval partitioning: lower bound on optimal solution

Def. The **depth** of a set of open intervals is the maximum number of intervals that contain any given time.

Key observation. Number of classrooms needed \geq depth.

Q. Does minimum number of classrooms needed always equal depth?

A. Yes! Moreover, earliest-start-time-first algorithm finds a schedule whose number of classrooms equals the depth.



Interval partitioning: analysis of earliest-start-time-first algorithm

Observation. The earliest-start-time first algorithm never schedules two incompatible lectures in the same classroom.

Theorem. Earliest-start-time-first algorithm is optimal.

Pf.

- Let d = number of classrooms that the algorithm allocates.
- Classroom d is opened because we needed to schedule a lecture, say j , that is incompatible with all $d - 1$ other classrooms.
- These d lectures each end after s_j .
- Since we sorted by start time, all these incompatibilities are caused by lectures that start no later than s_j .
- Thus, we have d lectures overlapping at time $s_j + \varepsilon$.
- Key observation \Rightarrow all schedules use $\geq d$ classrooms. ▀